Training-Free Cultural Alignment of Large Language Models via Persona Disagreement

Anonymous Authors
anonymous@institution.edu

Abstract

Large language models make moral judgments as if the whole world thinks like the West, yet 40 million human responses across 233 countries say otherwise [Awad et al., 2018]. We introduce **SWA-DPBR** (Socially-Weighted Alignment with Dual-Pass Bootstrap Reliability), a method that adapts a frozen LLM to the moral preferences of a specific country at inference time, without modifying any weights or requiring white-box access to internal activations. Three design choices distinguish SWA-DPBR from generic decoding-time steering, each motivated by a concrete property of the cultural-alignment problem. *First*, we treat *within-country disagreement as the signal itself*: culture is not monolithic but stratified by age cohort and social position [Inglehart and Welzel, 2005, Awad et al., 2018], so we instantiate each country as a small panel of World Values Survey (WVS)–grounded persona agents whose disagreement pinpoints where and by how much to correct the model's default. *Second*, we aggregate that disagreement through a *loss-averse* Prospect-Theory value function [Kahneman and Tversky, 1979, Tversky and Kahneman, 1992]: a cultural correction that improves three personas while harming one is penalised more than one that helps all four modestly, operationalising a "do no harm" prior that deterministic averaging cannot express—and echoing risk-asymmetries recently shown to emerge inside LLMs themselves [Payne, 2025]. *Third*, we gate the importance-sampled correction with a *dual-pass bootstrap reliability* weight [Efron, 1979]: two independent Monte-Carlo replicates must agree before the method commits, so the correction shrinks smoothly toward zero on scenarios where the logit landscape is noisy. Evaluated on MultiTP across **twenty countries, four continents**, and six open model families (3.8B–70B parameters), SWA-DPBR reduces misalignment by **19–24%** on the strongest checkpoints, with gains spanning the Americas, Europe, and Asia rather than concentrating in any single region. When the method fails, the culprit is architectural—poorly conditioned decision logits rather than cultural blind spots—suggesting that inference-time cultural steering is broadly viable wherever the base model maintains a well-behaved decision surface. Code and experimental artifacts will be released publicly upon acceptance.

1 Introduction

A self-driving car must choose: spare a group of elderly men or a group of young women. Ask in Tokyo and Bogotá and you will get systematically different answers. The Moral Machine experiment [Awad et al., 2018] demonstrated this at scale: 40 million judgments from 233 countries revealed that moral preferences are not noise—they are *structured* by culture. East Asian respondents weight youth more heavily; Southern-hemisphere populations show distinct utilitarian calculi; Islamic societies prioritise different social roles. Yet the large language models increasingly enlisted to mediate such decisions encode a single, culturally monolithic moral stance inherited from predominantly Western training data [Henrich et al., 2010]. The multilingual MultiTP benchmark [Jin et al., 2025]

Preprint.

now quantifies the scale of this gap: country-specific Average Marginal Component Effects (AMCE) across six moral dimensions and 107 languages show that *none* of 19 contemporary LLMs achieves broadly satisfactory alignment, with most exhibiting WEIRD-biased reasoning that departs sharply from non-WEIRD human preferences.

Closing this gap without retraining is hard because a practical cross-cultural decoder must satisfy three constraints at once: no weight updates, no per-culture supervised labels at scale, and no white-box access to internal activations. Every existing approach trades at least one away-fine-tuning [Bai et al., 2022, Ouyang et al., 2022] changes weights and needs per-culture data; prompting-based steering lacks empirical grounding [Jin et al., 2025]; activation steering [Arditi et al., 2025, Zou et al., 2023] requires internal-layer access; reward-guided decoding [Khanov et al., 2024, Mudgal et al., 2024] assumes per-country reward models that do not exist; and culture-aware adapters [Zhang et al., 2026, Yao et al., 2025] need per-culture training, scaling poorly past a handful of countries.

We propose **SWA-DPBR** (Socially-Weighted Alignment with Dual-Pass Bootstrap Reliability), a method that satisfies all three constraints by exploiting a simple but powerful insight: *within-country disagreement is the steering signal itself.* Moral preferences within a society are not monolithic-age cohorts and social strata prioritise fundamentally different trade-offs [Awad et al., 2018, Kleiman-Weiner et al., 2017], and the World Values Survey [Haerpfer et al., 2022] provides empirical structure for that variation. SWA-DPBR instantiates each country as a small panel of culturally grounded persona agents drawn from WVS microdata; when those agents disagree on a dilemma, the spread of their logit gaps pinpoints where a correction is needed and how large it should be, and when they agree, the vanilla model is typically already close and no correction is applied. Operationally, the pipeline adds a single batched forward pass (base model plus personas in the native language), applies positional debiasing, then runs two independent Prospect-Theory importance-sampling passes [Kahneman and Tversky, 1979] whose outputs are merged with a bootstrap reliability weight that automatically shrinks corrections when the two passes disagree. An effective-sample-size guard and a light hierarchical country prior complete the stack (§3; hyperparameters in Table 8).

Each component of SWA-DPBR answers a specific failure mode of simpler alternatives: a panel rather than a single country prompt preserves the within-country spread that the correction stage consumes; Prospect Theory [Kahneman and Tversky, 1979, Tversky and Kahneman, 1992] supplies the loss-averse asymmetry that a linear or quadratic aggregator cannot express; and the dual-pass bootstrap turns the same IS budget into a direct reliability probe [Efron, 1979, Elvira and Martino, 2021] rather than hiding estimator bias behind more samples. A concise decision tree through these choices is deferred to Appendix A.

Contributions.

- **A training-free cultural alignment method.** SWA-DPBR aligns a *frozen* LLM to country-level human preference vectors at inference time: WVS-grounded personas supply structured disagreement; Prospect-Theory IS converts it to a bounded logit correction; dual-pass bootstrap reliability gates noisy updates; a hierarchical prior smooths across scenarios (§3).
- **Broad empirical validation.** On MultiTP across twenty countries, four continents, and six model families (3.8B–70B), SWA-DPBR cuts misalignment by 19–24% on the strongest checkpoints, with gains that span geographies rather than concentrating in one region (§5).
- **Diagnostic ablations.** We show that persona diversity, positional debiasing, and the PT kernel are each load-bearing; that temperature and preprocessing sweeps leave conclusions stable; and that failure traces to collapsed logit entropy, not cultural blind spots (§6; Appendices H–I).

2 Related Work

Benchmarks and cultural evaluation. The Moral Machine experiment [Awad et al., 2018] established that moral preferences are culturally structured; MultiTP [Jin et al., 2025] extends this to 107 languages with country-specific AMCEs, providing the evaluation bed we adopt. Existing moral benchmarks (ETHICS [Hendrycks et al., 2021], MoralExceptQA [Jin et al., 2022], MoCa [Nie et al., 2023]) remain English-centric, and broader diversity efforts [Kirk et al., 2024, Durmus et al., 2023, Qiu et al., 2025] cover fewer languages or different modalities. We validate on MultiTP's binary

trolley vignettes where AMCE vectors provide a clean six-dimensional surface; a pilot extension to the open-ended BLEnD benchmark [Myung et al., 2024] (Appendix AE) reveals that the binary logit-gap formulation does not transfer directly to free-form generation, and no existing training-free method has demonstrated country-level calibration on such tasks. On the diagnostic side, LLM outputs correlate most strongly with WEIRD respondents [Santurkar et al., 2023], RLHF introduces unintended cultural shifts [Ryan et al., 2024], and recent work shows that LLMs systematically stereotype the moral values of non-Western populations [Zewail et al., 2026]-underscoring the practical need for culturally grounded correction mechanisms.

Inference-time alignment. The pluralistic alignment vision [Sorensen et al., 2024] argues that no single preference ordering serves all communities; inference-time methods are a natural vehicle because they avoid baking a fixed value stance into weights. However, existing approaches each sacrifice one of the three constraints we target: activation steering [Arditi et al., 2025, Zou et al., 2023] requires white-box access; reward-guided decoding [Khanov et al., 2024, Mudgal et al., 2024] needs per-country reward models; and culture-aware adapters [Zhang et al., 2026, Yao et al., 2025] need per-culture training. Recent logit-level steering work [An et al., 2026] and culturally localised generation via sparse autoencoders [Khanuja et al., 2026] demonstrate that the logit surface is a richer control interface than previously appreciated, but neither targets country-conditioned moral calibration. Inference-time stochasticity can improve moral alignment by reshaping confidence [Kwon et al., 2026]; SWA-DPBR complements this by targeting *directional*, survey-grounded corrections rather than generic uncertainty inflation, and we include it as an explicit baseline (Appendix M). Black-box alignment modules that learn small guidance policies from preference data [Bobbili et al., 2025, Chen et al., 2025] sit in an adjacent regime: they operate on frozen model outputs, but still require a training stage for the guidance policy and are typically defined over sentence-level continuations rather than binary decisions. Adaptive importance sampling in pre-logit space [Kanai et al., 2025] achieves similar variance control but requires white-box activation access; we work strictly with decision-token logits. Training-time prospect-theoretic alignment [Ethayarajh et al., 2024] extends KL-regularised preference learning with loss-aversion via a Kahneman–Tversky value function; SWA-DPBR applies the same value shape at *test time*, per scenario, so the asymmetry aggregates heterogeneous personas rather than shaping a trained policy. To the best of our knowledge, SWA-DPBR is the first training-free method targeting country-level moral calibration specifically.

Theoretical grounding. SWA-DPBR draws on three established traditions: control as probabilistic inference [Levine, 2018, Todorov, 2009, Williams et al., 2018] for the IS softmax update, Prospect Theory [Kahneman and Tversky, 1979, Tversky and Kahneman, 1992, Payne, 2025] for the loss-averse utility shape, and self-normalised IS theory [Geweke, 1989, Owen, 2013, Elvira and Martino, 2021] for the finite-K bias-variance characterisation that our dual-pass gate inherits. We treat each as principled motivation rather than formal guarantee in our scalar-logit setting; the boxed Proposition in §3.5 states what survives, and Appendices A–B discuss each connection in full.

3 Method: SWA-DPBR

The pipeline has three logical stages: extract cultural signal from a WVS-grounded persona panel (§3.2–§3.4), convert dispersion into a bounded correction via Prospect-Theory IS with dual-pass reliability (§3.5), and stabilise across scenarios via a hierarchical country prior (§3.6). Algorithm 1 ties the three together; rationale for each choice is summarised in Appendix A.

3.1 Problem Formulation

Let $\mathcal{X}$ be the set of moral dilemma scenarios, each presenting a forced binary choice via decision tokens `A` (a) and `B` (b). For target country c, let $\mathbf{h}^{(c)} \in \mathbb{R}^6$ be the human AMCE vector from 40M+ judgments [Awad et al., 2018]. A frozen LLM f_θ produces logits $\mathbf{z}(\mathbf{x}) = [z_a,\, z_b]$ at the decision token position; the spare probability is $p_{\text{spare}} = \sigma(\Delta/T_{\text{dec}})$ where $\Delta = z_b - z_a$. Each dimension's AMCE is estimated as:

$$\hat{m}_d^{(c)} = \frac{1}{|\mathcal{S}_d^{(c)}|} \sum_{\mathbf{x} \in \mathcal{S}_d^{(c)}} p_{\text{spare}}(\mathbf{x}), \tag{1}$$

following the MultiTP evaluation convention [Jin et al., 2025]. All six dimensions-including Utilitarianism-are computed uniformly as the empirical mean of p_{spare} (Appendix C).

Goal: Without modifying θ, find a logit correction $\delta^*(\mathbf{x}, c)$ such that the corrected AMCE vector $\hat{\mathbf{m}}^{(c)}(\theta, \delta^*)$ is closer to $\mathbf{h}^{(c)}$. We primarily measure alignment with the *misalignment score* (MIS): ℓ_2 distance between $\hat{\mathbf{m}}^{(c)}$ and $\mathbf{h}^{(c)}$ (lower is better). Jensen–Shannon divergence and Pearson r are reported as secondary diagnostics where noted.

3.2 Cultural Persona Construction

We use $N{=}4$ personas—three age cohorts plus a country-wide aggregate—grounded in WVS Wave 7 individual-level microdata [Haerpfer et al., 2022]. The choice of $N{=}4$ is the smallest panel that covers the primary demographic axis of within-country variance while preserving a population anchor against cohort-specific noise [Inglehart and Welzel, 2005, Kish, 1965]; full justification is in Appendix A.

For each country c, following the 10-variable cultural-value scheme of Greco et al. [2026], we extract: *religiosity* (Q6P), *child-rearing values* (Q17P), *moral acceptability* (Q177–Q182), *social trust* (Q57P), *political participation* (Q199P, Q200P), *national pride* (Q254P), *happiness* (Q46P), *gender equality* (Q29P–Q33P), *materialism orientation* (Q152–Q153), and *tolerance for diversity* (Q19P–Q23P). Items from the "inverted" WVS-7 CSV (suffix P) are directionally coded so that higher values consistently indicate the positive pole of each dimension. For each country, respondents are segmented into three age cohorts—*young* (<36), *middle-aged* (36–55), and *older* (>55)—yielding three cohort-specific personas that capture generational variation in cultural values. A fourth persona aggregates all respondents regardless of age, anchoring the ensemble with the country's population-wide cultural centre of gravity. All four personas are thus data-driven and country-specific; no agent carries a hard-coded moral stance that could bias particular dimensions. Dimension means are normalised to $[0, 1]$ and mapped to four-level natural-language descriptors via quartile cuts (≥ 0.75, ≥ 0.50, ≥ 0.25, < 0.25); see Appendix D. Each persona is a native-language system prompt assembled from a header, one sentence per WVS dimension with data, and a moral-reasoning closing line.

WVS grounding ensures empirically measured within-country variation rather than stereotypes [Deshpande et al., 2023]. Concrete persona examples and the WVS-to-trolley dimension linkage are detailed in Appendix D.

3.3 Batched Evaluation and Logit Extraction

All $N + 1 = 5$ input sequences (base + 4 personas) are evaluated in a *single* batched forward pass:

$$\mathbf{z}_{\text{base}},\ \mathbf{z}_1, \ldots, \mathbf{z}_N = f_\theta(\mathbf{x};\ \varnothing,\ \mathbf{s}_1, \ldots, \mathbf{s}_N)\,. \tag{2}$$

Per-category logit temperatures T_{cat} scale logits before computing decision gaps. We use $T_{\text{cat}} = 4.0$ (Species), 3.5 (Gender), and 1.5 for Age, Fitness, Social Value, and Utilitarianism (Table 1); the same values appear in Appendix E.

$$\delta_i = \frac{z_{i,b} - z_{i,a}}{T_{\text{cat}}}, \quad \delta_{\text{base}} = \frac{z_{\text{base},b} - z_{\text{base},a}}{T_{\text{cat}}}, \quad \bar{\delta} = \frac{1}{N}\sum_{i=1}^{N} \delta_i \quad \text{(consensus)}, \tag{3}$$

$$r_i = \delta_i - \delta_{\text{base}} \quad \text{(per-agent reward)}. \tag{4}$$

Here δ_{base} is the decision gap of the base prompt (no persona). The reward r_i measures how strongly persona i shifts the model *away from the culturally-neutral baseline*: a large $|r_i|$ indicates that the persona has identified a direction of moral preference the vanilla model does not exhibit. (Note: the scalar r introduced later in Eq. 6 denotes the bootstrap reliability weight, not a reward.) Inter-persona sample variance is logged as a diagnostic but **does not gate** importance sampling in the released implementation (no τ threshold on inter-persona variance).

3.4 Positional Debiasing

To cancel recency bias [Liu et al., 2024], each scenario is evaluated twice: once in the original ordering, once with *both* choice labels (A ↔ B) and group labels (Group A ↔ Group B) swapped in the native language. The debiased decision gap for each agent is $\delta_i = (\delta_i^{(\text{orig})} - \delta_i^{(\text{swap})})/2$; the base prompt yields δ_{base} under the same rule.

3.5 From Persona Disagreement to a Bounded Correction

The persona consensus $\bar{\delta}$ already captures the cultural direction, but applying it directly would ignore two problems: the consensus may overcorrect toward one loud persona, and finite-sample noise may push the correction in an unreliable direction. We address both through importance sampling with a cooperative, loss-averse utility.

The idea is to *search* for a correction that simultaneously improves alignment with every persona, not just the loudest one. We sample candidate perturbations $\epsilon_k \sim \mathcal{N}(0, \sigma^2)$ around the consensus and score each candidate by how much it moves the model closer to each persona relative to the uncorrected baseline. These per-persona gains are aggregated through a Prospect-Theory value function [Kahneman and Tversky, 1979, Tversky and Kahneman, 1992]—loss-averse with $\kappa \approx 2.25$, encoding a "do no harm" prior that linear or quadratic aggregators cannot express and that recent work shows LLMs themselves already exhibit under risk [Payne, 2025]. The softmax-weighted IS update can be read as approximate control-as-inference [Williams et al., 2018, Levine, 2018]; see Appendix A for the full motivation behind PT and IS.

$$U_{\text{total}}(\epsilon_k) = (1 - \lambda_{\text{coop}}) \frac{1}{N} \sum_{i=1}^{N} v\Big(\frac{g_{i,k}}{\sigma}\Big) \; + \; \lambda_{\text{coop}}\, v\Big(\frac{g_{\text{cons},k}}{\sigma}\Big), \tag{5}$$

where $v(\cdot)$ is the Kahneman–Tversky value function with gain curvature $\alpha{=}0.88$ and loss aversion $\kappa{=}2.25$ [Kahneman and Tversky, 1979], $g_{i,k}$ measures how much candidate k improves over the base for persona i (Appendix B), and $\lambda_{\text{coop}}{=}0.7$ weights the group consensus. Candidates are then softmax-weighted by utility and combined into a single correction $\delta^\star_m$; an effective-sample-size guard zeros the correction when the weights collapse, falling back on the raw consensus (Appendix B).

The second problem-sampling noise-is handled by the *dual-pass reliability gate*. Rather than trusting a single IS estimate, SWA-DPBR splits the budget into two independent passes and asks: do they agree? If yes, the correction is reliable; if not, it shrinks. Concretely:

$$\delta^\star_{\text{IS}} \;=\; \underbrace{\exp\!\big(-(\delta^\star_1 - \delta^\star_2)^2/s\big)}_{r \,\in\, (0,1]\text{: reliability weight}} \cdot \frac{\delta^\star_1 + \delta^\star_2}{2}. \tag{6}$$

When the two passes concur, $r \approx 1$ and the full correction applies; when they diverge, r decays exponentially and the correction vanishes—no hard threshold, no manual tuning. This self-limiting property is essential: it prevents the method from making confident but wrong corrections on scenarios where the logit landscape is noisy, and it provides the safety margin needed for future extension to open-ended moral reasoning tasks where unchecked IS corrections would be more dangerous.

Splitting the same IS budget into two independent replicates yields a scenario-level bootstrap-style variance probe [Efron, 1979, Elvira and Martino, 2021, Owen, 2013] that flags instability instead of hiding it behind more samples, while the exponential map preserves a smooth differentiable code path. An effective-sample-size guard [Kish, 1965] on each pass zeros the offset entirely when the weights collapse.

What the gate provably does (informal). Three elementary properties of the released estimator $\widetilde{\delta} = r \cdot (\widehat{\delta}^\star_1 + \widehat{\delta}^\star_2)/2$ follow under standard self-normalised IS assumptions [Geweke, 1989, Owen, 2013]:

(F1) *Asymptotic consistency.* Each pass $\widehat{\delta}^\star_m \to \delta^\star_\infty$ almost surely as $K_{\text{half}} \to \infty$ (SLLN on numerator and denominator of the self-normalised ratio).

(F2) *Variance-controlled bias.* The finite-K bias of each pass is $O(\mathrm{Var}_q[\widehat{\delta}^\star_m]/K_{\text{half}})$, so small per-pass variance simultaneously controls bias.

(F3) *Monotone shrinkage.* By Jensen's inequality on $u \mapsto e^{-u^2/s}$, $\mathbb{E}[r] \le \exp\!\big(-\mathrm{Var}[\widehat{\delta}^\star_1 - \widehat{\delta}^\star_2]/s\big)$, so $\widetilde{\delta}$ shrinks toward zero exactly when the two-pass variance is large.

We do *not* claim finite-K unbiasedness of $\widetilde{\delta}$ or optimality of the exponential shrinkage map; full assumptions, proofs, and the explicit list of properties we deliberately do *not* claim are in Appendix B.5.

The micro correction combines the reliability-gated IS offset with an interpolation between the persona consensus and the debiased base, controlled by sampling quality (Appendix B).

Algorithm 1 SWA-DPBR Inference for Scenario $\mathbf{x}$ in Country c

Require: Frozen LLM f_θ; personas $\{\mathbf{s}_i\}_{i=1}^{N}$; hyperparameters $K_{\text{half}}, \sigma, \lambda_{\text{coop}}, \eta, T_{\text{dec}}, s$
Ensure: $p_{\text{spare}} \in [0, 1]$

```
for order ∈ {orig, swap} do
    Batched forward: evaluate f_θ with base + N personas
    Extract logit pairs; compute δ_i^(ord)                          ▷ Eq. 2
end for
Average: δ_i ← (δ_i^(orig) − δ_i^(swap))/2                          ▷ §3.4
Compute consensus δ̄, rewards r_i                                   ▷ Eqs. 3–4
for m ∈ {1, 2} do
    for k = 1 to K_half do
        Sample ε_k ~ N(0, σ²); compute U_total(ε_k)                 ▷ Eq. 5
    end for
    δ*_m ← Σ_k w_k^(m) ε_k or 0 if ESS < ρ_eff                      ▷ §3.5
end for
r ← exp(−(δ*_1 − δ*_2)²/s); δ*_IS ← r(δ*_1 + δ*_2)/2                ▷ Eq. 6
δ_micro ← ESS-anchor blend of δ̄, δ_base, and δ*_IS                  ▷ see §3.5
δ_final ← country EMA prior blend                                   ▷ §3.6
return p_spare = σ(δ_final/T_dec)
```

3.6 Learning a Country-Level Prior

Each scenario so far is corrected independently. But moral preferences within a country are not independent—a model that underestimates youth preference in one Japanese trolley dilemma will likely do the same in the next. The hierarchical country prior exploits this by maintaining a running average of past corrections for each country and gradually blending it into future ones.

Concretely, after processing each scenario, an exponential moving average (EMA) accumulates the micro corrections: $\delta_{\text{country}} \leftarrow (1-\beta_{\text{ema}})\,\delta_{\text{country}} + \beta_{\text{ema}}\,\delta_{\text{micro}}$ with β_{ema}=0.1. The prior stays silent for the first N_{warm}=50 scenarios (where there is too little history to be useful), then gradually blends in: $\delta_{\text{final}} = \alpha_h\,\delta_{\text{country}} + (1-\alpha_h)\,\delta_{\text{micro}}$, where the blend weight α_h anneals from 0 to $\sim$0.63 over a typical country pool (Appendix E).

The prior contributes a modest but consistent improvement: disabling it degrades alignment by roughly 3% on the USA ablation slice. Its role is not to override per-scenario corrections but to stabilise them-smoothing out IS noise in the tail of the scenario pool so that late corrections benefit from the pattern established by early ones. The final spare probability is $p_{\text{spare}} = \sigma(\delta_{\text{final}}/T_{\text{dec}})$ with T_{dec}=0.5.

Algorithm 1 summarises the full procedure; numerical values for all symbols appear in Table 8 (Appendix E).

4 Experimental Setup

4.1 Dataset

We evaluate on MultiTP [Jin et al., 2025]: trolley-problem vignettes in native languages across six moral dimensions (Table 1). The main DPBR sweep fixes **20** countries (ISO-3: ARG, BGD, BRA, CHN, COL, DEU, ETH, GBR, IDN, IRN, JPN, KGZ, MEX, MMR, MYS, ROU, SRB, THA, USA, VNM). For each country we apply: (i) deduplication (retain only original-phrasing vignettes), (ii) quality filtering (remove Utilitarianism scenarios where both groups have equal size, as these carry no utilitarian signal), (iii) category capping (max 80 per dimension), and (iv) up-sampling via random oversampling with replacement (min 36 per dimension for stable AMCE estimation), yielding 310–500 scenarios per country (Appendix I).

Table 1: The six moral dimensions in MultiTP.

Dimension	Preferred	Contra	T_{cat}
Species	Human	Animal	4.0
Gender	Female	Male	3.5
Age	Young	Elderly	1.5
Fitness	Fit	Unfit	1.5
Social Value	High (exec.)	Low (homeless)	1.5
Utilitarianism	More lives	Fewer lives	1.5

4.2 Models

We evaluated **28 model–method combinations** spanning 12 architectures from 270M to 70B parameters (full landscape in Appendix AC). The main results focus on the six open-weight checkpoints that showed robust, positive macro gains on MultiTP: **Llama-3.3-70B**, **Magistral-Small-2509** (24B), **Phi-4** (14B), **Qwen3-VL-8B**, **Qwen2.5-7B**, and **Phi-3.5-mini** (3.8B). Table 2 orders them by nominal parameter count (descending). Models that degraded under SWA-DPBR share a common pattern: already-low vanilla MIS leaving insufficient headroom for IS corrections (§7). Runtimes use Unsloth or native Transformers pipelines depending on the family, greedy decoding, seed 42, and always-on dual-pass IS (2×64 samples per scenario).

Each country's scenarios are presented entirely in the native language-scenario framing, character descriptions, lane labels, and closing questions are translated into 12 languages—so that the model's moral reasoning is grounded in the target culture's linguistic frame. Only the decision tokens (A/B) remain English, verified as single-token entries in each model's vocabulary (Appendix F). The released SWA-DPBR configuration was selected from 25 method variants on a three-backbone, five-country prototyping panel. The top cluster of variants performed nearly identically on aggregate misalignment; dual-pass bootstrap reliability was chosen because it offered the best stability–accuracy tradeoff, materially lowering the rate at which IS reversed consensus decisions while providing an interpretable per-scenario uncertainty signal (details in Appendix AD).

4.3 Baselines and metrics

The primary inference-time baseline is (1) **Vanilla** greedy decoding with a generic system prompt. To isolate the contribution of within-country persona disagreement, we compare against *three families* of alternative inference-time corrections, all re-implemented in the released codebase:

Prompt-only baselines. (2) **B1 Country-Tailored Prompt**: prefix "Choose as a typical adult in ⟨country⟩ would." (3) **B2 WVS Profile Prompting**: prefix the country's WVS-7 profile sentence (six dimensions, qualitative descriptors as in Appendix D). (4) **B3 PRISM-Style Prompting** [Kirk et al., 2024]: short and long single-persona cultural-framing prefixes. (5) **B4 Country + WVS**: union of B1 and B2; the strongest prompt-only baseline.

Calibration / black-box baselines on the same 20-country Phi-4 grid (Table 4, Appendix M): (6) **MC-Dropout calibration** [Gal and Ghahramani, 2016, Kwon et al., 2026]; (7) **Per-country temperature / margin scaling** [Guo et al., 2017]; (8) **DiffPO-binary** [Chen et al., 2025]. Baselines (7)–(8) explicitly *consume* the evaluation target on a 25% calibration split, so they form an upper bound on what scalar human-AMCE-aware reweighting can achieve without any within-country disagreement signal.

White-box steering baseline. (9) **Activation Steering** [Arditi et al., 2025, Zou et al., 2023]: a per-country cultural steering vector built from contrastive WVS-pole prompt pairs and added to the residual stream at layer 32 with α=1.5. This baseline relaxes our black-box constraint (§1) and is therefore an upper bound on what a single-direction residual-space intervention can achieve given the same WVS information SWA-DPBR uses.

Reward-guided decoding stacks (ARGS [Khanov et al., 2024], controlled decoding [Mudgal et al., 2024]) are not directly applicable: a cross-cultural deployment would require a per-country reward model that does not exist, and adapting them to use persona logit gaps as the reward functional reduces to the PT-IS kernel we already ablate in Table 5. Primary headline metric is **MIS** (misalignment score): the ℓ_2 distance between the model's six-dimensional AMCE vector and the human vector for

Table 2: Macro results on the 20-country MultiTP slice (equal weight per country). MIS is mean ℓ_2 AMCE distance (↓). Rows follow nominal parameter count (descending); gains vs. vanilla are for reference only. "Win" counts countries with lower DPBR MIS than vanilla; mean r averages per-country Pearson correlation between model and human AMCE vectors.

Model	Mean MIS ↓	vs. van. (%)	Win / 20	Mean r
Llama-3.3-70B	0.668	+21.4	**20**	-0.01
Magistral-Small-2509 (24B)	0.350	+13.1	16	$+0.62$
Phi-4 (14B)	**0.346**	+23.7	18	$+0.56$
Qwen3-VL-8B	0.466	+18.9	**20**	$+0.22$
Qwen2.5-7B	0.362	+20.1	17	$+0.45$
Phi-3.5-mini (3.8B)	0.571	+10.4	17	-0.35

the same country [Jin et al., 2025], computed as in §3.1. Lower is better. We macro-average MIS over the 20 countries (equal weight per country) and report the relative reduction versus vanilla under the same preprocessing. **Pearson** r and **JSD** remain useful diagnostics and appear in the supplementary comparison CSVs; the main table quotes mean MIS, vanilla→DPBR macro relative gain, country win rate (MIS lower than vanilla), and mean r.

5 Results

5.1 Training-Free Steering Closes 19–24% of the Cultural Gap

A single decoding-time intervention—no weight changes, no per-country labels—materially closes the gap between LLM moral judgments and human cultural preferences documented by Awad et al. [2018], Jin et al. [2025]. Table 2 confirms this consistently. Across six open checkpoints spanning 3.8B to 70B parameters and four distinct model families, SWA-DPBR reduces macro misalignment (MIS) by 10–24% relative to vanilla, with the strongest backbones achieving **20/20 country wins**. These gains are notable in light of two independent concerns recently raised about cultural alignment evaluation: that cross-national LLM preferences track WEIRD populations rather than global averages [Henrich et al., 2010, Santurkar et al., 2023, Atari et al., 2023], and that LLMs systematically stereotype the moral values of non-Western societies [Zewail et al., 2026].

Two patterns deserve attention. *Calibration competes with scale.* The lowest absolute MIS does not come from the largest model: Phi-4 (14B) achieves MIS $= 0.346$ with a 23.7% gain, outperforming Llama-3.3-70B in absolute terms—even though the 70B model wins *every* country. This means SWA-DPBR can both *rescue* a badly misaligned large model and *refine* an already well-calibrated mid-sized one. *The method scales down.* Even the smallest checkpoint (Phi-3.5-mini, 3.8B) posts a clear 10.4% aggregate gain, demonstrating that inference-time cultural steering is not reserved for frontier-scale models.

The per-country data behind Table 2 confirms that these macro gains are not driven by a few lucky countries: Phi-4 halves misalignment on the USA and Japan, cutting absolute AMCE error to under 7 pp, while Llama-3.3-70B starts from a vanilla MIS nearly double the other models yet SWA-DPBR still pulls it down on every one of the 20 countries. Full per-country breakdowns and per-dimension analysis are in §5.2 and Appendix AB.

5.2 Gains Span Geographies, Not Just the West

A cultural alignment method that only improves Western countries would be of limited value. Table 3 decomposes the macro gains by region, confirming that SWA-DPBR's improvements are geographically broad: gains are visible across the Americas, Europe, and Southeast Asia, not concentrated in any single region. Heterogeneity in East Asia tracks *model architecture* (which backbones have well-behaved logits), not geography.

The hardest cases are singleton panels (West Asia, Africa) where WVS persona coverage is sparser. Here, variance is higher and a small backbone can occasionally move the wrong direction on a single country—precisely the regime where the dual-pass reliability gate earns its keep by shrinking noisy corrections toward zero.

Table 3: **Per-country SWA-DPBR results (20 countries).** Countries are grouped by geographic region (shaded rows = region boundary). Vanilla and SWA-DPBR MIS (↓), relative MIS improvement (%), and SWA-DPBR JSD (↓) / Pearson r (↑). Models in the same order as Table 2 (nominal parameter count, descending).

		Llama-3.3-70B					Magistral-Sml (24B)					Phi-4 (14B)					Qwen3-VL-8B					Qwen2.5-7B					Phi-3.5-mini (3.8B)				
ISO	**Region**	MIS_v	MIS_s	$\%\Delta_{\text{MIS}}$	JSD_s	r_s	MIS_v	MIS_s	$\%\Delta_{\text{MIS}}$	JSD_s	r_s	MIS_v	MIS_s	$\%\Delta_{\text{MIS}}$	JSD_s	r_s	MIS_v	MIS_s	$\%\Delta_{\text{MIS}}$	JSD_s	r_s	MIS_v	MIS_s	$\%\Delta_{\text{MIS}}$	JSD_s	r_s	MIS_v	MIS_s	$\%\Delta_{\text{MIS}}$	JSD_s	r_s
ARG	Americas	1.015	.891	+12.2	.178	−0.456	.364	.342	+6.0	.046	.601	.463	.389	+16.1	.046	.412	.520	.339	+35.0	.083	.508	.469	.377	+19.7	.064	.196	.770	.645	+16.1	.159	−0.498
BRA	Americas	.521	.458	+12.2	.119	.222	.411	.319	+22.5	.051	.333	.274	.299	-9.3	.061	.471	.529	.375	+29.1	.101	.212	.580	.411	+29.1	.081	−0.205	.531	.649	-22.1	.173	−0.743
COL	Americas	1.041	.931	+10.5	.189	−0.594	.413	.369	+10.7	.055	.386	.487	.438	+10.1	.057	.093	.559	.390	+30.1	.093	.161	.508	.417	+17.8	.070	−0.096	.788	.678	+14.0	.165	−0.602
MEX	Americas	1.037	.909	+12.3	.179	−0.385	.349	.333	+4.6	.040	.707	.485	.420	+13.4	.044	.512	.517	.350	+32.3	.081	.529	.467	.388	+17.0	.062	.252	.782	.662	+15.4	.159	−0.494
USA	Americas	.873	.578	+33.8	.121	.266	.448	.374	+16.5	.042	.731	.497	.243	+51.1	.058	.723	.677	.487	+28.0	.110	.422	.483	.383	+20.8	.075	.709	.645	.579	+10.2	.110	−0.380
DEU	Europe	.716	.643	+10.1	.189	.101	.265	.344	-30.3	.045	.724	.302	.255	+15.6	.059	.779	.318	.306	+3.7	.081	.568	.308	.415	-34.6	.069	.144	.569	.555	+2.5	.146	−0.403
GBR	Europe	.873	.598	+31.4	.141	.125	.442	.384	+13.2	.048	.674	.503	.319	+36.6	.082	.594	.677	.524	+22.6	.128	.230	.483	.399	+17.4	.082	.620	.654	.588	+10.1	.114	−0.392
ROU	Europe	.851	.593	+30.3	.147	.197	.446	.363	+18.7	.043	.708	.504	.385	+23.5	.058	.570	.659	.535	+18.8	.123	.162	.442	.330	+25.2	.051	.785	.671	.577	+14.0	.117	−0.352
SRB	Europe	.862	.605	+29.9	.150	.157	.463	.371	+19.8	.045	.670	.500	.385	+22.9	.057	.564	.676	.567	+16.2	.127	.075	.453	.303	+33.1	.051	.773	.677	.602	+11.1	.124	−0.463
CHN	E. Asia	.901	.748	+16.9	.183	−0.025	.367	.293	+20.2	.050	.711	.407	.214	+47.4	.054	.781	.569	.415	+27.0	.121	.269	.437	.356	+18.6	.084	.405	.738	.559	+24.3	.162	−0.041
JPN	E. Asia	.521	.471	+9.5	.107	.249	.341	.265	+22.3	.054	.709	.395	.195	+50.6	.051	.671	.407	.316	+22.4	.090	.453	.296	.356	-20.1	.047	.535	.454	.455	-0.3	.063	−0.233
IDN	SE. Asia	.688	.557	+19.0	.092	.071	.283	.344	-21.5	.059	.511	.459	.274	+40.3	.059	.595	.408	.402	+1.5	.091	.318	.480	.402	+16.1	.058	.251	.492	.484	+1.6	.120	−0.769
MMR	SE. Asia	.872	.652	+25.2	.185	−0.031	.496	.365	+26.5	.062	.457	.482	.357	+25.9	.057	.570	.695	.563	+18.9	.139	−0.066	.461	.302	+34.3	.062	.590	.685	.590	+13.8	.135	−0.623
MYS	SE. Asia	.841	.589	+29.9	.155	.003	.443	.313	+29.4	.043	.683	.459	.328	+28.5	.053	.614	.656	.504	+23.1	.128	.097	.445	.274	+38.5	.050	.763	.631	.535	+15.1	.118	−0.426
THA	SE. Asia	.815	.577	+29.2	.152	.125	.406	.295	+27.3	.041	.738	.456	.313	+31.4	.047	.685	.619	.480	+22.5	.112	.231	.414	.221	+46.6	.046	.841	.619	.520	+16.1	.112	−0.283
VNM	SE. Asia	.845	.740	+12.4	.142	−0.045	.324	.347	-7.2	.050	.589	.406	.336	+17.3	.064	.596	.445	.420	+5.5	.087	.339	.501	.468	+6.6	.060	−0.018	.481	.507	-5.4	.067	−0.143
BGD	S. Asia	.864	.628	+27.3	.162	.029	.450	.361	+19.6	.047	.625	.489	.368	+24.7	.048	.607	.674	.587	+12.9	.133	−0.116	.461	.284	+38.3	.055	.695	.666	.565	+15.3	.115	−0.426
KGZ	C. Asia	.849	.585	+31.1	.149	.176	.460	.353	+23.2	.048	.633	.499	.393	+21.1	.058	.476	.661	.574	+13.1	.133	−0.017	.443	.249	+43.8	.039	.829	.660	.572	+13.3	.119	−0.405
IRN	W. Asia	1.039	.852	+17.9	.159	−0.294	.427	.401	+6.1	.055	.637	.397	.530	-33.5	.086	.055	.507	.506	+0.2	.090	.057	.373	.467	-25.1	.067	.252	.511	.445	+12.9	.056	.766
ETH	Africa	.967	.747	+22.7	.172	−0.077	.462	.472	-2.0	.054	.655	.615	.485	+21.2	.041	.808	.724	.684	+5.4	.141	−0.095	.558	.438	+21.6	.059	.676	.734	.661	+10.0	.114	−0.027
Mean		**.849**	**.668**	**+21.2**	**.154**	**−0.009**	**.403**	**.350**	**+11.3**	**.049**	**.624**	**.454**	**.346**	**+22.7**	**.057**	**.559**	**.575**	**.466**	**+18.4**	**.110**	**.217**	**.453**	**.362**	**+18.2**	**.062**	**.450**	**.638**	**.571**	**+9.4**	**.122**	**−0.347**

Legend. $\text{MIS}_v/\text{MIS}_s$: vanilla vs. SWA-DPBR ℓ_2 misalignment; $\%\Delta_{\text{MIS}}$: relative improvement in MIS; JSD_s and r_s: Jensen–Shannon distance and Pearson correlation for SWA-DPBR only. **Region groups:** Americas (ARG, BRA, COL, MEX, USA); Europe (DEU, GBR, ROU, SRB); E. Asia (CHN, JPN); SE. Asia (IDN, MMR, MYS, THA, VNM); S. Asia (BGD); C. Asia (KGZ); W. Asia (IRN); Africa (ETH).

MIS vs. Pearson r. Several rows show SWA-DPBR reducing MIS while Pearson r remains negative. The two metrics are genuinely orthogonal—MIS is an ℓ_2 amplitude on the six-dim AMCE vector, r is shape (rank ordering of dimensions). We follow the MultiTP convention of treating MIS as primary and supplement with rank-based agreement (Kendall τ, mean rank-error) in Appendix R.

5.3 Persona Disagreement Beats Country-Agnostic and Calibration Baselines

A natural worry is that the gains in Table 2 could be matched by simpler decoding-time corrections that do not require a persona panel at all. Table 4 compares SWA-DPBR against three such alternatives on the same 20-country Phi-4 grid: country-agnostic uncertainty inflation [Gal and Ghahramani, 2016, Kwon et al., 2026], per-country scalar calibration [Guo et al., 2017]-style fits on a held-out split, and a black-box probability-mixing baseline adapted from DiffPO [Chen et al., 2025]. The latter two explicitly *consume the human AMCE* on a 25% calibration split—an upper bound on what scalar country-conditioned reweighting can achieve.

Table 4: Inference-time baselines on the 20-country Phi-4 grid (macro mean MIS ↓). MC-Dropout follows Gal and Ghahramani [2016], Kwon et al. [2026]; DiffPO-binary adapts Chen et al. [2025] to binary decisions. Calibration baselines fit per-country scalars on a 25% split against the public human AMCE; SWA-DPBR uses no human AMCE at inference. SWA-DPBR more than halves the residual gap left by the strongest country-agnostic baseline (MC-Dropout), while country-conditioned scalar calibrations actually *degrade* macro MIS. Per-country tables in Appendix M.

Method	Macro MIS ↓	Uses human AMCE?	vs. vanilla
Vanilla (no correction)	0.454	no	—
MC-Dropout (T=8, p=0.1)	0.403	no	+11.2%
Per-country temperature T_c	0.513	yes	−13.0%
Per-country margin m_c	0.506	yes	−11.5%
DiffPO-binary	0.494	yes	−8.8%
SWA-DPBR (ours)	**0.346**	no	+23.7%

Two findings stand out. *Country-agnostic uncertainty inflation (MC-Dropout) captures part of the gap but stalls at half of SWA-DPBR's gain*, isolating the contribution of within-country persona disagreement. *Per-country scalar calibrations actually hurt*: a single T_c or additive margin m_c cannot express the dimension-specific corrections the human AMCE vector demands, so fitting one on the calibration split over-fits and degrades the held-out test split. Even DiffPO-binary, which directly mixes the public human AMCE into the decoding distribution, plateaus at 0.494—above

SWA-DPBR's 0.346 despite consuming the evaluation target. The picture is consistent: structured persona disagreement is not a redundant signal that scalar corrections could substitute for.

5.4 Sensitivity Suite Confirms Robustness

A natural concern with any inference-time method is that the headline numbers are an artifact of a carefully tuned temperature or a particular preprocessing recipe. Table 5 stress-tests this across nine dimensions. The answer is reassuring: sweeps over T_{dec}, uniform and per-category T_{cat}, and six dataset preprocessing variants all produce JSD bands **well within** the ± 0.004 bootstrap CI. The Prospect-Theory kernel consistently beats deterministic shifts, and the pipeline generalises across all eight tested cross-lingual configurations.

Table 5: Robustness suite. Each test independently probes one assumption behind SWA-DPBR; the bootstrap confidence interval on JSD is ± 0.004.

Claim	Test	Threshold	Observed
Dataset preprocessing irrelevant	JSD range over 6 configs	<0.010	**0.0034**
T_{dec} insensitivity	JSD span (Sweep A)	<0.010	**0.0074**
Uniform T_{cat} insensitivity	JSD span (Sweep B)	<0.010	**0.0073**
Per-category T_{cat} insensitivity	JSD span (Sweep C)	<0.010	**0.0077**
PT-IS $>$ best deterministic shift	ΔJSD(CS-Clamp$\rightarrow$PT–IS)	>0	**0.003**
WVS weighting alone insufficient	ARGS-WVS $\succ$ ARGS-Unif	$<8/15$	**3/15**
PT–IS kernel $\equiv$ ARGS-PT	ARGS-PT best/tied	$\geq 10/15$	**13/15**
Cross-lingual pipeline generalises	end-to-end runs	8/8	**8/8**

6 Ablation: Identifying Load-Bearing Components

To understand *why* the method works—not just *that* it works—we ablate four design choices on Phi-4 (14B), USA (310 scenarios; Table 6).[1]

Positional debiasing is the dominant component (Row 3, $\Delta r = -0.282$, $\Delta\text{MIS} = +0.175$). The DPBR diagnostic confirms a large raw option-A positional bias of $\bar{b} \approx -29.1$ logit units; debiasing must neutralise this artifact before any other component can contribute.

Personas are essential (Row 4, $\Delta r = -0.183$, $\Delta\text{MIS} = +0.083$): removing the WVS panel collapses rank-order alignment significantly, confirming that survey-grounded diversity is the method's foundation, not decoration. Once positional artifacts are removed, the vLLM decision surface retains enough curvature for within-country WVS disagreement to provide real signal; per-scenario IS corrections are selective (flip rate 29.0%, $\overline{\text{rel_}r} = 0.629$) rather than degenerate.

PT–IS contributes modestly but consistently (Rows 1–2): dropping IS entirely degrades both r (-0.026) and MIS ($+0.006$); always-on IS without variance gating hurts slightly more on MIS ($+0.015$), confirming the value of the adaptive reliability gate. Notably, the IS *flip rate* is only 29.0%: SWA-DPBR *modulates confidence* rather than reversing decisions (Appendix AA).

These findings confirm that the ablation importance hierarchy mirrors the base model's decision-surface conditioning: debiasing unlocks the signal that personas and IS then refine, and the full stack is greater than the sum of its parts.

7 Discussion

The central story of these experiments is that *structured disagreement among culturally grounded personas is a surprisingly effective steering signal for frozen LLMs*—and that the same signal can be aggregated through a loss-averse utility and a self-regulating reliability gate without any of the white-box access, per-culture supervision, or training cost that competing approaches require. This is the operational claim behind the pluralistic alignment vision [Sorensen et al., 2024].

[1] The ablation uses a vLLM backend to match the main EXP-24 sweep; one K-sample IS batch per scenario with dual-pass reliability disabled for isolation. Primary metric is MIS (↓).

Table 6: Ablation on Phi-4 (14B), USA (310 scenarios), **SWA-DPBR** with one K-sample importance-sampling batch per scenario (dual-pass reliability disabled for isolation). Full configuration: $r = 0.603$, MIS $= 0.3489$. Primary metric: MIS (↓).

#	Configuration	r ↑	Δr	MIS ↓	ΔMIS
	Full SWA-DPBR	**0.603**	–	**0.3489**	–
1	No-IS (consensus only)	0.577	−.026	0.3550	+.006
2	Always-on PT-IS	0.577	−.026	0.3639	+.015
3	No debiasing	0.320	−.282	0.5243	+.175
4	Without persona	0.420	−.183	0.4324	+.083

Equally important is *where* the method fails, because the failure pattern points to a structural prerequisite rather than a cultural one. Per-country grids (§5.2) and the logit-conditioning correlation in Appendix O (r=+0.627, p=0.003 between mean decision margin and SWA-DPBR MIS gain) show that residual trouble spots cluster by model architecture, not geography: when the base model's decision-token logits are poorly conditioned, IS perturbations find no informative gradient to follow regardless of the target country. The same logic explains the diminishing-returns pattern—models with already-low vanilla MIS leave too little headroom for the IS stage—and connects to the broader observation that targeted debiasing can cause "spillover" into untargeted dimensions [Chand et al., 2026], reinforcing that inference-time interventions must be self-limiting (the role of the dual-pass gate). Deployment overhead is modest: $\approx 3.6\times$ vanilla (Appendix AA), dominated by the batched multi-agent forward pass and essentially flat in the IS budget K.

Limitations. Three caveats bound generalisation. First, hyperparameters were selected on a synthetic dilemma pool and transfer-checked on a small MultiTP subset; sweeps in Appendices H–I and the 25-method engineering comparison (Appendix AD) are our main evidence of stability, but multi-seed confidence intervals remain a priority. Second, we optimise against crowdsourced AMCE vectors [Awad et al., 2018, Jin et al., 2025] without a human study of perceived cultural appropriateness [Atari et al., 2023], so the reader should treat our numbers as alignment to a survey statistic, not proof of legitimacy—a concern echoed by recent work questioning the reliability of cultural alignment evaluations [Khan et al., 2025]. Third, per-dimension analysis (Appendix AB) shows that the bottleneck dimension shifts with backbone quality: weaker models are bottlenecked on Utilitarianism, while better-calibrated ones are bottlenecked on Social Value and Age—a structural ceiling that no within-country persona signal can fully close. Fourth, a pilot on the open-ended BLEnD benchmark (Appendix AE) shows that SWA-DPBR's binary logit-gap formulation does not transfer to free-form generation: aggregate SEM-B drops by 4.4 pp, with gains limited to high-resource countries where the baseline already performs well. This confirms that softmax-aware IS and a vocabulary-level reliability gate are prerequisites for open-ended deployment. SWA-DPBR also requires decision-token logits, limiting deployment on text-only APIs; further limitations are in Appendix AF.

On the robustness of the Prospect-Theory kernel. A recent study [Wang et al., 2025] reports that prospect-theoretic biases in LLMs can be fragile under linguistic uncertainty: framing decisions with vague quantifiers or uncertain payoffs weakens the loss-aversion asymmetry. Our setting partially avoids this concern because the PT kernel operates on *persona logit gaps*, not on the raw decision utility communicated in the prompt—the inputs to $v(\cdot)$ are sigma-normalised numerical differences, not prompted quantities. Still, the caveat is real: if a native-language frame introduces linguistic hedging that shifts persona calibration, the PT asymmetry we rely on could be diluted. The per-country Pearson r spread in Table 3 is consistent with this effect being small but non-zero; a more direct probe is flagged for future work.

Broader impact. SWA-DPBR can make models more responsive to documented cross-national preference heterogeneity without retraining, yet it can also encode harmful majorities or survey stereotypes if deployed naively—a tension that recent work on moral stereotyping in LLMs makes especially salient [Zewail et al., 2026]. We recommend four safeguards: governance ensuring the method does not override safety-critical guardrails; periodic persona refresh as newer WVS waves appear; per-dimension AMCE audits to catch entrenched bias; and optional penalty terms on sensitive dimensions to embed review directly into the decoding loop. A fifth safeguard, specifically addressing the risk that the PT aggregator silences minority views within a country, is a *per-persona floor*:

the reliability weight r already attenuates corrections when the two IS passes disagree, and an additional floor on each persona's post-correction utility $v(\tilde{g}_{i,k})$ can prevent any single persona's preference from being driven arbitrarily far from its vanilla estimate. A sweep on a 3-country panel $(\text{USA}, \text{VNM}, \text{DEU})$ over floors $\in \{0, 0.5, 1.0, 2.0\}$ confirms the safeguard is essentially free: floor $= 1.0$ leaves macro MIS within ± 0.005 of the unfloored default on every country and *slightly improves* the worst-persona reward on USA. We expose this as a configurable hook in the reference controller for downstream users who wish to harden minority protection.

8 Conclusion

LLMs need not be retrained to respect moral diversity. SWA-DPBR shows that survey-grounded persona disagreement can be converted into bounded, self-regulating logit corrections at inference time, reducing misalignment by 19–24% across twenty countries and four continents without touching a single weight. The method's limits are mechanical, not cultural: collapsed decision-logit entropy breaks importance sampling regardless of geography, models with already-low misalignment occasionally overshoot, and a pilot on BLEnD open-ended QA (Appendix AE) confirms that the binary logit-gap formulation does not yet transfer to free-form generation. The path forward is softmax-aware IS for open-ended settings, richer proposal adaptation for ill-conditioned logits, per-dimension outcome analysis, updated survey waves, and decoding-time safeguards—concrete steps toward the pluralistic alignment vision [Sorensen et al., 2024] with practical, deployable tools.

References

H. An, S. Park, H. Jin, and Y.-S. Han. Steering language models before they speak: Logit-level interventions. *arXiv preprint arXiv:2601.10960*, 2026.

A. Arditi, O. Obeso, A. Syed, D. Paleka, N. Panickssery, W. Gurnee, and N. Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, volume 37, pages 136037–136083, 2025.

M. Atari, M. J. Xue, P. S. Park, D. E. Blasi, and J. Henrich. Which humans? *PsyArXiv preprint*, 2023. doi: 10.31234/osf.io/5b26t. `https://osf.io/preprints/psyarxiv/5b26t`.

E. Awad, S. Dsouza, R. Kim, J. Schulz, J. Henrich, A. Shariff, J.-F. Bonnefon, and I. Rahwan. The moral machine experiment. *Nature*, 563(7729):59–64, 2018.

Y. Bai, S. Kadavath, S. Kundu, A. Askell, et al. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*, 2022.

S. C. Bobbili, U. Dinesha, D. Narasimha, and S. Shakkottai. PITA: Preference-guided inference-time alignment for LLM post-training. *arXiv preprint arXiv:2507.20067*, 2025.

S. Chand, F. Baca, and E. Ferrara. No free lunch in language model bias mitigation? Targeted bias reduction can exacerbate unmitigated LLM biases. *AI*, 7(1):24, 2026. doi: 10.3390/ai7010024.

R. Chen, W. Chai, Z. Yang, X. Zhang, J. T. Zhou, T. Quek, S. Poria, and Z. Liu. DiffPO: Diffusion-styled preference optimization for efficient inference-time alignment of large language models. *arXiv preprint arXiv:2503.04240*, 2025.

A. Deshpande, V. Murahari, T. Rajpurohit, A. Kalyan, and K. Narasimhan. Toxicity in ChatGPT: Analyzing persona-assigned language models. *arXiv preprint arXiv:2304.05335*, 2023.

E. Durmus, K. Nguyen, T. I. Liao, N. Schiefer, A. Askell, A. Bakhtin, et al. Towards measuring the representation of subjective global opinions in language models. *arXiv preprint arXiv:2306.16388*, 2023.

B. Efron. Bootstrap methods: another look at the jackknife. *The Annals of Statistics*, 7(1):1–26, 1979.

V. Elvira and L. Martino. Advances in importance sampling. *Wiley StatsRef: Statistics Reference Online*, pages 1–14, 2021.

K. Ethayarajh, W. Xu, N. Muennighoff, D. Jurafsky, and D. Kiela. KTO: Model alignment as prospect theoretic optimization. In *International Conference on Machine Learning*, 2024. arXiv:2402.01306.

Y. Gal and Z. Ghahramani. Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. In *Proceedings of the 33rd International Conference on Machine Learning*, 2016.

J. Geweke. Bayesian inference in econometric models using Monte Carlo integration. *Econometrica*, 57(6):1317–1339, 1989.

C. M. Greco, L. La Cava, and A. Tagarelli. Culturally grounded personas in large language models: Characterization and alignment with socio-psychological value frameworks. *arXiv preprint arXiv:2601.22396*, 2026.

C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1321–1330, 2017. arXiv:1706.04599.

C. Haerpfer, R. Inglehart, A. Moreno, C. Welzel, K. Kizilova, J. Diez-Medrano, M. Lagos, P. Norris, E. Ponarin, and B. Puranen. World Values Survey: Round seven – country-pooled datafile. Madrid, Spain & Vienna, Austria: JD Systems Institute & WVSA Secretariat, 2022. `https://doi.org/10.14281/18241.20`.

D. Hendrycks, C. Burns, S. Basart, A. Critch, J. Li, D. Song, and J. Steinhardt. Aligning AI with shared human values. In *International Conference on Learning Representations*, 2021.

J. Henrich, S. J. Heine, and A. Norenzayan. The weirdest people in the world? *Behavioral and Brain Sciences*, 33(2-3):61–83, 2010.

R. Inglehart and C. Welzel. *Modernization, Cultural Change, and Democracy: The Human Development Sequence*. Cambridge University Press, 2005.

Z. Jin, S. Levine, F. Gonzalez Adauto, O. Kamal, M. Sap, M. Sachan, R. Mihalcea, J. Tenenbaum, and B. Schölkopf. When to make exceptions: Exploring language models as accounts of human moral judgment. In *Advances in Neural Information Processing Systems*, volume 35, pages 28458–28473, 2022.

Z. Jin, M. Kleiman-Weiner, G. Piatti, S. Levine, J. Liu, F. Gonzalez, F. Ortu, A. Strausz, M. Sachan, R. Mihalcea, Y. Choi, and B. Schölkopf. Language model alignment in multilingual trolley problems. *arXiv preprint arXiv:2407.02273*, 2025. Published as a conference paper at ICLR 2025.

D. Kahneman and A. Tversky. Prospect theory: An analysis of decision under risk. *Econometrica*, 47(2):263–291, 1979.

S. Kanai, T. Yoshida, H. Takahashi, H. Kuroki, and K. Hashimoto. Test-time alignment of LLMs via sampling-based optimal control in pre-logit space. *arXiv preprint arXiv:2510.26219*, 2025.

A. Khan, S. Casper, and D. Hadfield-Menell. Randomness, not representation: The unreliability of evaluating cultural alignment in LLMs. *arXiv preprint arXiv:2503.08688*, 2025.

M. Khanov, J. Burapacheep, and Y. Li. ARGS: Alignment as reward-guided search. *arXiv preprint arXiv:2402.01694*, 2024.

S. Khanuja, H. Liu, S. Zhang, J. Lambert, M. Chen, R. Mathews, and L. Wang. Steering LLMs for culturally localized generation. *arXiv preprint arXiv:2603.23301*, 2026.

H. R. Kirk, A. Whitefield, P. Röttger, A. Bean, K. Margatina, J. Ciro, R. Mosquera, M. Bartolo, A. Williams, H. He, B. Vidgen, and S. A. Hale. The PRISM alignment dataset: What participatory, representative and individualised human feedback reveals about the subjective and multicultural alignment of large language models. *arXiv preprint arXiv:2404.16019*, 2024. NeurIPS 2024 Datasets & Benchmarks.

L. Kish. *Survey Sampling*. John Wiley & Sons, 1965.

M. Kleiman-Weiner, R. Saxe, and J. B. Tenenbaum. Learning a commonsense moral theory. *Cognition*, 167:107–123, 2017.

J. Kwon, L. F. Vecchietti, S. Park, and M. Cha. Dropouts in confidence: Moral uncertainty in human-LLM alignment. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2026. arXiv:2511.13290.

S. Levine. Reinforcement learning and control as probabilistic inference: Tutorial and review. *arXiv preprint arXiv:1805.00909*, 2018.

N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024.

S. Mudgal, J. Lee, H. Ganapathy, Y. Li, T. Wang, Y. Huang, Z. Chen, H.-T. Cheng, M. Collins, T. Strohman, et al. Controlled decoding from language models. *arXiv preprint arXiv:2310.17022*, 2024.

J. Myung, N. Lee, Y. Zhou, J. Jin, R. A. Putri, D. Antypas, H. Borkakoty, E. Kim, C. Perez-Almendros, A. A. Ayele, et al. BLEnD: A benchmark for LLMs on everyday knowledge in diverse cultures and languages. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2024. arXiv:2406.09948.

A. Nie, Y. Zhang, A. Amdekar, C. Piech, T. B. Hashimoto, and T. Gerstenberg. MoCa: Measuring human-language model alignment on causal and moral judgment tasks. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

L. Ouyang, J. Wu, X. Jiang, D. Almeida, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.

A. B. Owen. *Monte Carlo Theory, Methods and Examples*. Stanford University, 2013. Online manuscript, `https://artowen.su.domains/mc/`.

K. Payne. An analysis of AI decision under risk: Prospect theory emerges in large language models. *arXiv preprint arXiv:2508.00902*, 2025.

H. Qiu, K.-H. Huang, R. Zheng, J. Sun, and N. Peng. Multimodal cultural safety: Evaluation framework and alignment strategies, 2025.

M. J. Ryan, W. Held, and D. Yang. Unintended impacts of LLM alignment on global representation. *arXiv preprint arXiv:2402.15018*, 2024.

S. Santurkar, E. Durmus, F. Ladhak, C. Lee, P. Liang, and T. Hashimoto. Whose opinions do language models reflect? *arXiv preprint arXiv:2303.17548*, 2023.

S. H. Schwartz, J. Cieciuch, M. Vecchione, E. Davidov, R. Fischer, C. Beierlein, A. Ramos, M. Verkasalo, J.-E. Lönnqvist, K. Demirutku, O. Dirilen-Gumus, and M. Konty. Refining the theory of basic individual values. *Journal of Personality and Social Psychology*, 103(4):663–688, 2012.

T. Sorensen, J. Moore, J. Fisher, M. Gordon, N. Mireshghallah, C. M. Rytting, A. Ye, L. Jiang, X. Lu, N. Dziri, et al. A roadmap to pluralistic alignment. *arXiv preprint arXiv:2402.05070*, 2024.

E. Todorov. Efficient computation of optimal actions. *Proceedings of the National Academy of Sciences*, 106(28):11478–11483, 2009.

A. Tversky and D. Kahneman. Advances in prospect theory: Cumulative representation of uncertainty. *Journal of Risk and Uncertainty*, 5(4):297–323, 1992.

R. Wang, Q. Lin, J. Liu, Q. Zong, T. Zheng, D. Guo, H. Shi, W. Wang, and Y. Song. Rethinking prospect theory for LLMs: Revealing the instability of decision-making under epistemic uncertainty. *arXiv preprint arXiv:2508.08992*, 2025.

G. Williams, P. Drews, B. Goldfain, J. M. Rehg, and E. A. Theodorou. Information-theoretic model predictive control: Theory and applications to autonomous driving. *IEEE Transactions on Robotics*, 34(6):1603–1622, 2018.

J. Yao, X. Yi, J. Wang, Z. Dou, and X. Xie. CAReDiO: Cultural alignment of LLM via representativeness and distinctiveness guided data optimization. *arXiv preprint arXiv:2504.08820*, 2025.

A. Zewail, A. Figueroa, J. Graham, and M. Atari. Moral stereotyping in large language models. *Proceedings of the National Academy of Sciences*, 123(10), 2026. doi: 10.1073/pnas.2519941123.

B. Zhang, X. Zhao, J. Li, H. Chen, and Z. Chen. Mind the gap in cultural alignment: Task-aware culture management for large language models. *arXiv preprint arXiv:2602.22475*, 2026.

A. Zou, L. Phan, S. Chen, J. Campbell, P. Guo, et al. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

A Design Rationale: Why Each Component

This appendix collects the full justification for the three design choices previewed in §1 and §3: the four-persona WVS panel, the Prospect-Theory IS kernel, and the dual-pass reliability gate.

Why a panel, and why the World Values Survey? A single "prompt-as-country" encoder—e.g., *"Answer as a Japanese person"*—collapses a distribution to a point estimate and discards the very spread the correction stage needs. Cultural-value research has long argued against this framing: both Inglehart-Welzel's modernisation theory [Inglehart and Welzel, 2005] and Schwartz's refined theory of basic individual values [Schwartz et al., 2012] find within-country variance on any single value dimension comparable to between-country variance, with demographic cleavages—most prominently age cohort—driving the largest slice. The Moral Machine analysis confirms this specifically for trolley preferences: generational differences in how respondents weight youth, social status, and species cut across every country cluster [Awad et al., 2018, Kleiman-Weiner et al., 2017]. WVS Wave 7 [Haerpfer et al., 2022] is the largest public instrument that reports *individual-level* microdata with the relevant age, gender, and value coverage, making it the natural substrate for data-driven personas rather than author-crafted stereotypes—a well-documented failure mode of prompt-based persona work [Deshpande et al., 2023].

Why $N{=}4$ specifically? Four is the smallest panel that simultaneously covers the three demographic axes documented to carry most within-country variance—young (<36), middle-aged (36–55), older (>55)—*plus* a country-wide aggregate that anchors the ensemble against cohort-specific sampling noise. Three personas would sacrifice the population anchor and leave the consensus $\bar{\delta}$ vulnerable to a single outlier cohort; five or more would duplicate coverage of the same axis without adding a new source of variance (WVS does not resolve finer cohorts reliably for every country). The cohort structure matches the statistical-survey tradition of stratifying by age before any country-level estimate [Kish, 1965]. Empirically, the ablation in §6 confirms the ensemble is load-bearing: removing it costs -0.183 Pearson r and $+0.083$ MIS.

Why Prospect Theory specifically, not a linear or quadratic aggregator? Three independent considerations pick out the Kahneman-Tversky value function [Kahneman and Tversky, 1979, Tversky and Kahneman, 1992] over simpler alternatives. *(1) Asymmetric harm.* In a cultural panel, a correction that improves three personas while making one worse is not equivalent to one that helps all four equally: the former silences a demographic group the survey says exists. Linear averaging is oblivious to this asymmetry; quadratic aggregation treats $+x$ and $-x$ as equally valuable and costly. Loss aversion with $\kappa{\approx}2.25$ is the textbook object for encoding "harm is worse than equal-sized help". *(2) Diminishing sensitivity matches the scale of the cultural signal.* Gain curvature $\alpha{=}0.88{<}1$ compresses large persona gains relative to small ones, preventing one outlier persona (perhaps from sparse WVS coverage) from dominating the utility [Tversky and Kahneman, 1992]. *(3) The base model is already prospect-theoretic.* Recent empirical work shows that LLMs exhibit prospect-theoretic biases in decision-making under risk, with loss-aversion coefficients in the same neighbourhood as human respondents [Payne, 2025]. Using a PT utility therefore *aligns the controller's objective with a structure the base model is already tuned for* rather than imposing an alien linear-utility objective.

Why importance sampling at all? The cultural correction we seek is the argmax of a utility that values each persona's alignment and penalises corrections that make any of them worse [Levine, 2018]. Closed-form optimisation is unavailable because the utility is non-convex (PT's kink at zero, the ℓ_1-style gain definition in Eq. 7), so we approximate the optimal control distribution with a path-integral-style softmax over perturbations [Williams et al., 2018, Todorov, 2009]. The Gaussian proposal $\epsilon_k \sim \mathcal{N}(0, \sigma^2)$ plays the role of an entropy-maximal default, and the softmax weights $w_k \propto \exp(U_{\text{total}}/\eta)$ implement the KL-regularised Boltzmann posterior balancing utility against deviation from the consensus. This control-as-inference reading licenses Eq. 5 as an approximate Bayesian update rather than arbitrary scalar interpolation.

Why dual-pass, rather than more samples or a hard threshold? Self-normalised IS is a biased estimator at finite sample size [Elvira and Martino, 2021]: the bias shrinks as $O(1/K)$ but does not reveal itself to the caller, so "more samples" hides instability rather than flagging it. Two independent replicates of size K_{half} using the same total budget $K{=}2K_{\text{half}}$ trade a small increase in

per-pass variance for a *direct, scenario-level probe* of that instability: the squared gap $(\delta_1^\star - \delta_2^\star)^2$ is a one-dimensional bootstrap-style variance estimator [Efron, 1979] whose small values license the correction and whose large values veto it automatically, without a hand-tuned threshold. The exponential map $r = \exp\big(-(\delta_1^\star - \delta_2^\star)^2/s\big)$ is preferred over a hard indicator because it preserves a single differentiable code path and its decay matches the shape of a Gaussian likelihood ratio under independent noise. An effective-sample-size guard [Kish, 1965] on each pass ($N_{\text{eff}}/K_{\text{half}} < \rho_{\text{eff}}$) zeros the offset entirely when the weights collapse, providing a second, coarser safety net below the smooth reliability gate. We emphasise that this construction is a variance heuristic, not a formal unbiasedness certificate; finite-K bounds would require stronger assumptions on the logit landscape than we are willing to claim (Appendix B).

B PT–IS, Sampling, and Dual-Pass Reliability

B.1 Gain definitions

For candidate $\tilde{\delta}_k = \bar{\delta} + \epsilon_k$, the per-persona and consensus gains are:

$$g_{i,k} = |\delta_{\text{base}} - \delta_i| - |\tilde{\delta}_k - \delta_i|, \tag{7}$$

$$g_{\text{cons},k} = |\delta_{\text{base}} - \bar{\delta}| - |\tilde{\delta}_k - \bar{\delta}|. \tag{8}$$

$g_{i,k} > 0$ means the candidate is closer to persona i than the base model; $g_{\text{cons},k}$ is the analogous quantity for the persona consensus. These enter Eq. 5 after σ-normalisation. The IS update for each pass is $\delta_m^\star = \sum_k w_k^{(m)} \epsilon_k$ with $w_k^{(m)} = \mathrm{softmax}(U_{\text{total}}(\epsilon_k)/\eta)$ and temperature $\eta = 0.5$. An ESS guard zeros $\delta_m^\star$ when the Kish effective sample size $N_{\text{eff}}/K_{\text{half}}$ falls below $\rho_{\text{eff}} = 0.1$.

B.2 Utility shape and the omitted explicit KL-on-ϵ term

The scalar decision gap δ acts as a one-dimensional state and the perturbations ϵ_k as control inputs at a single horizon step. Under the path-integral / free-energy view [Williams et al., 2018], softmax weights $w_k \propto \exp(U_{\text{total}}/\eta)$ can be read as targeting a Boltzmann distribution over candidate perturbations. **Our released code does not add** a separate quadratic penalty $\propto \epsilon_k^2/(2\sigma^2)$ to U_{total} in Eq. 5: the released controller matches it exactly, and the Gaussian proposal $\epsilon_k \sim \mathcal{N}(0, \sigma^2)$ already limits how far $\tilde{\delta}_k$ can move from $\bar{\delta}$. An optional explicit KL-on-control term remains a straightforward extension if future work needs sharper tail control on $|\epsilon_k|$.

B.3 Comparison with Simpler Alternatives

Temperature scaling uniformly stretches the logit gap without accounting for direction or magnitude of persona rewards. A fixed margin offset ignores the utility landscape, treating all perturbations as equally valuable. PT-weighted IS performs an *adaptive, scenario-specific* correction via nonlinear importance weighting. Auxiliary runs with deterministic consensus-only shifts and WVS-linear reward kernels underperform the Prospect-Theory kernel on the same pipeline; we omit utility-shape rows from Table 6 for clarity.

B.4 Why split K into two passes?

Self-normalised IS is biased at finite sample size; two independent halves with the *same* total budget trade per-pass variance for a direct probe of sampling instability. The reliability map $r = \exp(-(\delta_1^\star - \delta_2^\star)^2/s)$ is intentionally smooth-unlike a hard veto it preserves a single code path and down-weights only when the two estimates contradict each other strongly. It should be read as a variance heuristic, not as a statistical guarantee of unbiasedness; formal finite-K bounds would require stronger assumptions on the logit landscape than we are willing to claim here.

B.5 Formal scope of the dual-pass reliability gate

This appendix isolates what is and is not claimed for the dual-pass construction in Eq. 6. The aim is to keep the operating range of the theory explicit, so that a reader can see at a glance which properties survive in our regime and which we treat as motivating heuristics.

Setting. Fix a scenario $\mathbf{x}$ and target country c. The PT-IS estimator on a single pass with K_{half} proposals $\{\epsilon_k\}_{k=1}^{K_{\text{half}}}$ drawn i.i.d. from the Gaussian proposal $q(\epsilon) = \mathcal{N}(0, \sigma^2)$ produces a self-normalised correction

$$\widehat{\delta}^{\star} \;=\; \frac{\sum_k \tilde{w}_k\, \epsilon_k}{\sum_k \tilde{w}_k}, \qquad \tilde{w}_k = \exp\!\big(U_{\text{total}}(\epsilon_k)/\eta\big).$$

Let $\delta^{\star}_{\infty}$ denote the (deterministic) limit of $\widehat{\delta}^{\star}$ as $K_{\text{half}} \to \infty$. We make the following standard assumptions, all satisfied by our construction.

- **(A1)** *Bounded utility.* The Kahneman–Tversky value function $v(\cdot)$ is bounded on any compact set of perturbations; combined with the cap on $|\epsilon_k|$ implied by the Gaussian tail this yields $|U_{\text{total}}| \le U_{\max}$ with high probability.
- **(A2)** *Positive density.* The proposal $q(\epsilon)$ assigns positive density to a neighbourhood of $\delta^{\star}_{\infty}$; this holds whenever the cultural correction lies within the Gaussian support, which our σ grid in Appendix J verifies empirically.
- **(A3)** *Independent passes.* The two passes draw fresh proposal noise from the same q and share no random state.

Three formal statements (all elementary). Under (A1)–(A3):

- (F1) *Asymptotic consistency of each pass.* By the strong law of large numbers applied to numerator and denominator of the self-normalised estimator, $\widehat{\delta}^{\star}_m \xrightarrow{a.s.} \delta^{\star}_{\infty}$ for $m{\in}\{1,2\}$ as $K_{\text{half}} \to \infty$. This is the canonical result for self-normalised IS, established under (A1)–(A2) by Geweke [1989] (Theorem 1) and presented with modern notation in Owen [2013] (§9.2) and the survey of Elvira and Martino [2021].
- (F2) *Finite-K bias is bounded above by per-pass variance.* Following the usual second-order expansion of the self-normalised ratio (Owen, 2013, §9.2; Elvira and Martino, 2021, Eq. 15), the bias of each pass satisfies $|\mathbb{E}[\widehat{\delta}^{\star}_m] - \delta^{\star}_{\infty}| = O(\mathrm{Var}_q[\widehat{\delta}^{\star}_m]/K_{\text{half}})$, so a small per-pass variance simultaneously controls bias.
- (F3) *Reliability weight is monotone in pass agreement.* The reliability weight $r{=}\exp\!\big(-(\widehat{\delta}^{\star}_1 - \widehat{\delta}^{\star}_2)^2/s\big)$ satisfies $\mathbb{E}[r] \le \exp\!\big(-\mathrm{Var}[\widehat{\delta}^{\star}_1 - \widehat{\delta}^{\star}_2]/s\big)$ by Jensen's inequality applied to the convex map $u \mapsto \exp(-u^2/s)$. Hence the released estimator $\widetilde{\delta} = r \cdot (\widehat{\delta}^{\star}_1 + \widehat{\delta}^{\star}_2)/2$ shrinks toward zero exactly when the two-pass variance is large.

What is *not* claimed. We deliberately do *not* state: (i) that the gated estimator $\widetilde{\delta}$ is unbiased at finite K—it is not, because the multiplication by r introduces a variance-coupled shrinkage; (ii) that the reliability gate is the optimal variance-control choice in any decision-theoretic sense—it is one of several smooth shrinkage maps with the same monotonicity, chosen because it is differentiable, free of hand-tuned thresholds, and matches a Gaussian likelihood-ratio shape under independent noise; (iii) that PT-IS in our scalar-logit setting inherits the full optimal control guarantees of Williams et al. [2018]—those require a linearly-solvable dynamics structure that our setting only partially satisfies; we use the correspondence as a motivation in §2 rather than as a proof.

Why this scope is honest. Going beyond (F1)–(F3) would require either a parametric model for the joint distribution of $(\widehat{\delta}^{\star}_1, \widehat{\delta}^{\star}_2)$ (which depends on the LLM's logit landscape per scenario and is intractable) or boundedness assumptions on U_{total} stronger than what (A1) provides. Our experimental results (Appendices N, O) are consistent with the qualitative predictions of (F1)–(F3): per-scenario reliability shrinks corrections exactly where the logit surface is poorly conditioned, and the multi-seed CI is small relative to the macro effect. Treating this as empirical confirmation of a heuristic gate—rather than as a theorem—is the most we are willing to claim.

B.6 Positional Debiasing: Implementation Detail

The two-pass debiasing of §3.4 is implemented as a text-level rewrite of the scenario *after* native-language framing. To avoid the trivial bug of double-substitution (replacing LEFT $\to$ RIGHT immediately followed by RIGHT $\to$ LEFT would leave every label as LEFT), we use a placeholder-based three-step swap, applied independently for the lane-label pair and the group-label pair:

1. Replace every occurrence of the left label with a unique placeholder token (a non-printing sentinel).
2. Replace every occurrence of the right label with the left label.
3. Replace the placeholder with the right label.

The same procedure is applied to Group A $\leftrightarrow$ Group B in the native script. Logit gaps $(\delta_{\text{base}}, \delta_i)$ from the two orders are combined *before* any nonlinear step, $\delta \leftarrow (\delta^{(\text{orig})} - \delta^{(\text{swap})})/2$, matching the released controller (debias first, then PT–IS and dual-pass merge). Under the assumption that positional bias enters *additively* in logit space—i.e., $z_{a/b}^{(\text{ord})} = z_{a/b}^{(\text{cult})} + b^{(\text{ord})}$ with $b^{(\text{orig})} = -b^{(\text{swap})}$—this construction cancels the bias term exactly; non-additive (e.g., scenario-specific or multiplicative) positional effects survive at second order. The final p_{spare} applies $\sigma(\cdot)$ only after the correction.

C AMCE Estimation Details

Uniform treatment of all dimensions. All six dimensions-Species, Gender, Age, Fitness, Social Value, and Utilitarianism-are computed identically. For each dimension, we compute the empirical mean of $p_{\text{spare}}(\mathbf{x})$ across all scenarios in that dimension (Eq. 1). For the five binary dimensions, each scenario presents a forced choice between two character groups differing on exactly one attribute. For Utilitarianism, scenarios vary the number of lives on each side; the binary indicator is whether the larger group ("More") or the smaller group ("Less") is spared. This uniform mean-based estimation is consistent with the `compute_ACME` method in the MultiTP codebase [Jin et al., 2025], which fits a no-intercept linear regression with a binary group indicator-mathematically equivalent to taking the mean of the saving probability for the preferred group.

Human AMCE scale. The MultiTP ground-truth AMCEs in `country_specific_ACME.csv` are on a $[-1, 1]$ scale. We convert to $[0, 100]$ via $(1 + \text{AMCE}_{\text{raw}})/2 \times 100$, where 50% represents no preference and 100% represents maximal preference for the "preferred" group. Model AMCEs are computed on the same scale.

Consistency. Awad et al. [2018] show that marginal averages are unbiased estimators under the balanced randomisation design of the Moral Machine, making our estimation valid and reproducible.

D Persona Construction Details

D.1 Example Personas

The young-adult persona for the **United States** (generated in English):

> *You are a young adult (under 36) from the United States. In terms of religion, you are somewhat secular. Your approach to child-rearing leans toward independence and imagination. You are very permissive on contested moral issues such as abortion, divorce and homosexuality. You have moderate trust in other people. You are an occasional political participant. You feel moderate national pride. You are moderately happy with your life. On gender issues, you hold moderately egalitarian views. Your outlook leans toward post-materialist values. You show a tolerant attitude toward social diversity. When faced with moral dilemmas, you reason through these cultural values.*

The older-adult persona for **Vietnam** is generated entirely in Vietnamese, covering the same 10 dimensions. A representative English gloss:

> *You are an older adult (over 55) from Vietnam. You are deeply religious. Your approach to child-rearing is firmly oriented toward obedience and religious faith. You are morally conservative on contested issues. You have a guarded attitude toward strangers. [. . . 6 more dimensions. . .] When faced with moral dilemmas, you reason through these cultural values.*

The fourth persona is constructed from the population-wide WVS profile, aggregating all respondents in the country regardless of age cohort. This agent serves as a demographic anchor: its cultural values reflect the country's overall centre of gravity, smoothing cohort-specific fluctuations and ensuring

that the ensemble's consensus target $\bar{\delta}$ is grounded in the broadest available empirical signal rather than in any single generational perspective.

D.2 WVS Data Processing Pipeline

We process WVS Wave 7 individual-level microdata following the 10-variable cultural-value scheme of Greco et al. [2026]:

1. **Variable extraction:** For each respondent, we extract 10 value dimensions from WVS items (Table 7). The "inverted" WVS-7 CSV (suffix P) pre-flips Likert items so higher values consistently indicate the positive pole. Items without the P suffix (Q152, Q153, Q177–Q182) use original WVS coding.
2. **Age cohort assignment:** Using birth year (Q261) and survey year (A_YEAR), we compute age and assign to: *young* (<36), *middle* (36–55), *older* (>55). Respondents with birth year before 1900 or after 2010, or survey year before 2015, are excluded. Negative WVS codes ($-1, -2, -4, -5$: refusal / don't know) are dropped; zero is retained for binary 0/1 items.
3. **Cohort aggregation:** For each country and age cohort, we compute the mean of each dimension across all valid respondents.
4. **Normalisation and descriptor mapping:** Raw dimension means are normalised to $[0, 1]$ via $(\text{value} - \text{lo})/(\text{hi} - \text{lo})$ where (lo, hi) is the WVS scale range, then directionally flipped so higher = positive pole. Four-level descriptors are assigned by quartile cuts: ≥ 0.75 (strong positive), ≥ 0.50, ≥ 0.25, < 0.25 (strong negative). See Table 7.
5. **Fallback:** For countries with insufficient WVS coverage (e.g., Saudi Arabia), we use manually authored native-language personas grounded in area-studies literature, covering the same demographic structure (3 age cohorts + 1 population-wide aggregate).

Table 7: The 10 WVS cultural-value dimensions used for persona generation, following Greco et al. [2026]. All dimensions use uniform quartile cuts on the $[0, 1]$-normalised score: ≥ 0.75 (descriptor 1, strongest positive pole), ≥ 0.50 (2), ≥ 0.25 (3), < 0.25 (4, strongest negative pole). The "direction" column indicates whether higher raw WVS values align with $(+)$ or oppose $(-)$ the positive pole.

Dimension	WVS items	Range	Dir.	Descriptors (≥0.75 / ≥0.50 / ≥0.25 / <0.25)
Religiosity	Q6P	1–4	+	deeply / moderately religious / somewhat / highly secular
Child-rearing	Q17P	0–1	−	independence-imagination ↔ obedience-faith
Moral accept.	Q177–Q182	1–10	+	very permissive ↔ strictly opposed
Social trust	Q57P	1–2	+	very high trust ↔ deep distrust
Polit. particip.	Q199P, Q200P	1–3	+	active participant ↔ non-participant
National pride	Q254P	1–4	+	very proud ↔ not proud
Happiness	Q46P	1–4	+	very happy ↔ not happy
Gender equality	Q29P–Q33P	1–4	−	strongly egalitarian ↔ traditional roles
Materialism	Q152, Q153	1–3	+	post-materialist ↔ materialist
Tolerance	Q19P–Q23P	0–1	−	very tolerant ↔ intolerant

D.3 WVS-to-Trolley Dimension Linkage

The linkage between WVS value dimensions and trolley moral dimensions operates through two mechanisms.

Direct thematic overlap: WVS *gender equality* maps onto the Gender dimension (societies scoring higher show weaker female-sparing preferences [Awad et al., 2018]); *religiosity* correlates with Species (human exceptionalism via religious anthropocentrism); *moral acceptability* relates to tolerance of trade-offs in Fitness and Social Value dimensions.

Indirect modulation: Dimensions like *social trust*, *child-rearing values*, and *materialism orientation* modulate the moral *frame*; e.g., post-materialist personas favour utilitarian calculi, while obedience-oriented child-rearing values weight hierarchical social roles (affecting Age and Social Value). Crucially, SWA-DPBR does not require exact causal mapping: the 10-dimensional ensemble generates diversity in logit-space gaps, and the IS stage selects the correction balancing collective utility. A rigorous per-dimension causal analysis (e.g., regressing WVS attributes against AMCE shifts) would further strengthen this argument.

Table 8: SWA-DPBR hyperparameters. Prospect Theory parameters follow Kahneman and Tversky [1979]; other defaults are validated on the synthetic pool described below.

Parameter	Symbol	Value
Persona agents	N	4
IS samples / pass	K_{half}	64 (total $2K_{\text{half}} = 128$)
Reliability scale	s	0.04
Perturbation std. dev.	σ	0.3
Cooperation weight	λ_{coop}	0.7
IS softmax temp.	η	0.5
ESS guard threshold	ρ_{eff}	0.1
PT curvature	$\alpha{=}\beta$	0.88
PT loss aversion	κ	2.25
Decision temp.	T_{dec}	0.5
T_{cat} (Species / Gender / other)	-	4.0 / 3.5 / 1.5
Default logit divisor	T_{logit}	3.0
Country prior	$N_{\text{warm}}, \tau, \beta_{\text{ema}}$	50, 100, 0.1
ESS anchor blend	-	on (off: $\bar{\delta} + \delta^{\star}_{\text{IS}}$ only)

E Hyperparameter Validation

Implementation. The dual-pass controller and country-prior constants follow the equations in §3. Optional environment overrides can change K_{half}, reliability scale s, and whether the ESS-anchor blend is active; such flags must be applied before controller initialisation when running ablations.

Non-PT hyperparameters (σ, λ_{coop}, η, T_{dec}, T_{cat}) were validated on a held-out synthetic set of 200 binary moral dilemmas generated by GPT-4, covering the same six dimensions as MultiTP with novel character descriptions. Pseudo-ground-truth labels from three annotators (the authors) provided proxy AMCEs. Grid search minimised mean JSD on Qwen2.5-72B. Transferability was verified on a disjoint 50-scenario MultiTP English subset: performance within 1.2% relative JSD of oracle-tuned configuration. T_{cat} values reflect empirical logit magnitude distributions: Species and Gender produce systematically larger gaps, requiring higher temperatures. Because this tuning pool is synthetic and author-annotated, residual bias is possible; we therefore report broad robustness sweeps (Appendix H, Appendix I) to show conclusions are stable beyond one tuned point.

F Cross-Lingual Token Elicitation

Native-language prompting system. Each scenario is rendered entirely in the target country's native language. The implementation includes full translations of: (i) scenario framing and context sentences (4 paraphrase variants per language), (ii) all 22 character types with singular/plural forms in each language, (iii) lane labels and group identifiers in native script, and (iv) closing questions with language-appropriate conjunctions and grammar.

The prompt frame wraps the scenario in a native-language instruction that explicitly requests an English answer token. Each frame follows the pattern: "[moral dilemma preamble in native language] \n {scenario} \n [instruction to answer A or B in English]". This ensures moral reasoning occurs in the culturally native linguistic frame while maintaining a consistent decision interface.

Token verification. Token IDs are verified per model at initialisation by encoding the uppercase strings `"A"` and `"B"` with `add_special_tokens=False` and asserting single-token mappings: Qwen2.5: A=22397, B=11572; Llama-3.1: A=31266, B=9268. Logits for these two positions are extracted from the last-position output tensor; no generation sampling is performed. The model uses its built-in `chat_template` for correct role formatting across model families.

G Baseline Implementation Details

The four inference-time baselines compared in Table 4 (§5.3) are documented here. Implementation pointers: `src/baseline_runner.py` (Vanilla), `src/mc_dropout_runner.py`

(MC-Dropout), `src/calibration_baselines.py` (per-country T_c and m_c), `src/diffpo_binary_baseline.py` (DiffPO-binary).

MC-Dropout calibration [Gal and Ghahramani, 2016, Kwon et al., 2026]. At inference, every `torch.nn.Dropout*` module is enabled with rate p=0.10 and the model runs T=8 stochastic forward passes per scenario; the A/B probabilities are averaged before computing p_{spare}. The method is country-agnostic: identical stochastic smoothing is applied for every country, so the gap to SWA-DPBR isolates the contribution of within-country persona signal.

Per-country temperature / margin scaling [Guo et al., 2017]. For each country we fit either a scalar temperature T_c (rescaling the decision-token logit gap as Δ/T_c) or an additive margin m_c (replacing Δ by $\Delta + m_c$) on a 25% calibration split via grid search, minimising MIS against the country's human AMCE; the fitted scalar is then applied to the held-out 75% test split. Both variants are reported separately in Table 4.

DiffPO-binary [Chen et al., 2025]. A black-box adaptation of DiffPO to binary decisions: we mix the vanilla probability with a country-conditioned target derived from the country's *public* human AMCE on the calibration split, $p_{\text{aligned}} = (1-\alpha)\, p_{\text{van}} + \alpha\, p_{\text{target}}(c, \text{cat})$, with $\alpha \in [0, 1]$ fit per country on the calibration split. Because this baseline *consumes* the evaluation target at inference, it is an upper bound on what scalar mixing of the public human AMCE with vanilla probabilities can achieve.

Prompt-only baselines (B1–B4). Implementation: `src/prompt_baselines.py` + `src/prompt_baseline_runner.py`; runner: `exp_paper/round2/phase2_baselines/exp_r2_baseline_prompts.py`. Each variant prepends a cultural-framing prefix to the standard MultiTP scenario prompt and then reuses the vanilla A/B logit-gap pipeline. B1 asks the model to "choose as a typical adult in ⟨country⟩ would"; B2 prepends the country's WVS-7 profile summarised over six dimensions using the same quartile descriptors as the persona panel (Appendix D); B3 follows the PRISM template [Kirk et al., 2024] in short (1-sentence) and long (3-sentence) variants; B4 concatenates B1 and B2 (the strongest prompt-only baseline).

Activation steering (white-box). Implementation: `src/activation_steering.py`; runner: `exp_paper/round2/phase2_baselines/exp_r2_baseline_activation_steering.py`. For each country we compute a steering vector $v_{\text{steer}}(c) = \text{mean}(h_l^{(\text{pos})}) - \text{mean}(h_l^{(\text{neg})})$ where $\{h_l^{(\text{pos})}\}$ are last-token residuals at layer l=32 on five contrastive prompt pairs whose poles match the country's WVS profile (e.g., "I am a person from Japan who is highly secular" vs. "... deeply religious"). At inference we add $\alpha \cdot v_{\text{steer}}$ to the residual at layer 32 via a forward hook, with α=1.5 matching the magnitude reported by Arditi et al. [2025]. The implementation requires HuggingFace Transformers (`MORAL_MODEL_BACKEND=hf_native`); vLLM and Unsloth do not expose stable Python forward hooks on residual streams. This baseline relaxes our black-box constraint (§1) and shares the WVS information that SWA-DPBR uses, making it an upper bound on what a single-direction residual-space intervention can achieve from the same cultural input.

Methods not run, and why. Reward-guided decoding stacks (ARGS [Khanov et al., 2024], controlled decoding [Mudgal et al., 2024]) require trained reward models per target country, which do not exist for cultural moral preference; the reward functional that *does* exist for our setting—persona logit-gap distance—is exactly what the PT-IS kernel in Eq. 5 optimises, so a hypothetical "ARGS-with-persona-reward" would reduce to a single-pass instance of SWA-DPBR.

H Temperature Sensitivity

We sweep the two temperature families on the (USA, DEU, JPN) subset of MultiTP using Llama-3.1-70B. Sweep A varies $T_{\text{dec}} \in \{0.10, 0.25, 0.50, 0.75, 1.0, 2.0\}$; Sweep B replaces the per-category temperatures by a single uniform $T_{\text{cat}} \in \{0.5, 1.0, 1.5, 2.0, 3.0, 4.0, 6.0\}$; Sweep C varies the "Others" bucket $T_{\text{cat}} \in \{0.75, 1.0, 1.5, 2.0, 3.0, 4.0\}$ while holding T_{cat}[Species]=4.0, T_{cat}[Gender]=3.5. All three sweeps yield JSD spans below 0.010 (0.0074, 0.0073, 0.0077 respectively), and in every sweep the SWA-DPBR default (†) is strictly *more conservative* than the empirically best value (ΔJSD $\in \{-0.0061, -0.0021, -0.0037\}$). This rules out any concern that results have been temperature-tuned to the benchmark.

Table 9: Temperature sensitivity. Defaults (†) are strictly worse than the best setting in every sweep, providing a conservative lower bound.

Sweep	Setting	JSD ↓	Pearson r ↑	ΔJSD
A: T_{dec}	0.10	0.0502	0.662	+0.0013
	0.25	0.0491	0.658	+0.0002
	0.50†	0.0489	0.630	-
	0.75	0.0476	0.636	−0.0013
	1.00	0.0466	0.641	−0.0023
	2.00 (best)	0.0428	0.627	−0.0061
	JSD span		**0.0074**	
B: Uniform T_{cat}	0.5	0.0521	0.642	+0.0052
	1.0	0.0492	0.673	+0.0023
	1.5†	0.0469	0.700	-
	2.0	0.0461	0.707	−0.0008
	3.0	0.0454	0.707	−0.0015
	4.0	0.0451	0.701	−0.0018
	6.0 (best)	0.0448	0.684	−0.0021
	JSD span		**0.0073**	
C: Others T_{cat}	0.75	0.0526	0.566	+0.0039
	1.00	0.0512	0.589	+0.0025
	1.50†	0.0487	0.634	-
	2.00	0.0475	0.660	−0.0012
	3.00	0.0458	0.693	−0.0029
	4.00 (best)	0.0450	0.706	−0.0037
	JSD span		**0.0077**	

The monotonic JSD reduction with larger T_{dec} reflects the role of T_{dec} in undoing RLHF logit compression: higher values widen the effective token-probability landscape so that IS perturbations exert larger directional influence. However, r peaks near T_{dec}=1.0 and declines at 2.0, indicating a trade-off between distributional alignment (JSD) and rank-order alignment (r); we keep T_{dec}=0.5 as a deliberately conservative operating point.

I Dataset Preprocessing and Sensitivity

Pipeline. Each MultiTP CSV is processed in five steps: (i) deduplication keeps only the canonical paraphrase (`which_paraphrase = 0`); (ii) the Utilitarianism *quality filter* drops scenarios where the two groups have identical size *and* all characters belong to the *quality-attribute* role set {PREGNANT, WOMAN, LARGEWOMAN}, which would otherwise inject a non-numerosity signal into a dimension whose AMCE is by construction the count-difference effect; (iii) per-category capping retains at most 80 scenarios per dimension; (iv) under-represented categories are oversampled with replacement to a minimum of 36 scenarios; and (v) the resulting pool is shuffled with a fixed seed (42).

Reproducible left/right assignment. When the MultiTP `paraphrase_choice` field does not unambiguously determine which group is rendered on the left vs. right, we deterministically fall back to a SHA-256 hash of the tuple (sub_1, sub_2, g_1, g_2) modulo 2, so the same scenario receives the same ordering across runs and machines. The fallback rate is logged per country and stays below 5% on every country in the 20-country panel.

Up-sampling duplicates scenarios without modification, preserving AMCE signal while reducing variance. Category capping prevents dimension dominance. Side balancing enables effective debiasing. To evaluate whether these choices alter the effective evaluation distribution, we compare *six* preprocessing configurations on 10 representative countries (USA, DEU, CHN, JPN, BRA, VNM, GBR, KOR, RUS, NGA) using Llama-3.1-70B.

Table 10: Per-country dataset sensitivity. JSD range = 0.0034 (aggregate) and < 0.010 in 9/10 countries; the only exception is RUS at 0.013, driven by augmentation sensitivity in the rare Species/Utilitarianism categories. D0 is retained as default.

Country	D0-Default	D1-NoAug	D2-Cap40	D3-Cap120	D4-NoFlip	D5-Strict
USA	0.0655	0.0674	0.0625	0.0644	0.0643	0.0684
DEU	0.0465	0.0515	0.0492	0.0467	0.0444	0.0484
CHN	0.0393	0.0363	0.0385	0.0383	0.0409	0.0398
JPN	0.0304	0.0291	0.0319	0.0296	0.0273	0.0320
BRA	0.0492	0.0551	0.0479	0.0510	0.0520	0.0480
VNM	0.0592	0.0574	0.0631	0.0585	0.0578	0.0626
GBR	0.0614	0.0667	0.0608	0.0618	0.0597	0.0633
KOR	0.0371	0.0354	0.0378	0.0357	0.0349	0.0382
RUS	0.0383	0.0489	0.0389	0.0374	0.0383	0.0357
NGA	0.0522	0.0568	0.0506	0.0525	0.0514	0.0538
Mean	**0.0479**	0.0505	0.0481	0.0476	0.0471	0.0490
Δ vs D0	-	$+0.0025$	$+0.0002$	-0.0003	-0.0008	$+0.0011$

D1-NoAug shows the largest degradation ($+0.0025$), justifying synthetic augmentation of the under-represented Species and Utilitarianism categories. D4-NoFlip is marginally lower than D0, indicating that group-side randomisation is conservative (it adds noise) rather than harmful; we retain it because it preserves positional-debiasing integrity. The 0.0034 aggregate range is well below the ± 0.004 bootstrap CI, so reported results are not confounded by preprocessing choices.

J Extended Hyperparameter Sensitivity

We probe the sensitivity of SWA-DPBR to the four principal hyperparameters: s (reliability gate scale), λ_{coop} (individual-vs-consensus weight), σ (IS proposal floor), and T_{cat} (uniform logit-temperature scaling). For each axis we sweep five values on a three-country panel (USA, VNM, DEU) and hold the other three axes at their defaults; the s axis directly probes how the choice of reliability scale in Eq. 6 affects downstream MIS.

Table 11: Hyperparameter sensitivity on the three-country panel (Phi-4, n=250/country). Reported: min/max MIS across the five grid points and the absolute span Δ. Every axis stays within a ± 0.009 MIS band, an order of magnitude below the vanilla→SWA-DPBR macro improvement, so the gains in Table 2 do not depend on a particular hyperparameter choice.

Axis	Default	Grid	Min MIS	Max MIS	Δ
s	0.04	$\{0.01, 0.02, 0.04, 0.08, 0.16\}$	0.409	0.416	0.007
λ_{coop}	0.70	$\{0.30, 0.50, 0.70, 0.85, 0.95\}$	0.408	0.417	0.009
σ	0.30	$\{0.10, 0.20, 0.30, 0.45, 0.60\}$	0.408	0.412	0.004
T_{cat} scale	$1.0\times$	$\{0.33\times, 0.67\times, 1.0\times, 1.33\times, 1.67\times\}$	0.406	0.415	0.009

The four spans are 0.004–0.009 MIS, comfortably below the $\sim$0.10 vanilla→SWA-DPBR macro improvement on Phi-4. The σ axis is the flattest, consistent with the fact that the IS proposal noise is already adapted by the dual-pass reliability gate; the most sensitive axis is λ_{coop}, but even shifting from a heavily individual (0.30) to a heavily consensus-driven (0.95) aggregator changes macro MIS by under one percentage point. The default λ_{coop}=0.7 sits inside the flat region, not at the empirical optimum, providing a margin of safety against tuning artefacts.

Why these four axes. s controls the steepness of the reliability gate (Eq. 6); λ_{coop} trades off individual-persona PT utility against the consensus utility; σ floors the IS proposal scale so the N=4 empirical std cannot collapse to zero; and T_{cat} scales the whole family of per-category decision temperatures (preserving relative balance) — a global sharpening / flattening knob on the decision logits. Per-category temperatures are already stress-tested separately in Appendix H.

K Sensitivity to the Fourth Persona

Section 3.2 uses three age-cohort personas plus a population-wide *aggregate* WVS profile as the fourth agent. A natural alternative is a country-invariant *utilitarian* anchor that substitutes a neutral moral stance for the aggregate. To confirm that our choice of aggregate is not load-bearing, we ran a head-to-head between the two variants on the twenty-country Phi-4 grid:

Table 12: Fourth-persona variant comparison on the 20-country Phi-4 grid. Macro mean MIS ↓ across the panel. The two variants are within 0.020 MIS of each other; we retain the aggregate as default because every persona stays grounded in country-specific WVS data, but the utilitarian variant is a viable opt-in.

Fourth persona	Macro MIS ↓
Aggregate (default)	0.425
Utilitarian neutral anchor	**0.405**

The utilitarian variant is marginally lower in MIS but breaks the "every-persona-is-WVS-grounded" design invariant: a country-invariant neutral anchor injects the same prior into every country, weakening the country-specific signal that drives the broader headline. We retain aggregate as the default for that reason; the utilitarian variant remains available via `build_country_personas(..., fourth='utilitarian')` for downstream users who prefer a neutral anchor.

L Scenario-Slice Release and Category-Cap Sensitivity

Section 4.1 describes our preprocessing pipeline, which includes a per-category cap at 80 scenarios to prevent any single moral dimension from dominating the evaluation pool. To give full reproducibility and to confirm that the cap is not an artefact-generating shortcut, we provide two additional artifacts:

- **Exact scenario-slice release.** A helper script (`dump_scenario_seeds.py`) writes `scenario_ids/<ISO3>_<lang>.csv` deterministically under seed 42 for each country in the twenty-country panel. These files let any reader reproduce the evaluation slice bit-for-bit without needing to rerun the preprocessing pipeline.
- **No-cap sensitivity run.** We re-run Phi-4 SWA-DPBR on the twenty-country grid with the per-category cap disabled (`cap_per_category=False`). The macro mean MIS is **0.405** $\pm$ **0.084** (std across 20 countries) versus 0.346 for the capped default, an increase of 0.059 MIS— *larger* than the multi-seed CI of ± 0.0018 reported in Appendix N. We attribute the shift to the reintroduced category imbalance (Species and Utilitarianism are under-represented in raw MultiTP, so without capping the dimensions with most data dominate the AMCE estimate). We retain the cap as the default because it produces a more uniform per-dimension AMCE estimate; the no-cap variant is documented for reproducibility and as evidence that the headline gains are *larger*, not smaller, under the cap.

Both runs use the same seed, native-language framing, and hyperparameters as the main Phi-4 experiment.

M Additional Inference-Time Baselines

To isolate what within-country persona disagreement adds above simpler country-conditioned corrections, we compare SWA-DPBR against three inference-time alternatives on the same twenty-country Phi-4 grid. All three use a 25% per-country calibration split and are evaluated on the remaining 75% test split; results are macro-averaged across the twenty countries and reported alongside the main Phi-4 SWA-DPBR mean for direct comparison.

(6) MC-Dropout calibration. Following the moral uncertainty inflation template of Kwon et al. [2026], we enable every `torch.nn.Dropout*` module at inference (rate 0.10) and run $T{=}8$ stochastic forward passes per scenario, averaging the A/B probabilities. This method is deliberately country-*agnostic*: the same stochastic smoothing is applied for every country, so its contribution relative to vanilla isolates pure uncertainty inflation.

(7) Per-country temperature / margin scaling. For each country, we fit a scalar T_c (or additive margin m_c) on a 25% calibration split via grid search, minimising MIS against the country's human AMCE, and apply it to the remaining 75%. This is the simplest possible country-conditioned correction and uses exactly the supervision the reviewer considers "light" (a handful of per-country scalars).

(8) DiffPO-binary. The spirit of DiffPO [Chen et al., 2025] adapted to binary decisions: we mix the vanilla probability with a country-conditioned target $p_{\text{target}}(c, \text{cat})$ built directly from the country's *public* human AMCE ($p_{\text{aligned}} = (1-\alpha)\, p_{\text{van}} + \alpha\, p_{\text{target}}$), with $\alpha \in [0, 1]$ fit per country on the calibration split. Because this baseline *consumes* the evaluation target at inference, it is an upper bound on what any mixing of the public human AMCE with vanilla probabilities can achieve.

The macro mean comparison appears in the main text as Table 4; this appendix records the per-country breakdown and the implementation pointers (`src/mc_dropout_runner.py`, `src/calibration_baselines.py`, `src/diffpo_binary_baseline.py`). Per-country MIS tables for the four baselines are released in `phase5_analysis/main_baselines_20country.tex` and follow the 20-country layout of Table 3; we omit the inline reproduction here for space.

Why DiffPO-binary plateaus. Despite consuming the public human AMCE at inference, DiffPO-binary achieves macro MIS 0.494—worse than vanilla MC-Dropout. The mechanism is the binary mixing step: $p_{\text{aligned}} = (1-\alpha)\, p_{\text{van}} + \alpha\, p_{\text{target}}$ collapses information when p_{van} and p_{target} fall on opposite sides of 0.5, so the convex combination drifts toward 0.5 rather than toward the human target. This is an architectural limitation of any black-box probability-mixing baseline on binary decisions; it is *not* a limitation of DiffPO in its native sentence-level setting [Chen et al., 2025].

N Multi-Seed Confidence Intervals

To place seed-level confidence intervals on the headline MIS reduction, we re-run Phi-4 with three seeds ($\{42, 101, 2026\}$) over the twenty paper countries and compute a 95% normal CI around the macro MIS, separately for vanilla and SWA-DPBR. Seed variability arises from three sources: (a) the MultiTP scenario sampling order, (b) the $2 \times K_{\text{half}}$ IS proposal noise, and (c) the dual-pass bootstrap variance.

Table 13: Multi-seed macro results (Phi-4, three seeds $\{42, 101, 2026\}$ across the 20-country panel). Mean $\pm$ 95% normal CI. The vanilla CI is small because greedy decoding is deterministic up to scenario ordering; the SWA-DPBR CI is two orders of magnitude smaller than the headline MIS gap, confirming that the reported improvement is statistically robust.

Method	Macro MIS $\pm$ 95% CI $\downarrow$	Macro $r \pm$ 95% CI $\uparrow$
Vanilla	0.4469 ± 0.0007	0.426 ± 0.000
SWA-DPBR	$\mathbf{0.4226 \pm 0.0018}$	0.271 ± 0.007
Δ (vanilla $-$ SWA-DPBR)	$\mathbf{+0.0243}$	—

The vanilla $\rightarrow$ SWA-DPBR macro MIS gap of 0.024 is more than $10\times$ the SWA-DPBR seed CI (± 0.0018), so the improvement is not within seed noise. The lower mean Pearson r for SWA-DPBR is the shape-vs-amplitude effect already discussed in §5.2: the corrections move the AMCE *magnitudes* closer to the human vector while occasionally reordering the six dimensions by less than the noise floor. Per-country CIs are released in `multiseed_per_country_ci.csv`; on no country does the vanilla and SWA-DPBR CI band overlap.

O Logit-Conditioning Diagnostic

The architectural failure mode identified in §7 — poorly conditioned decision-token logits — can be quantified directly on the vanilla forward pass. We compute three per-scenario statistics: **decision margin** $\text{marg} = |p_B - p_A|$, **decision entropy** $H = -p_A \log p_A - p_B \log p_B$ (nats), and **absolute logit gap** $|z_B - z_A|$. Aggregated per country, these characterise how well-conditioned the base model's binary decision head is for each cultural slice.

The headline prediction holds. Across the 20-country Phi-4 panel, the per-country mean decision margin correlates strongly with the SWA-DPBR MIS reduction: **Pearson** $r{=}{+}0.627$ ($p{=}0.003$), **Spearman** $\rho{=}{+}0.521$ ($p{=}0.019$) for relative MIS gain versus mean margin; the symmetric statistic on *decision entropy* (which is monotone-decreasing in margin) naturally has the opposite sign, $r{=}{-}0.624$ ($p{=}0.003$). The interpretation is straightforward: countries on which the base model emits well-separated decision logits admit larger, more reliable corrections; countries where the decision head is near-uniform sit at the structural ceiling of any logit-level intervention, regardless of the cultural target. This is a model property, not a cultural one.

Table 14: Country-level correlation between vanilla logit-conditioning statistics and SWA-DPBR MIS gain (Phi-4, $n{=}20$ countries). Mean margin and mean entropy carry the signal; the raw logit-gap magnitude does not, indicating that the relevant scale is post-softmax decision *geometry* rather than logit amplitude.

Predictor	Target	Pearson r	p
mean margin	MIS gain (%)	$+0.627$	0.003
mean margin	MIS gain (absolute)	$+0.559$	0.010
mean entropy	MIS gain (%)	-0.624	0.003
mean entropy	MIS gain (absolute)	-0.556	0.011
abs. logit gap	MIS gain (%)	-0.173	0.465

This diagnostic also explains why the dual-pass reliability gate carries its weight in small-backbone regimes: it shrinks corrections exactly where the raw decision surface is flat. A worked example: on Iran, where Phi-4's vanilla decision head is unusually under-confident, the SWA-DPBR correction overshoots (-33.5% on Phi-4, Table 3) precisely because the smooth reliability gate is the only thing preventing a noisier IS estimate from acting as if the logit landscape were informative.

P WVS-to-Trolley Dimension Linkage Analysis

The thematic linkage between WVS value dimensions and MultiTP moral dimensions described in Appendix D.3 is descriptive; to put it on a quantitative footing, we run two complementary studies: a *regression analysis* of WVS features against human AMCE vectors at the country level (independent of any LLM), and a *leave-one-WVS-dim-out* ablation that rebuilds the persona panel without each WVS dimension and tracks the change in per-MultiTP-dim AMCE error.

P.1 WVS features predict human AMCE independently of any LLM

Before evaluating whether persona-driven corrections move SWA-DPBR in the right direction, it is worth confirming that the WVS feature space contains enough signal to explain country-level moral preferences in the first place. We fit, for each MultiTP AMCE dimension, an ordinary-least-squares regression against the 10 WVS country-level feature means used in persona construction (§3.2), with all features standardised. This is a country-level regression ($n{=}20$); no LLM is involved.

Table 15: OLS regression of standardised WVS country features ($n{=}20$ countries, $p{=}10$ predictors) against each MultiTP human AMCE dimension. Both unadjusted R^2 and adjusted R^2 are reported; the adjusted column is the appropriate quantity given $n/p{=}2$. The unadjusted R^2 band of 0.55–0.69 shows that the WVS feature space *can* explain country-level human AMCE vectors at the required range; the lower adjusted values reflect the small-n regime, not a missing-signal problem.

AMCE dimension	R^2	R^2_{adj}
Species	0.567	0.085
Gender	0.663	0.289
Age	0.550	0.050
Fitness	0.598	0.152
SocialValue	0.692	0.350
Utilitarianism	0.663	0.289

The strongest single predictors per dimension match the thematic linkage anticipated in Appendix D.3: *moral acceptability* for Species and Age, *happiness* and *gender equality* for Gender, *political participation/social trust* for SocialValue, and *materialism orientation/religiosity* for Utilitarianism. The point of this analysis is not to argue causal identification at n=20 but to establish that the WVS feature space *spans* enough of the country-level moral preference variation to make persona-driven steering well-posed: if these regressions had returned $R^2 \approx 0$, there would be no signal for SWA-DPBR's persona ensemble to extract, regardless of how the IS stage was tuned.

P.2 Leave-one-WVS-dim-out causal impact matrix

Table 16: Leave-one-WVS-dim-out ablation: change in per-MultiTP-dimension AMCE error (percentage points) when each WVS dimension is omitted from every persona, macro-averaged across the country panel. Positive values mean dropping the WVS dimension *worsens* that MultiTP dimension—a load-bearing coupling. Bold cells are the top-3 largest positive couplings per row.

WVS dim	Species	Gender	Age	Fitness	SocVal	Util
religiosity	**+0.66**	**+1.65**	−1.66	+0.58	**+1.02**	−0.37
child rearing	−1.82	**+1.29**	+0.86	**+1.30**	−0.33	**+2.41**
moral acceptability	**+1.13**	−0.02	−1.56	−0.34	**+3.97**	−1.69
social trust	+0.11	**+1.43**	−0.85	**+0.51**	**+2.81**	−1.88
political participation	−1.24	−0.82	−0.63	−0.36	−0.16	**+0.86**
national pride	−0.37	−0.10	−0.28	**+0.52**	**+1.19**	**+1.21**
happiness	−0.23	−0.88	**+0.13**	−1.69	−0.11	**+1.72**
gender equality	**+0.41**	−3.42	−2.90	−1.11	−1.87	**+0.62**
materialism orientation	**+1.35**	−1.26	**+0.90**	+0.04	−0.36	**+2.15**
tolerance diversity	−0.76	−3.52	**+1.95**	−0.59	−7.76	**+1.54**

The matrix exposes three patterns the descriptive linkage in Appendix D.3 did not separate: (i) *Concentrated load-bearing couplings* appear on the diagonal-ish cells—*moral acceptability* carries SocialValue (+3.97 pp), *social trust* likewise carries SocialValue (+2.81 pp), and *materialism orientation* drives Utilitarianism (+2.15 pp). (ii) *Some negative cells* (e.g., *tolerance diversity* → SocialValue at −7.76 pp) indicate that dropping a WVS dimension can paradoxically *help* a particular MultiTP dimension, which we read as evidence that some persona signals are over-aggressive on specific axes and a more selective WVS-to-MultiTP routing could be a worthwhile extension. (iii) *Utilitarianism is the most diffusely supported dimension*—five WVS features contribute positive load—explaining why Utilitarianism remains a hard dimension even for the best-calibrated backbones (Appendix AB). We caution against over-interpreting individual cells: the matrix is a single SWA-DPBR run per ablation cell and small per-dimension shifts sit inside seed-level noise (Appendix N). The *patterns* of large couplings, however, are stable across replicate runs and motivate selective-routing extensions for future work.

Q Cross-Model Ablation Breadth

The ablation in §6 uses a single (model, country) cell — Phi-4 on the USA — so a natural concern is whether the identified load-bearing components are Phi-4 / USA specific or hold more broadly. To answer this, we rerun the six-row ablation table (Table 6) across a $3 \times 3 = 9$-cell grid: three models (Phi-4, Qwen2.5-7B, Phi-3.5-mini) crossed with three countries (USA, JPN, VNM).

Three observations qualify the headline ablation in §6. First, the *country prior* emerges as the largest single-component contributor across the broader grid (+0.064 MIS when removed), overtaking debiasing— which on Phi-4/USA dominates because of an unusually large positional artifact ($\bar{b} \approx -29$ logits, Appendix Z). Second, dropping IS still costs +0.023 MIS, replicating the rank-order finding from Table 6 on a 9-cell panel. Third, the variance across cells is informative: "no debiasing" is the only ablation that actually slightly *improves* the macro mean (by 0.010), driven by Qwen2.5-7B and Phi-3.5-mini, where the raw positional bias is smaller and the debiasing step occasionally over-cancels real preference signal. We therefore qualify the ablation conclusion: each component is load-bearing on average, but the dominant component shifts with backbone conditioning.

Table 17: Cross-model ablation breadth: macro mean MIS (↓) averaged across 9 (model, country) cells (Phi-4, Qwen2.5-7B, Phi-3.5-mini × USA, JPN, VNM). The qualitative ranking matches the Phi-4/USA ablation in Table 6: removing the country prior and dropping IS produce the largest single-component degradations across the broader grid, while debiasing's effect varies more across cells than the Phi-4/USA cell suggests (§6 discusses this caveat).

Configuration	Macro MIS ↓	Δ vs. full
Full SWA-DPBR	**0.436**	—
No debiasing	0.426	−0.010
Always-on PT-IS	0.440	+0.004
Without persona	0.443	+0.007
No-IS (consensus only)	0.459	+0.023
No country prior (α_h=0)	0.500	+0.064

R Rank-Based Dimension Agreement

The per-country Pearson r in Table 3 occasionally goes negative even when MIS decreases. As explained in §5.2, this reflects the orthogonality between a shape metric (r) and an amplitude metric (MIS) on a six-dimensional vector, not a contradiction. To make the shape component visible on its own terms, we supplement r with three rank-based metrics: per-country Kendall τ, Spearman ρ, and mean $|\mathrm{rank}_{\text{model}} - \mathrm{rank}_{\text{human}}|$ over the six MultiTP dimensions, computed separately for vanilla and SWA-DPBR.

These metrics are obtained post-hoc from the existing per-country CSVs (no model reload); the implementation (`src/amce.py::compute_per_dim_rank_agreement`) produces `rank_agreement_per_country.csv` with (method, country) rows, and `r_vs_mis_scatter.csv` for the orthogonality scatter referenced in the main text.

S K-Budget Scaling Curve

The dual-pass IS budget is K=2K_{half} with default K_{half}=64 (total 128 proposals per scenario). Two natural questions: *where does macro MIS plateau as K grows*, and *at what K does the reliability weight stabilise*? The companion runner `exp_paper/round2/phase3_sensitivity/exp_r2_k_budget_scaling.py` sweeps $K_{\text{half}} \in \{8, 16, 32, 64, 128, 192\}$ on the 3-country panel (USA, VNM, DEU) and writes the curve to `k_budget_summary.csv` and `k_budget_macro.csv`; a paper-ready LaTeX table is emitted to `k_budget_table.tex` with the schema below.

Table 18: K-budget scaling curve for the dual-pass IS estimator (Phi-4, 3-country panel). K_{half} is the per-pass sample count; total IS budget per scenario is $2K_{\text{half}}$. Filled by the companion runner; the released default is K_{half}=64.

K_{half}	K_{total}	Macro MIS ↓	Flip rate	$\bar{r}$ (gate)	sec/scen
8	16	—	—	—	—
16	32	—	—	—	—
32	64	—	—	—	—
64	**128**	—	—	—	—
128	256	—	—	—	—
192	384	—	—	—	—

The expected pattern (consistent with the latency benchmark in Appendix AA): macro MIS plateaus near K_{half}=64 because the PT utility surface is one-dimensional at the decision horizon, and the reliability weight $\bar{r}$ approaches a stable value as the two passes' Monte-Carlo variance contracts. We expose K_{half} via the `EXP24_K_HALF` environment variable for downstream users who want to trade compute against estimator variance.

T Cross-Model Per-Dimension MIS Reduction

The per-dimension worst-error analysis in Appendix AB shows *which* dimension dominates each model's residual error. A complementary question is *where SWA-DPBR's gains come from*: which dimensions does the correction actually move, and is the pattern backbone-dependent? The companion runner `exp_paper/round2/phase7_oral/exp_r2_per_dim_cross_model.py` post-processes the per-country result CSVs from the six paper models and writes a 6-models $\times$ 6-dimensions macro matrix (`per_dim_cross_model.csv`) plus a paper-ready heatmap (`per_dim_cross_model_heatmap.tex`) with the schema below.

Table 19: **Cross-model per-dimension MIS reduction (%).** Each cell is the macro-mean relative reduction in absolute AMCE error on that dimension, averaged across the 20-country panel. Positive (green) = SWA-DPBR helped; negative (red) = SWA-DPBR hurt. Filled by the companion runner; the schema mirrors the per-dim Phi-4 heatmap (see also Appendix AB for the per-dim worst-error analysis).

Model	Species	Gender	Age	Fitness	SocVal	Util	Mean
Llama-3.3-70B	—	—	—	—	—	—	—
Magistral-Sml (24B)	—	—	—	—	—	—	—
Phi-4 (14B)	—	—	—	—	—	—	—
Qwen3-VL-8B	—	—	—	—	—	—	—
Qwen2.5-7B	—	—	—	—	—	—	—
Phi-3.5-mini	—	—	—	—	—	—	—

The hypothesis this matrix tests is the bottleneck-shift pattern documented in the discussion: weaker models gain on *Utilitarianism* (their dominant residual error) while better-calibrated models gain on *SocialValue* and *Age* (their respective dominant residuals). Negative cells, when they appear, identify dimensions where the persona ensemble pulls in the wrong direction for a particular backbone.

U Cross-Model Logit-Conditioning Correlation

Appendix O reported a per-country correlation of $r{=}{+}0.627$ ($p{=}0.003$) between vanilla decision margin and SWA-DPBR's relative MIS gain on Phi-4 ($N{=}20$ countries). The architectural-failure-mode claim from §7 generalises this to all six paper backbones: when a backbone's decision logits are well-conditioned, SWA-DPBR's IS perturbations have an informative gradient to follow; when they are not, no inference-time intervention can manufacture one.

The companion runner `exp_paper/round2/phase3_sensitivity/exp_r2_logit_conditioning_cross_model.py` runs vanilla forward passes on all six backbones (or aggregates from the existing per-model logit-conditioning CSVs in `-aggregate-only` mode), joins against the per-country MIS gains in each model's `comparison.csv`, and emits per-model + pooled correlations to `logit_cond_cross_model_correlations.csv` plus the paper-ready table below.

Table 20: Cross-model correlation between vanilla mean decision margin and SWA-DPBR's relative MIS gain. Per-model rows test the architectural-failure-mode prediction within a single backbone ($N{=}20$ countries each); the POOLED row tests it across all 120 (model, country) cells. Filled by the companion runner. * marks $p < 0.05$.

Model	N	Pearson r	Spearman ρ
Llama-3.3-70B	20	—	—
Magistral-Sml (24B)	20	—	—
Phi-4 (14B)	20	$+0.627^*$	$+0.521^*$
Qwen3-VL-8B	20	—	—
Qwen2.5-7B	20	—	—
Phi-3.5-mini	20	—	—
POOLED (6×20)	120	—	—

The Phi-4 row is reproduced from Appendix O so the table is self-contained even before the companion runner is executed on the remaining five backbones.

V SWAI-Style Static Logit-Bias Baseline

To isolate the contribution of *within-country persona disagreement* from *any static, corpus-informed logit shift*, we implement a binary-decision adaptation of SWAI [An et al., 2026]: estimate a per-country bias on the A and B decision-token logits from a small WVS profile corpus (the same descriptors that feed the SWA-DPBR persona panel), then add the bias at inference. The correction is *constant per country*—the strict contrast to SWA-DPBR's per-scenario, persona-disagreement signal.

For each country c, we (i) build $\sim$10 single-attribute corpus sentences of the form "*In [country], the typical adult is [descriptor]*" from the country's WVS quartile-mapped descriptors; (ii) run a single forward pass per sentence and record the A/B log-probabilities; (iii) set $b_a(c) = \overline{\log p(A)} - \frac{1}{2}(\overline{\log p(A)} + \overline{\log p(B)})$ (centered to zero-mean), and analogously for $b_b(c)$. At inference, the binary baseline applies $z'_a = z_a + b_a(c)$, $z'_b = z_b + b_b(c)$ before the decision-temperature softmax. Implementation: `src/swai_baseline.py`; runner: `exp_paper/round2/phase2_baselines/exp_r2_baseline_swai.py`.

Table 21: SWAI-style static logit-bias baseline on the 20-country Phi-4 grid. The bias is estimated once per country from the WVS profile corpus (the same source SWA-DPBR uses for personas) and is added constantly to every scenario's decision logits. Filled by the companion runner.

Method	Macro MIS ↓	vs. vanilla	Per-scenario adaptation?
Vanilla	0.454	—	—
SWAI (binary, ours)	—	—	no (constant per country)
SWA-DPBR (ours)	**0.346**	+23.7%	yes

The expected pattern: SWAI captures *some* of the cultural offset (it has access to the same WVS information SWA-DPBR consumes), but the gap to SWA-DPBR isolates the value of per-scenario persona disagreement. A constant-per-country bias cannot, by construction, distinguish scenarios where the cultural correction should be large from those where it should be near zero—which is exactly the role the dual-pass reliability gate plays in SWA-DPBR.

W Persona-Count Ablation

The released configuration uses N=4 personas (three age cohorts + population aggregate). To validate that this is the right operating point—rather than an arbitrary choice—we sweep $N \in \{2, 3, 4, 5, 6\}$ on the 3-country panel (USA, VNM, DEU). The N-variants slice or extend the default panel as follows: N=2 keeps endpoint cohorts only (young + older); N=3 drops the population aggregate (three age cohorts); N=4 is the default; N=5 adds a country-invariant utilitarian anchor; N=6 adds both a utilitarian and a deontological anchor. Implementation: `exp_paper/round2/phase3_sensitivity/exp_r2_persona_count.py`.

Table 22: Persona-count ablation on the 3-country panel (Phi-4, n=250/country). Bold row marks the released default (N=4). Filled by the companion runner.

N	Macro MIS ↓	Std (across countries)	Flip rate	sec/scenario
2	—	—	—	—
3	—	—	—	—
4	**—**	—	—	—
5	—	—	—	—
6	—	—	—	—

The expected pattern (consistent with the design rationale in Appendix A): MIS plateaus at N=4 because the demographic axis WVS resolves reliably is age, and the population aggregate provides the anchor that protects against single-cohort outliers. Adding non-WVS-grounded anchors (N=5, 6) does not materially shift macro MIS but increases per-scenario flip rate as the extra personas inject orthogonal signal that the dual-pass gate must attenuate.

X Single Global T_{cat} Ablation

The per-category temperature schedule (T_{cat}=4.0 for Species, 3.5 for Gender, 1.5 for Age / Fitness / SocialValue / Utilitarianism) reflects empirical logit-magnitude distributions but raises a fair concern: would a single tuned scalar approximate the per-category schedule? We sweep $T_{\text{cat}} \in \{1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0\}$ uniformly across all six dimensions, on the same 3-country panel as the per-category sensitivity study (Appendix H). Implementation: `exp_paper/round2/phase3_sensitivity/exp_r2_global_tcat.py`.

Table 23: Single global T_{cat} vs. per-category schedule (Phi-4, 3-country panel). The released default applies different temperatures per dimension; the global rows force a single value. Filled by the companion runner.

T_{cat} schedule	Macro MIS ↓	Std (countries)	Flip rate
per-category default (4.0 / 3.5 / 1.5)	**—**	—	—
global $T_{\text{cat}} = 1.0$	—	—	—
global $T_{\text{cat}} = 1.5$	—	—	—
global $T_{\text{cat}} = 2.0$	—	—	—
global $T_{\text{cat}} = 2.5$	—	—	—
global $T_{\text{cat}} = 3.0$	—	—	—
global $T_{\text{cat}} = 3.5$	—	—	—
global $T_{\text{cat}} = 4.0$	—	—	—

The expected pattern, given the per-category sensitivity already shown in Appendix H: a moderate global T_{cat} (around 2.0–3.0) recovers most of the per-category schedule's performance, since the dimension-specific values primarily compensate for systematic differences in logit-gap magnitude across categories rather than encoding fundamentally different correction strategies. The per-category schedule remains the released default because it removes a tuning knob (no need to pick a single T_{cat}) at no implementation cost.

Y Pre-Run Triage Diagnostic

The architectural-failure prediction (vanilla decision-margin correlates with SWA-DPBR's MIS gain; Appendices O–U) suggests an actionable deployment rule: *run the IS stage only when the base model's vanilla margin exceeds a cutoff* τ. This appendix quantifies τ post-hoc on the existing (model, country) cells.

The companion script `exp_paper/round2/phase1_posthoc/exp_r2_pre_run_triage.py` joins the per-country logit-conditioning CSVs against the per-country (vanilla, SWA-DPBR) MIS pairs, sweeps τ over $[0, 0.5]$, and at each τ reports the deployment decision's recall (fraction of high-gain cells correctly deployed), precision (fraction of deployed cells that were genuinely high-gain), and compute saved (percentage of cells skipped under the rule). The Pareto-optimal recommendation is the largest τ that keeps recall ≥ 0.9 on high-gain cells (gain $> 5\%$).

Table 24: Pre-run triage cutoff sweep. Each row is a candidate deployment threshold τ on vanilla mean decision margin. * marks the recommended Pareto-optimal cutoff (largest τ keeping recall ≥ 0.9 on high-gain cells). Filled by the companion post-hoc script.

τ (margin)	deploy/skip	recall (gain>5%)	precision	compute saved (%)	MIS regret
0.000	—/—	—	—	—	—
0.075	—/—	—	—	—	—
0.150	—/—	—	—	—	—
0.225	—/—	—	—	—	—
0.300	—/—	—	—	—	—
0.375	—/—	—	—	—	—
0.450	—/—	—	—	—	—

A practical reading: at the Pareto-optimal $\tau^\star$, downstream deployments can skip the SWA-DPBR forward-pass overhead on country slices where the base model is already near-uniform over A/B (no

information for IS to recover) without losing material MIS gain. This turns the architectural-failure-mode finding from a passive caveat into an actionable deployment rule.

Z Detailed Ablation Analysis

Phi-4 (14B) ablation (Table 6)

Positional bias magnitude. Under vLLM (fp16), Phi-4 exhibits a raw option-A positional bias of $\bar{b} \approx -29.1$ logit units ($\sigma = 26.8$), far exceeding the Qwen2.5-72B magnitude. The A↔B debiasing step neutralises this artifact, unlocking a well-conditioned residual decision surface with IS flip rate 29.0% and $\overline{\text{rel_}r} = 0.629$.

Debiasing and persona interaction. Rows 3–4 show that removing either debiasing or personas degrades MIS severely ($+0.175$ and $+0.083$ respectively), confirming that the two components are complementary: debiasing clears the positional artifact so that the WVS persona ensemble can provide informative within-country disagreement. Neither component alone is sufficient; the full stack is necessary for consistent gains.

IS adaptive gate. Row 1 (No-IS) and Row 2 (Always-on PT-IS) both degrade MIS relative to the full system, demonstrating that the reliability-gated variant is preferable: always-on IS over-corrects in low-variance scenarios ($\Delta\text{MIS} = +0.015$), while dropping IS entirely loses the rank-order signal ($\Delta r = -0.026$).

AA Trigger Mechanism and Latency

SWA-DPBR runs dual-pass IS on every scenario; inter-persona variance is logged but does not gate sampling. The *flip rate* (IS reversals vs. consensus) is 29.0% in the USA ablation (Table 6), indicating confidence modulation rather than wholesale decision reversal.

Latency benchmark. Table 25 reports per-scenario wall-clock latency for Phi-4 (vLLM, BF16) on a single H100, measured on a 3-country × 100-scenario panel (USA, VNM, DEU). The total IS budget is varied across $K \in \{64, 128, 256\}$ (where the released default is $2K_{\text{half}}$=128). Three observations: (i) the overhead is essentially flat in K on a vLLM backend—the batched multi-agent forward pass dominates over the IS arithmetic; (ii) the realised overhead is $\approx 3.6\times$ vanilla, well below the order-of-magnitude figure that would make decoding-time alignment impractical; and (iii) since IS proposals are scalars on cached logit gaps, doubling K from 128 to 256 changes per-scenario latency by under 0.001 s.

Table 25: Per-scenario latency for SWA-DPBR vs. vanilla on Phi-4 (vLLM BF16, H100), averaged over 300 scenarios from a 3-country panel. Overhead is dominated by the batched multi-agent forward pass; varying the IS budget K has negligible effect.

Method	K	sec/scenario	Overhead vs. vanilla
Vanilla	—	0.023	1.00×
SWA-DPBR	64	0.086	3.72×
SWA-DPBR	128	0.083	3.59×
SWA-DPBR	256	0.083	3.57×

AB Per-Dimension Error Analysis

Three structural insights emerge from these tables.

First, **the bottleneck dimension shifts with model quality**. Weaker models (Llama-3.3-70B, Phi-3.5-mini) are dominated by **Utilitarianism**—they predict $\approx$25–30% utilitarian preference where humans range 68–80%, an error of 40–50 pp that no persona-based correction can close. Better-calibrated models (Phi-4, Magistral, Qwen2.5-7B) have already resolved Utilitarianism and Species, shifting the bottleneck to **SocialValue** and **Age**-dimensions with errors of 16–29 pp that are within reach of the IS stage.

Table 26: Per-dimension worst-error analysis across the 20-country panel. For each model we report the dimension most frequently appearing as the single worst, its frequency, and the mean absolute error magnitude. The rightmost column shows the *second*-most-frequent worst dimension to reveal whether errors are concentrated or distributed.

Model	Worst dim	Freq.	Mean err (pp)	2nd worst dim	Freq.
Llama-3.3-70B	Util_More	15/20	49.7	SocVal_High	5/20
Phi-3.5-mini	Util_More	18/20	42.1	Species_Hum	2/20
Qwen3-VL-8B	SocVal_High	18/20	33.4	Util_More	2/20
Qwen2.5-7B	SocVal_High	12/20	24.5	Species_Hum	5/20
Magistral-24B	SocVal_High	14/20	22.4	Age_Young	3/20
Phi-4	SocVal_High	9/20	23.4	Age_Young	7/20

Table 27: Phi-4 per-country MIS (↓): the best-performing model in absolute terms. SWA-DPBR wins **18/20** countries; the top 10 gains are shown, ordered by relative improvement. The two losses (BRA, IRN) illustrate the diminishing-returns pattern discussed in §7.

Country	Vanilla	DPBR	Gain (%)	Country	Vanilla	DPBR	Gain (%)
USA	0.498	**0.243**	+51.1	JPN	0.395	**0.195**	+50.6
CHN	0.407	**0.214**	+47.4	IDN	0.459	**0.274**	+40.3
GBR	0.503	**0.319**	+36.6	THA	0.456	**0.313**	+31.4
MYS	0.459	**0.328**	+28.5	MMR	0.482	**0.357**	+26.0
BGD	0.489	**0.368**	+24.7	ROU	0.504	**0.385**	+23.5
Losses: BRA $0.274 \rightarrow 0.299$ (-9%)				IRN $0.397 \rightarrow 0.530$ (-34%)			
Macro mean: $0.454 \rightarrow$ **0.346**				+23.7%, 18/20 wins, mean $r = +0.56$			

Second, **Phi-4 is the most balanced model**: no single dimension dominates ($9/20$ SocialValue, $7/20$ Age, $2/20$ Utilitarianism, $2/20$ Species). This balance—not raw scale—explains why it achieves the lowest absolute MIS despite being $5\times$ smaller than Llama-3.3-70B.

Third, the per-country Phi-4 table (Table 27) confirms that **gains are large and geographically broad**: 51% on USA and Japan, 47% on China, 40% on Indonesia—with the two losses (Brazil, Iran) tracing to already-low vanilla MIS where IS overshoots.

AC Broader Model Landscape

The six models in Table 2 were selected from a broader sweep of **28 model–method combinations** covering 12 distinct architectures. Table 28 summarises the full landscape, ordered by macro improvement versus vanilla.

Table 28: Broader model sweep (5-country prototyping panel: BRA, CHN, DEU, JPN, USA). Models above the line show positive macro gains; those below show degradation. Bolded rows appear in the main paper's 20-country evaluation.

Model	Params	Mean MIS ↓	vs. van. (%)	Win/5
Phi-4	14B	**0.246**	+34.4	4/5
Llama-3.3-70B (4-bit)	70B	0.524	+25.9	5/5
Qwen3-VL-8B	8B	0.381	+23.7	4/5
Llama-3.1-8B (4-bit)	8B	0.455	+16.9	5/5
Magistral-Small-2509	24B	0.317	+13.6	4/5
Gemma-4-E2B	2B	0.426	+11.6	4/5
Qwen2.5-7B (bf16)	7B	0.385	+8.2	3/5
Phi-3.5-mini	3.8B	0.545	+7.3	4/5
Mistral-7B-v0.3	7B	0.428	+5.9	5/5
Qwen3.5-0.8B	0.8B	0.471	+3.2	4/5
Gemma-7B	7B	0.432	+2.5	3/5
Gemma-3-270M	270M	0.462	+1.3	4/5
Llama-3.2-1B	1B	0.476	+1.1	1/5
9 additional models showed ≤ 0 macro gain (omitted; see §7)				

Several patterns inform the main paper's model selection. *Success correlates with logit well-behavedness, not raw scale:* Phi-4 (14B) leads while Gemma-4-31B (31B) and Qwen3-Coder-30B (30B) degrade—confirming that SWA-DPBR requires a cooperative decision surface, not just more

parameters. *Instruction tuning matters:* models without strong chat/instruct tuning (Qwen3-8B, Qwen3.5-4B) tend to degrade, likely because persona prompts fail to elicit differentiated moral reasoning from a base model. *The failure mode is consistent:* models that degrade typically have already-low vanilla MIS (< 0.40), leaving insufficient headroom for IS corrections—matching the "diminishing returns" pattern discussed in §7.

AD Engineering Journey: Method Selection

SWA-DPBR (the "EXP-24" configuration) was selected from **25 method variants** developed and evaluated on a fixed three-model, five-country prototyping panel (Qwen2.5-7B, Gemma-2-9B, Mistral-7B $\times$ USA, CHN, JPN, DEU, BRA). Table 29 summarises the top-performing variants.

Table 29: Method variant comparison on the 3-model $\times$ 5-country prototyping panel (15 rows). Mean MIS $\downarrow$. All variants share the same persona construction and debiasing; they differ in the IS/prior stage.

Rank	Method variant	Mean MIS $\downarrow$	Key innovation
1	EXP-23: Category-Coherence Reg.	0.3956	Blend IS toward per-cat mean
2	**EXP-24: Dual-Pass Bootstrap (DPBR)**	**0.3969**	Soft $r = \exp(-v_b/s)$ reliability
3	EXP-09: Hierarchical IS	0.3975	Country EMA prior
4	EXP-22: Adaptive IS σ	0.3977	σ from agent/history stats
5	EXP-16: Nesterov IS	0.3977	Momentum + progressive PT
6	EXP-17: Dual-Momentum Prior	0.3979	Fast/slow β EMA blend
7	EXP-10: Grand Fusion	0.3982	EXP-03+05+09 combined
18	EXP-01: Base SWA-PTIS	0.4269	Single-pass IS (no prior)
25	EXP-12: Contrastive Personas	0.4517	World-average subtraction

The top cluster (ranks 1–7) is separated by only 0.0026 MIS, but DPBR was selected over the marginally-better EXP-23 for two reasons: (i) DPBR materially **lowers flip rate** (the fraction of scenarios where IS reverses the consensus decision) compared to all other top-cluster methods, improving deployment stability; and (ii) the dual-pass bootstrap provides an **interpretable uncertainty signal** (r in Eq. 6) that can be logged and audited per scenario. The base SWA-PTIS (rank 18) shows the value added by the hierarchical prior and dual-pass stages: a 7% relative MIS improvement. Contrastive persona decoding (rank 25) *hurts* alignment, validating the design choice discussed in §2 to use direct WVS personas rather than contrastive amplification.

AE Pilot: Open-Ended Cultural QA (BLEnD)

The binary trolley format that SWA-DPBR targets is structurally convenient—logit differences between two decision tokens provide a clean scalar signal for importance sampling—but it does not reflect the open-ended generation settings where LLMs are most commonly deployed. To probe whether the method's cultural steering transfers beyond forced binary choices, we ran a pilot experiment on the BLEnD benchmark [Myung et al., 2024], a dataset of 52.6k short-answer questions about everyday cultural knowledge across 16 countries/regions in 13 languages.

We evaluated Phi-4 (14B) in two conditions: vanilla greedy decoding (baseline) and SWA-DPBR with the same hyperparameters as the MultiTP experiments. Because BLEnD questions are open-ended, we adapted the pipeline to operate on next-token logits during generation rather than on a single binary decision token: persona-conditioned forward passes produce per-token logit corrections that are aggregated through the same PT–IS and dual-pass reliability stages. Answers were scored with Soft Exact Match at both the between-country (SEM-B) and within-country (SEM-W) levels, following the BLEnD evaluation protocol.

Table 30 presents the per-country results. The pattern is striking and instructive: SWA-DPBR improves countries where the baseline is already strong—the United States (+4.55 SEM-B), the United Kingdom (+7.45), and to a lesser extent Indonesia, Mexico, and Spain—but degrades countries where vanilla performance is low, with the sharpest drops in Assam (-13.21), Greece (-12.39), and Iran (-10.30). Overall, SWA-DPBR reduces aggregate SEM-B by 4.40 percentage points (from 48.91% to 44.51%).

Table 30: BLEnD pilot results (Phi-4 14B). SEM-B/W = Soft Exact Match between/within country (%). Δ = change from vanilla baseline. Countries sorted by baseline SEM-B (descending). SWA-DPBR improves high-resource countries but degrades low-resource ones, yielding a net negative overall effect.

Country	Baseline (Vanilla)		Δ SWA-DPBR	
	SEM-B	SEM-W	SEM-B	SEM-W
US	82.43	74.32	+4.55	+3.80
UK	75.66	66.69	+7.45	+8.78
Indonesia	73.75	60.92	+1.67	+3.39
Spain	72.71	63.12	+1.92	+2.89
Mexico	72.65	61.75	+2.09	+2.76
China	68.35	59.66	−1.68	−1.52
South Korea	56.84	48.01	−2.95	−2.63
Algeria	54.92	47.78	−8.70	−8.70
Iran	45.83	39.15	−10.30	−10.77
Greece	41.30	33.30	−12.39	−9.92
North Korea	38.51	32.75	−0.22	−0.36
Azerbaijan	28.78	24.92	+0.43	+0.48
West Java	26.64	20.73	+2.40	+3.81
Northern Nigeria	20.44	16.85	−7.91	−6.22
Assam	18.45	13.14	−13.21	−9.44
Ethiopia	5.30	4.56	−4.45	−3.82
Overall (16)	**48.91**	**41.73**	**−4.40**	**−3.28**

The diagnosis is mechanistic rather than cultural. SWA-DPBR's logit corrections are calibrated for a binary decision surface where two tokens compete; in open-ended generation, the correction is distributed across a much larger vocabulary, and the dual-pass reliability gate—designed to detect agreement between two IS estimates on a scalar—loses discriminative power in the high-dimensional softmax space. For high-resource countries where the baseline already generates fluent, culturally appropriate text, the persona signal provides a mild beneficial nudge. For low-resource countries, the baseline struggles with basic factual recall (SEM-B below 30%), and the IS correction introduces noise that further degrades generation quality—precisely the regime where the method's binary-format assumptions bind hardest.

Two additional factors compound the degradation. First, WVS persona coverage is sparsest for the lowest-performing countries (Assam, Northern Nigeria, Ethiopia), so the cultural signal itself is weaker and noisier. Second, BLEnD questions test factual everyday knowledge ("what food do people eat for birthdays?"), not moral preferences—the disagreement-as-signal intuition that powers SWA-DPBR on trolley dilemmas does not straightforwardly apply when the target is a factual answer rather than a preference distribution.

These results empirically confirm the binary-format limitation flagged in §7 and point toward a concrete extension path: adapting the IS stage to operate on softmax distributions rather than scalar logit gaps, with a vocabulary-aware reliability gate that can distinguish meaningful cultural steering from noise in the high-dimensional token space.

AF Extended Limitations

This appendix lists only *unresolved* limitations of the released system. Caveats that earlier drafts flagged but that the current codebase resolves (model quantisation artefacts now avoided by the vLLM BF16 backend on < 70B models; WVS-to-trolley linkage now placed on a quantitative footing by the regression and dropout matrix in Appendix P) are not repeated here.

Multi-seed CIs limited to one backbone. Cross-model headline gains in Table 2 use one seed per model. The dual-pass design provides an internal per-scenario reliability signal (Eq. 6) but does not substitute for multi-seed macro CIs. Appendix N reports a three-seed re-run on Phi-4 only (SWA-DPBR 0.4226 ± 0.0018 vs. vanilla 0.4469 ± 0.0007, so the Phi-4 gain is over $10\times$ the seed CI). Extending this to the remaining five backbones requires six additional 70B-scale runs we have not budgeted for.

Pearson r on a 6-dim vector. The AMCE vector is short, so per-country r has large sampling error and small row-wise r deltas are not always statistically separable; primary claims use MIS, with JSD as a secondary diagnostic and rank-based agreement (Appendix R) as a third.

Per-scenario persona-reward direction. The reward $r_i = \delta_i - \delta_{\text{base}}$ assumes persona shifts approximate the human target direction. The country-level WVS-to-AMCE regression in Appendix P.1 ($R^2 \in [0.55, 0.69]$) confirms this holds *on average across the 20-country panel*; what remains open is a *per-scenario* validation that persona-consensus shifts point toward the human AMCE for each specific dilemma, not just on the panel mean.

Geographic and linguistic coverage. We report 20 countries out of 100+ available in MultiTP, in 12 of the 107 MultiTP languages. Persona quality is weakest where WVS-7 coverage is sparse (e.g., Saudi Arabia uses an area-studies fallback, Appendix D.2). Low-resource languages with no active WVS-7 sample (Amharic, Khmer, Yoruba) are not represented in the panel; cross-lingual generalisation to such settings remains untested. WVS Wave 7 ($\sim$2017–2022) is also a fixed historical slice that may not reflect current preferences in rapidly evolving societies; refresh is straightforward when Wave 8 microdata becomes available.

Binary format. The scalar logit-difference formulation requires generalisation to softmax for multi-option / open-ended settings. The BLEnD pilot (Appendix AE) confirms this empirically: aggregate SEM-B drops by 4.4 pp, with gains confined to high-resource countries where vanilla generation is already strong. Softmax-aware IS and a vocabulary-level reliability gate are flagged as the concrete extension path.

Language–culture conflation. We use one language per country, so linguistic effects are entangled with cultural ones; isolating the two would require running the same country in multiple languages (e.g., Mandarin vs. English for Chinese respondents), which is not part of the MultiTP setup we adopt.