
Training-Free Cultural Alignment of Large Language Models via Persona Disagreement

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 LLMs are increasingly deployed in decisions that require culture-dependent moral
2 judgements, yet they answer as if the whole world thinks with a Western mind-
3 set. The Moral Machine experiment [Awad et al., 2018] showed this is wrong at
4 scale: 40 million judgments across 233 countries reveal that moral preferences are
5 systematically structured by culture, and a model that ignores this variation does
6 not merely underperform, but also imposes one society’s intuitions on all others.
7 Existing fixes do not scale to global deployment, as fine-tuning needs per-country
8 preference data and GPU budgets, reward-guided decoding needs per-country re-
9 ward models, and activation steering needs access to model internals that black-
10 box APIs do not expose. In this work, we focus on this realistic inference time
11 regime, with no weight updates, no training data, and no internal access. The key
12 observation is that within-country demographic *disagreement*, not consensus, is
13 the steering signal. When culturally grounded personas agree, the base model is
14 already calibrated. But when they disagree, the spread tells us what to fix and
15 how. We propose DISCA, which instantiates each country as a panel of four
16 World-Values-Survey-grounded persona agents [Haerpfer et al., 2020], converts
17 their disagreement into a bounded, loss-averse correction, whose magnitude is set
18 by the panel’s variance, and shrinks the correction toward zero when the estimate
19 is unreliable. Across 20 countries and 7 open-weight backbones from five model
20 families, DISCA reduces cultural misalignment on MultiTP [Jin et al., 2025] by
21 10–24% on the six $\geq 3.8\text{B}$ backbones, 3.4% on the smallest 2B model on binary
22 moral dilemmas, and 2–7% on open-ended scenarios. Furthermore, a 14B back-
23 bone with DISCA reaches lower absolute misalignment than a vanilla 70B model.

24 1 Introduction

25 Large language models increasingly mediate decisions that turn on moral judgement: content mod-
26 eration policies, clinical prioritisation guidelines, autonomous-vehicle behaviour specifications, and
27 high-stakes recommendation pipelines. In each of these settings, the model’s implicit moral pref-
28 erences become the system’s moral preferences. Yet a growing body of evidence shows that these
29 implicit preferences are not culturally neutral. The Moral Machine experiment collected 40 million
30 judgments across 233 countries and showed that moral preferences are systematically structured by
31 culture [Awad et al., 2018]; large language model (LLM) outputs correlate most strongly with re-
32 spondents from Western, educated, industrialised, rich, and democratic populations [Henrich et al.,
33 2010, Santurkar et al., 2023]; and post-training alignment introduces further unintended cultural
34 shifts [Ryan et al., 2024, Zewail et al., 2026]. A model that ignores this variation does not merely
35 underperform on a benchmark, but also imposes one society’s intuitions on every user it serves.

36 The cultural gap is now well-documented at the model level. The MultiTP benchmark [Jin et al.,
37 2025] extends the Moral Machine protocol to 107 languages with country-specific Average Marginal
38 Component Effects (AMCEs), and reports consistently weak country-level alignment across all 19
39 evaluated LLMs. The asymmetries are substantial, e.g., Japan ranks first among countries on pref-

erence for sparing pedestrians, while China ranks 116th, and current models do not capture either extreme. Concurrent work shows that sociodemographic persona prompts shift LLM moral decisions two to four times more than they shift human decisions [Kim et al., 2025], indicating that the models are not insensitive to cultural framing, but that the framing they receive at deployment time does not match the cultural distribution they are asked to serve.

Closing this gap is harder than it looks. Three families of methods have been proposed, and each demands resources that do not scale to the long tail of countries served by a single deployed model. Fine-tuning and culture-aware adapters require per-country preference datasets and GPU budgets for every target population [Zhang et al., 2026, Yao et al., 2025]. Reward-guided decoding requires a separate trained reward model per country [Khanov et al., 2024, Mudgal et al., 2024], and activation steering requires write-access to model internals [Arditi et al., 2025, Turner et al., 2024] that black-box APIs do not expose. Any deployable solution must therefore work at inference time, must not assume per-country reward models or labelled preference data, and must not assume write access to model internals beyond what an API exposes. We refer to this as the *black-box, public-data-only regime*: the method we propose requires only API-level access to decision-token log-probabilities (exposed by every open-weight backbone we evaluate and by major commercial APIs that return logprobs), and we argue this is the *only regime where cultural alignment can scale from a handful of curated countries to the hundreds a globally-deployed model actually serves*.

We observe that when certain culturally grounded persona prompts representing the same country face the same moral dilemma, the dispersion of their decoding-time logit gaps is itself informative. If they agree, the base model is already calibrated for that country and no correction is needed. But if they disagree, the spread of their gaps encodes both the direction and the magnitude of the needed correction. Our method, DISCA (Disagreement-Informed Steering for Cultural Alignment), instantiates each country as a panel of four persona agents grounded in World Values Survey microdata [Haerpfer et al., 2020], evaluates them in a single batched forward pass, and converts the panel’s variance into a bounded, loss-averse logit correction whose magnitude is automatically attenuated when the underlying estimate is unreliable. We show that this attenuation is principled: the underlying oracle correction is the MSE-optimal scalar shrinkage of the consensus, whose weight depends only on the within-panel variance (Proposition 1), and we approximate it with an empirical variance-aware shrinkage heuristic.

Across 20 countries spanning four continents and seven open-weight backbones (2B–70B parameters) from five model families, DISCA reduces cultural misalignment on MultiTP by 10–24% on the six ≥ 3.8 B backbones (and 3.4% on the smallest 2B model) on binary moral dilemmas, and 2–7% on open-ended scenarios. The strongest absolute alignment is achieved by Phi-4 (14B), which surpasses a vanilla Llama-3.3-70B backbone despite using one fifth of the parameters; this suggests that, in the regime we study, calibration competes with scale rather than merely complementing it. An evaluation on the BLEND factual cultural QA benchmark [Myung et al., 2024] reveals a clean scope boundary: DISCA steers values, where the decision reduces to a scalar logit gap, but does not transfer to factual recall, where the decision is a single token in a large vocabulary. Failure modes are architecturally rather than culturally determined and are diagnosable from the base model’s vanilla logit conditioning alone. Primary metrics (MIS and diagnostics) are defined in Section 4. Our contributions are summarised as follows:

- An inference-time cultural alignment method that requires no weight updates, no per-country reward models, and no white-box access.
- A formal characterisation of disagreement-driven shrinkage. DISCA is the first inference-time alignment method to use independent-run disagreement as a reliability signal.
- An empirical evaluation across 20 countries, four continents, seven open-weight backbones, and three evaluation formats (binary moral dilemmas, open-ended scenarios, factual QA), demonstrating that a 14B model with the proposed recipe surpasses a 70B baseline.

2 Related Work

Benchmarks and cultural evaluation. The Moral Machine experiment [Awad et al., 2018] established that moral preferences are culturally structured, and MultiTP [Jin et al., 2025] extends it to 107 languages with country-specific AMCEs; LLM alignment to these AMCEs varies widely across model families [Takemoto, 2024, Ahmad and Takemoto, 2025] and correlates most with WEIRD respondents [Santurkar et al., 2023, Henrich et al., 2010], with RLHF adding unintended cultural shifts [Ryan et al., 2024, Zewail et al., 2026]. Kim et al. [2025] further show that sociodemographic

96 personas shift LLM moral decisions 2–4 $\times$ more than they shift humans (a “partisan sorting” phe-
 97 nomenon absent in human respondents). This last finding directly motivates our design: rather than
 98 using personas to steer toward a single viewpoint, DISCA uses them to *measure* within-country
 99 disagreement and converts that into a bounded correction.

100 **Inference-time alignment.** The pluralistic alignment vision [Sorensen et al., 2024] argues that
 101 inference-time methods are a natural vehicle for cultural adaptation, but existing approaches each
 102 sacrifice one of the three constraints we target: activation steering [Arditi et al., 2025, Turner et al.,
 103 2024] requires white-box access; reward-guided decoding [Khanov et al., 2024, Mudgal et al., 2024]
 104 needs per-country reward models; culture-aware adapters [Zhang et al., 2026, Yao et al., 2025] need
 105 per-culture training. Training-time prospect-theoretic alignment [Ethayarajh et al., 2024] uses the
 106 same PT value shape we adopt, but at training time; DISCA applies it at *test time*, per scenario,
 107 aggregating heterogeneous personas rather than shaping a trained policy.

108 **Distinction from superficially similar paradigms.** DISCA may resemble ensemble methods ar-
 109 chitecturally, but not algorithmically. In self-consistency [Wang et al., 2023] and LLM debate [Du
 110 et al., 2024, Liang et al., 2024], diversity is noise to be eliminated via majority vote; standard en-
 111 sembles average outputs to reduce *prediction* variance. In DISCA, diversity *is* the signal: the *spread*
 112 of persona logit gaps, not the mode, sets the correction magnitude. We show (Proposition 1) that the
 113 MSE-optimal scalar shrinkage of the consensus correction is a closed-form function of within-panel
 114 variance, making within-panel variance a *sufficient statistic for correction reliability*. Standard cali-
 115 bration (temperature scaling, MC-Dropout [Gal and Ghahramani, 2016, Kwon et al., 2026]) adjusts
 116 confidence isotropically; DISCA applies anisotropic, loss-averse corrections. DISCA is the first
 117 training-free method targeting country-level moral calibration; the method is grounded in control-
 118 as-inference, Prospect Theory, and importance-sampling variance control.

119 3 Disagreement-Informed Steering for Cultural Alignment

120 DISCA aligns the model by creating an adjustment derived from the agents’ consensus, and the
 121 agreement among the agents controls how much of it actually applies. The magnitude of the adjust-
 122 ment is not fixed: when the agents agree on what the preference, we apply the adjustment in full;
 123 when they disagree, we apply only a fraction of it, scaled by how much the agents disagree. This
 124 inverts the standard combination structure used by self-consistency [Wang et al., 2023], multi-agent
 125 debate [Du et al., 2024, Liang et al., 2024], and ensembles, where N candidate outputs are aggre-
 126 gated by majority or averaging and disagreement among them is treated as noise to be suppressed.
 127 DISCA instead treats disagreement as the quantity that controls how much adjustment to apply.

128 3.1 Setup

129 We consider a frozen language model f_θ on a forced-choice task: each input x asks the model to
 130 choose between two options via decision tokens A and B . Let $z(x) = [z_a, z_b]$ denote the model’s
 131 logits at the decision-token position, and let $\delta(x) = z_b - z_a$ be the resulting decision gap. For target
 132 country c , we assume access only to f_θ as black-box and a public characterisation of country c ’s
 133 population (in our case, World Values Survey microdata [Haerpfer et al., 2020]); we do not assume
 134 per-country preference labels, reward models, or write access to θ . Our goal is to compute, from f_θ
 135 alone, a per-scenario logit adjustment $\delta^*(x, c)$ such that the adjusted gap $\delta(x) + \delta^*(x, c)$ produces
 136 decisions better aligned with country c ’s preferences. Alignment is measured against a country-level
 137 reference vector; the benchmark and metric are described in Section 4. The remainder of Section 3
 138 focuses on how δ^* should depend on the disagreement, what formal property this dependence has,
 139 and how to estimate the correction.

140 3.2 The disagreement principle

141 Suppose we evaluate the same scenario under N culturally grounded prompts representing the same
 142 target country, each producing a logit gap δ_i . Each δ_i is a noisy estimate of what country c ’s
 143 population-level preference would imply for this scenario. If the gaps cluster tightly, the N esti-
 144 mates converge on the same answer and the consensus is a reliable signal. If they are dispersed, the
 145 consensus is averaging over disagreeing estimates and is unreliable as a steering target. The princi-
 146 ple is that within-agent dispersion, not consensus magnitude, controls how much of the consensus
 147 adjustment we should actually apply.

148 Let $\delta_1, \dots, \delta_N$ be the logit gaps produced by N agents on a single scenario, and let δ_h denote the
 149 unobserved logit gap consistent with country c ’s population-level preference on that scenario. We

SWA_DPBR/figures/overview_pipeline_169.pdf

Figure 1: **DISCA overview.** Stage 1 builds WVS-grounded persona prompts for a trolley scenario in country c ; Stage 2 runs a frozen large language model (LLM) on the base prompt and each persona, aggregates persona-level signals in logit space, and applies *Prospect-Theory importance sampling* (PT-IS; §3.3.2) together with a dual-pass reliability gate to obtain the final sparing probability. Pseudocode and the six MultiTP attribute-temperature pairs provided in Appendix A1.

150 model the agents as $\delta_i = \delta_h + \eta_i$, $\mathbb{E}[\eta_i] = 0$, $\text{Var}(\eta_i) = \tau^2$, with the η_i assumed independent across
 151 agents and identically distributed with common variance τ^2 . This is a standard hierarchical setup:
 152 δ_h is a fixed-but-unknown population parameter, and each agent is a noisy view of it. The i.i.d.
 153 structure is a modelling simplification, since all agents share base model f_θ and the scenario. Define
 154 the consensus $\bar{\delta} = \frac{1}{N} \sum_i \delta_i$, the within-agent variance $D^2 = \frac{1}{N-1} \sum_i (\delta_i - \bar{\delta})^2$, the consensus
 155 adjustment $\Delta = \bar{\delta} - \delta_{\text{base}}$, and the oracle adjustment $\Delta_h = \delta_h - \delta_{\text{base}}$, where δ_{base} is the logit gap
 156 produced by the unconditioned base prompt (without persona or country-related prompts) on the
 157 same scenario. We will apply a scaled version of Δ as our adjustment: $\hat{\Delta}(\gamma) = \gamma \cdot \Delta$ for some
 158 $\gamma \in [0, 1]$. We choose γ by minimising the mean squared error against the unobserved oracle:
 159 $\text{MSE}(\gamma) = \mathbb{E}[(\hat{\Delta}(\gamma) - \Delta_h)^2]$.

160 **Proposition 1** (Variance-aware shrinkage - Proof in App. A2). *Under the agent model above:*

- 161 (i) D^2 is unbiased for τ^2 , and the variance of the consensus adjustment is $\text{Var}(\Delta) = \tau^2/N$.
 162 (ii) Among shrinkage estimators of the form $\hat{\Delta}(\gamma) = \gamma \cdot \Delta$ with $\gamma \in [0, 1]$, the MSE-minimising
 163 one has $\gamma^* = \frac{\Delta_h^2}{\Delta_h^2 + \tau^2/N}$, which lies in $[0, 1]$ and is monotone-decreasing in τ^2 .

164 3.3 The DISCA estimator

165 Proposition 1 tells us what shape the shrinkage should have but not how to compute it: γ^* depends
 166 on the unobservable Δ_h . In particular, the proposition’s two parts work together to motivate the
 167 method. Part (ii) gives the optimal γ^* in closed form, but it is not directly computable: γ^* depends
 168 on τ^2 and on Δ_h , neither of which we observe. Part (i) gives us a way to estimate τ^2 from data
 169 we already have: the within-agent variance D^2 has expected value equal to τ^2 . As Δ_h has no
 170 analogous estimator from the agents’ outputs, we cannot compute γ^* exactly. Instead, we construct
 171 an estimator δ^* that is computable from the agents’ outputs alone, matches the shape of γ^* , bounded
 172 in $(0, 1)$, & monotone-decreasing in the variance estimate.

We use $N=4$ agents per target country: three age cohorts (young, middle-aged, older) plus a country-wide aggregate, each instantiated as a native-language system prompt built from cohort- and country-level aggregates of the WVS cultural profile of Greco et al. [2026] (full prompts in App. A11). The choice $N=4$ is the smallest panel that simultaneously covers the three demographic axes documented to carry most within-country variance [Inglehart and Welzel, 2005] and includes a country-wide aggregate that anchors the ensemble against cohort-specific sampling noise; it is also the empirical optimum of a sweep over $N \in \{2, 3, 4, 5, 6\}$ on a three-country panel (Tab. 21), where $N=4$ minimises macro MIS, $N=2$ leaves the consensus under-determined, and $N \geq 5$ duplicates demographic coverage without adding new variance. To remove the order bias of forced-choice decoding, we evaluate every scenario under both (A, B) and (B, A) orderings and take the order-symmetrised gap $\delta_i = \frac{1}{2}(\delta_i^{(AB)} - \delta_i^{(BA)})$ as the agent’s output.

3.3.1 Loss-averse importance sampling

The plain consensus adjustment $\Delta = \bar{\delta} - \delta_{\text{base}}$ moves the base model toward the average direction because $\bar{\delta}$ is the average consensus. The problem is that an average direction can hide a cohort that strongly disagrees: three cohorts mildly in favour and one cohort strongly against will average out to "mildly in favour," and the consensus will move the model in a direction that one cohort actively rejects. We want a steering rule that avoids this. One that prefers a direction nobody strongly dislikes over a direction with a higher average but a strong dissenter.

We frame this as a cooperative-bargaining problem among the cohorts [Nash, 1950, Kalai and Smorodinsky, 1975]. Each cohort is a player with its own utility over correction directions, each is loss-averse, which means a small misalignment with one’s preferred direction is felt more than a same-sized improvement, and we want to pick a direction that all players can collectively accept. This framing organises three design choices: where to look for candidates, how each player scores a candidate, and how all scores combine into a group verdict.

Rather than committing to Δ , we draw K candidate perturbations $\epsilon_k \sim \mathcal{N}(0, \sigma^2)$ in a neighbourhood of the consensus, defining the candidate state $\tilde{\delta}_k = \bar{\delta} + \epsilon_k$. Each candidate is scored by how much closer to cohort i ’s preferred direction it lands, in absolute ℓ_1 distance, against two reference points: the cohort gap δ_i and the panel consensus $\bar{\delta}$,

$$g_{i,k} = |\delta_{\text{base}} - \delta_i| - |\tilde{\delta}_k - \delta_i|, \quad g_{\text{cons},k} = |\delta_{\text{base}} - \bar{\delta}| - |\tilde{\delta}_k - \bar{\delta}|, \quad (1)$$

so positive g means the candidate moved closer to the relevant target than the base prompt did, negative means it moved the model away.

We need a scoring rule $v(\cdot)$ that is concave on the gain side so additional gain beyond what’s already aligned is worth less and steeper on the loss side so any cohort loss matters more than the same-size gain elsewhere. The canonical such rule in the decision-theory literature is the Kahneman–Tversky value function [Kahneman and Tversky, 1979, Tversky and Kahneman, 1992]:

$$v(z) = \begin{cases} z^\alpha & z \geq 0, \\ -\kappa(-z)^\alpha & z < 0, \end{cases} \quad (2)$$

with curvature $\alpha \in (0, 1]$ and loss-aversion $\kappa \geq 1$. The asymmetric kernel v acts as an aggregator that penalises misalignment more than rewarding incremental alignment.

Cooperative bargaining sits between two extremes: a purely *utilitarian* verdict averages the players’ utilities and ignores how the gains are distributed, while a purely *collective* verdict asks how the deal serves the group as a whole. We use a blend of the two:

$$U_{\text{total}}(\epsilon_k) = (1 - \lambda_{\text{coop}}) \frac{1}{N} \sum_{i=1}^N v\left(\frac{g_{i,k}}{\sigma}\right) + \lambda_{\text{coop}} v\left(\frac{g_{\text{cons},k}}{\sigma}\right), \quad (3)$$

where $g_{\text{cons},k}$ measures the deal’s gain against the consensus direction $\bar{\delta}$ rather than against any single cohort. The first term is the utilitarian verdict over loss-averse players; the second is the collective verdict scored by the same loss-averse rule. The coefficient $\lambda_{\text{coop}} \in [0, 1]$ controls the bargaining position, with $\lambda_{\text{coop}} = 0$ purely utilitarian and $\lambda_{\text{coop}} = 1$ purely collective. We use $\lambda_{\text{coop}} = 0.7$, on the collective side; a sensitivity sweep over $\lambda_{\text{coop}} \in [0, 1]$ is in App. A10. The division by σ inside v scores each gain in units of the perturbation scale, so the search radius σ and the value-function shape (α, κ) can be set independently.

219 The natural choice, returning the highest-utility candidate, is unsafe here. The utility $U_{\text{total}}(\epsilon_k)$ is
 220 itself a noisy quantity. It depends on the random draw ϵ_k and on the agent gains $g_{i,k}$, both of
 221 which carry stochastic error. Picking the argmax over K noisy scores selects the candidate that
 222 drew favourable noise on this sample, not the candidate with the highest true utility. We therefore
 223 aggregate all candidates with weights proportional to their utility:

$$\delta_{PT-IS} = \frac{\sum_{k=1}^K w_k \epsilon_k}{\sum_{k=1}^K w_k}, \quad w_k \propto \exp(U_{\text{total}}(\epsilon_k)/\eta), \quad (4)$$

224 where $\eta > 0$ is an IS softmax temperature ($\eta=0.5$ default; Table 18). This form is not ad hoc:
 225 it is the path-integral solution to the stochastic optimal control problem of selecting a correction
 226 direction under a noisy utility function [Williams et al., 2018, Levine, 2018]. The high-utility can-
 227 didate dominates the average, but near-winners contribute as well, and stochastic error in any one
 228 candidate’s utility is suppressed by the contributions of the others. We refer to the full procedure as
 229 *loss-averse importance sampling*, or PT-IS for short.

230 3.3.2 Dual-pass reliability gate

231 PT-IS gives us a single correction δ_{PT-IS} for each scenario. The natural next step is to apply it: add
 232 δ_{PT-IS} to the model’s base logits and read off the steered decision. The problem is that δ_{PT-IS} is a single
 233 number combining cohort opinions. When all cohorts agree, that number reflects a unanimously
 234 backed direction. When the cohorts disagree, the same number can come out small even though no
 235 cohort actually wants the small push it represents.

236 Concretely, imagine two cohorts prefer direction $+0.5$ and two prefer -0.5 . PT-IS with loss-
 237 aversion will not pick either side; it returns a compromise, say $\delta_{PT-IS} = +0.05$, close to zero because
 238 pushing in either direction would hurt half the cohorts. Applying $+0.05$ to the model’s logits nudges
 239 the model in a direction no cohort actually asked for. This is the failure mode Proposition 1 warned
 240 us about: γ^* contracts toward zero when cohort disagreement is high, because the appropriate re-
 241 sponse to disagreement is to apply less correction, not a more carefully averaged correction. We
 242 need a way to identify contested scenarios at run time and contract δ_{PT-IS} accordingly.

243 The obvious approach is to increase K . With infinitely many candidates, δ_{PT-IS} converges to a
 244 deterministic limit and run-to-run variation vanishes. But this does not solve the problem: the $+0.05$
 245 compromise becomes a more precise compromise. More samples sharpen the average; they do not
 246 change whether the average deserves to be applied at full strength.

247 We run PT-IS twice on the same scenario, on disjoint sub-budgets of size $K/2$, producing indepen-
 248 dent corrections $\delta_{PT-IS}^{(1)}$ and $\delta_{PT-IS}^{(2)}$. When the cohorts agree, both runs find the same compromise and
 249 the two corrections are close. When the cohorts disagree, the runs are pulled by competing cohort
 250 camps and land on different compromises. The squared inter-run gap

$$V_r = (\delta_{PT-IS}^{(1)} - \delta_{PT-IS}^{(2)})^2$$

251 is the diagnostic: zero when the runs agree, growing as they diverge. V_r is the half-sample vari-
 252 ance estimator: a single sample is split in two, and variance is estimated from the squared inter-half
 253 gap. The technique originates in Mahalanobis [?] and was generalised as balanced repeated replica-
 254 tion [?]; it is canonical in settings that need a qualitative variance signal on a constrained compute
 255 budget. Two runs also keep total compute matched to a single full PT-IS run.

256 We multiplicatively shrink the correction by

$$r = \exp(-V_r/s) \in (0, 1], \quad (5)$$

257 with bandwidth $s > 0$ controlling the decay rate. The exponential is the simplest function satisfy-
 258 ing the three shape properties Proposition 1 requires: bounded in $(0, 1]$, monotone-decreasing, and
 259 smooth. The final DISCA correction is the gated average of the two runs:

$$\delta^* = r \cdot \frac{\delta_{PT-IS}^{(1)} + \delta_{PT-IS}^{(2)}}{2}. \quad (6)$$

260 δ^* matches the qualitative shape of $\gamma^* \cdot \Delta$: bounded by the average run magnitude, monotone-
 261 decreasing in inter-run disagreement (mirroring γ^* in τ^2), and smooth in all inputs. The empirical
 262 contribution of the gate is bundled inside the full-method ablation in Table 4.

3.4 Extension to Open-Ended Ethical Scenarios

To extend DISCA to open-ended scenarios, the model generates free-form responses under both the original prompt and four persona-specific prompts. A separate large language model (LLM) judge then parses these outputs to return a {choice, confidence} pair. The continuous decision signal is represented by a pseudo-logit gap (the A/B logit distance) extracted by that judge. The dual-pass Prospect-Theory importance sampling (PT-IS) mechanism is applied similarly to the binary case.

4 Experimental Results

We test DISCA on MultiTP [Jin et al., 2025], the multilingual extension of the Moral Machine trolley problem [Awad et al., 2018]. Each scenario shows the model two groups of people that an autonomous vehicle could spare, and the model has to choose one of them by emitting a single decision token, A or B . Every scenario isolates one of six moral attributes (species, gender, age, fitness, social value, and the number of lives saved) while matching the two groups on everything else. An age scenario, for example, pits a young group against an older one and holds the rest of the description constant; the rate at which the model spares the younger group, averaged across many such scenarios, is then read off as its preference on age. Following the Moral Machine convention we always assign decision token B to the side the AMCE measures toward; this is a measurement convention from conjoint analysis that fixes the sign of the preference, not a normative claim about which side ought to be saved. The frozen model emits logits $z(x) = [z_a, z_b]$ at the decision position, the decision gap is $\delta(x) = z_b - z_a$, and the sparing probability under decoding temperature T_{dec} is $p_{\text{spare}}(x) = \sigma(\delta(x)/T_{\text{dec}})$. Averaging the sparing probability across the scenarios for one attribute and one country gives that country’s Average Marginal Component Effect (AMCE):

$$\hat{m}_d^{(c)} = \frac{1}{|\mathcal{S}_d^{(c)}|} \sum_{x \in \mathcal{S}_d^{(c)}} p_{\text{spare}}(x). \quad (7)$$

Stacking the six attributes gives a six-dimensional vector $\hat{m}^{(c)}$ that we compare to the human AMCE vector $h^{(c)}$ aggregated from real Moral Machine judgments [Awad et al., 2018, Jin et al., 2025]. Our headline metric is the *misalignment score* $\text{MIS}(c) = \|\hat{m}^{(c)} - h^{(c)}\|_2$, which is just the straight-line distance between the model’s preference vector and the country’s human one; lower is better. We report it macro-averaged over a 20-country panel that spans four continents. The country list, the pre-processing pipeline (deduplication, quality filter, per-attribute caps, oversampling, and deterministic shuffling), the per-country slice sizes, and a cap-disabled sensitivity rerun all live in Appendix A13. Every scenario is presented in the country’s native language; only the decision tokens A and B stay in English so that token-ID extraction is consistent across model families. We also report Pearson r for shape agreement and Jensen–Shannon divergence (JSD) as secondary diagnostics.

We tested 28 model–method combinations across 12 architectures from 270M to 70B parameters, with the full landscape in Appendix A6, and report seven open-weight backbones with the most robust gains: Llama-3.3-70B, Magistral-24B, Phi-4 (14B), Qwen3-VL-8B, Qwen2.5-7B, Phi-3.5-mini (3.8B), and Gemma-4-E2B (2B). The few models that degrade share a single pattern: their vanilla MIS is already low enough that there is no headroom left for any correction (we return to this in §4.2). The released hyperparameter configuration was validated on a small three-backbone, five-country prototyping panel before being frozen for the headline sweep; how we elicit the A and B tokens consistently across model families is described in Appendix A12.

We compare DISCA against two tiers of baselines on the same 20-country Phi-4 grid (Table 1; full implementations in Appendix A9). The first tier is training-free and never sees the human AMCE: vanilla decoding, a WVS Profile Prompt that summarises the country’s WVS values, a PRISM-style cultural framing prompt [Kirk et al., 2024], a fixed logit offset, activation steering [Arditi et al., 2025, Zou et al., 2023], and MC-Dropout [Gal and Ghahramani, 2016, Kwon et al., 2026]. We exclude reward-guided decoding [Khanov et al., 2024, Mudgal et al., 2024] because it assumes per-country reward models, which simply do not exist for the long tail of countries we care about. The second tier is the opposite: oracle baselines that fit one per-country scalar (temperature or additive margin) directly on the country’s full human AMCE [Guo et al., 2017], serving as upper bounds on what any 1-d-per-country oracle correction can do.

4.1 Main results

A single decoding-time intervention, with no weight changes and no per-country labels, materially closes the gap between large language models and human moral preferences. Across seven back-

Table 1: Inference-time baseline comparison on the Phi-4 20-country grid (macro MIS, lower is better). Oracle baselines (below dashed line) use human AMCE during calibration; even so they fail to match DISCA.

Method	AMCE?	MIS ↓
DISCA (ours)	No	0.346
PRISM-Style Prompt [Kirk et al., 2024]	No	0.384
MC-Dropout [Gal and Ghahramani, 2016]	No	0.403
Activation steering [Arditi et al., 2025]	No	0.430
Fixed logit offset	No	0.439
WVS Profile Prompt	No	0.453
Vanilla decoding	No	0.454
Margin scaling [Guo et al., 2017]	Yes	0.506
Temp. scaling [Guo et al., 2017]	Yes	0.513

Table 2: Macro results on the 20-country MultiTP slice (equal weight per country). MIS is mean ℓ_2 AMCE distance ($\pm$ std over three seeds {42, 101, 2026} chosen under a fixed compute budget; per-cell std median 0.008, full multi-seed analysis in §A8.1); “Win” = countries with lower DISCA MIS than vanilla.

Model	MIS ↓	Gain (%)	Win/20	r
Llama-3.3-70B	.668 $\pm$.006	+21.3 $\pm$ 0.8	20	−0.01
Magistral-24B	.350 $\pm$.005	+13.0 $\pm$ 0.7	16	+0.62
Phi-4 (14B)	.346 $\pm$.004	+23.6 $\pm$ 0.6	18	+0.56
Qwen3-VL-8B	.466 $\pm$.006	+18.8 $\pm$ 0.9	20	+0.22
Qwen2.5-7B	.362 $\pm$.005	+20.0 $\pm$ 0.8	17	+0.45
Phi-3.5-mini	.571 $\pm$.006	+10.3 $\pm$ 1.0	17	−0.35
Gemma-4-E2B (2B)	.479	+3.4	14	+0.46

Table 3: Open-ended track summary (20 countries, 310 scenarios each). **VAN** = vanilla decoding (no DISCA). **DISCA** = full pipeline. **MIS** = misalignment score (§4). $\Delta\%$ reports mean relative MIS reduction across four backbones; per-country breakdown in Appendix A3.1.

Model	Vanilla	DISCA	$\Delta\%$
Llama-3.3-70B	0.471	0.439	+6.85%
Phi-4 (14B)	0.324	0.302	+6.67%
Qwen2.5-7B	0.321	0.306	+4.55%
Phi-3.5-mini	0.521	0.510	+2.13%

bones from five model families, DISCA reduces macro MIS by roughly ten to twenty-four percent relative to vanilla decoding (Table 2), and two of the seven backbones (Llama-3.3-70B and Qwen3-VL-8B) improve on every one of the twenty countries we test. A 2D PCA of the AMCE vectors tells the same story visually (Figure 3, Appendix A5): every country point migrates from the vanilla cloud toward the human cloud, none in the wrong direction. No simpler intervention comes close. On the same Phi-4 grid (Table 1) DISCA outperforms every training-free baseline by a clear margin, and even the two oracle baselines that get to peek at each country’s full human AMCE during calibration end up worse than DISCA, because a single per-country scalar (whether a temperature or an additive margin) cannot capture all six AMCE dimensions simultaneously.

The per-country breakdown (Table 6, Appendix A3) confirms that the macro gains are not concentrated in a handful of lucky cells. Geographic patterns, large single-country swings, and the harder regions (such as those with sparser WVS coverage) are all detailed there. To rule out artefacts, we also ran nine independent robustness checks and three random seeds; every conclusion sits comfortably within the bootstrap noise floor (Appendices A10 and A8.1).

Calibration competes with scale The single most striking finding is a crossing in the curve of MIS against model size (Appendix A5): a 14B model with DISCA (Phi-4) reaches lower absolute misalignment than a 70B model without it (Llama-3.3-70B). Per-dimension error analysis (Table 9, Appendix A4) explains why. Different models hit a wall on different attributes. Weaker models are dominated by Utilitarianism errors that are simply too large for any persona-based correction to close. Better-calibrated models have already resolved Utilitarianism and Species, so their bottleneck shifts to Social Value and Age, two attributes that DISCA’s importance-sampling stage actually moves the needle on. Phi-4 turns out to be the most balanced of the seven across all dimensions, and this, rather than raw parameter count, brings its top performance despite being five times smaller than the baseline.

Generalisation beyond binary dilemmas The same disagreement signal also helps in open-ended ethical generation (§3.4), not just in forced binary choices. All four backbones we tested in this regime improve on average, with gains up to roughly seven percent (Table 3). A small number of country–model cells regress slightly (six of eighty cells in Table 7, all with $|\Delta| \leq 3.1\%$); the gate keeps these residual losses bounded but cannot reverse them. A logit-noise stress test (Appendix A3.2) shows that the full gated pipeline maintains a 10–12 percent MIS advantage over an ungated variant—the same PT-IS stack with the agreement-based shrinkage disabled—across noise levels $\sigma \in [0, 2]$, and the gap widens as the noise grows.

4.2 Ablation and Discussion

The results so far establish that DISCA works across model families, geographies, and evaluation formats. The harder question is why it works, where it stops working, and what it should not be expected to do. To find out which part of the pipeline carries the gain, we ablate four design choices across three 20-country backbones (Table 4).¹ A clean and consistent hierarchy emerges across all three. Persona diversity is the irreplaceable core: removing the WVS-grounded personas causes the largest degradation and hurts the majority of countries, which confirms that within-country disagreement, not consensus, is what carries most of the alignment value. The controlled correction through PT-IS contributes a stable second-order gain, and the adaptive reliability gate matters more than raw correction strength. Positional debiasing has the smallest macro effect of the four but remains necessary as a conditioning step, because without it prompt-order contamination corrupts the cultural signal upstream and the rest of the pipeline never sees clean inputs.

Table 4: Cross-backbone ablation on all 20 paper countries. “ n hurt” = number of countries with strictly higher MIS than Full DISCA.

Configuration	Qwen2.5-7B (BF16)			Phi-3.5-mini-Instruct			Magistral-Sml (24B)		
	MIS ↓	Δ vs Full	n hurt	MIS ↓	Δ vs Full	n hurt	MIS ↓	Δ vs Full	n hurt
Full DISCA	0.362	–	–	0.571	–	–	0.334	–	–
Without persona	0.417	+0.055	16	0.618	+0.047	17	0.395	+0.061	18
Always-on PT-IS	0.379	+0.017	14	0.586	+0.015	14	0.355	+0.021	15
No-IS (consensus)	0.371	+0.009	12	0.579	+0.008	13	0.345	+0.011	13
No debiasing	0.363	+0.001	12	0.573	+0.002	11	0.338	+0.004	10

A trustworthy inference-time controller has to know when not to act, and DISCA does. On the vast majority of scenarios the dual-pass reliability gate passes the correction through unattenuated, and on the small fraction where the two passes disagree it suppresses aggressively (the full audit is in Appendix A8). Replacing the bounded, loss-averse aggregation with simple consensus averaging barely changes the mean gain, but it nearly quadruples the number of harmed country–model cells and triples worst-case degradation (Table 14, Appendix A8.2). The bounded aggregation is therefore primarily a safety mechanism rather than an average-case booster; it is what lets us claim the gains without the long tail of regressions a naive average would introduce.

When DISCA does fail, the cause is always architectural rather than cultural. The 28-model sweep (Appendix A6) shows that success correlates with how well-behaved the decision logits are, not with how many parameters the model has, and per-scenario decision margin, entropy, and absolute logit gap (Appendix A8) flag likely-to-degrade configurations from the vanilla forward pass alone. Phi-4 leads the sweep at 14B while several 30B+ models degrade [Chand et al., 2026], which suggests a simple deployment-time check on logit conditioning is enough to flag likely failures before running the full pipeline. The price we pay for the gains is moderate, roughly four times the latency of vanilla decoding (Appendix A14).

The disagreement signal does have a clear scope boundary. On the BLEnD factual cultural QA benchmark [Myung et al., 2024] (Appendix A7) it does not help, and the reason is illuminating in itself. Value alignment in MultiTP operates through a single scalar gap between two logits, which is exactly the kind of low-dimensional decision surface our importance sampling navigates well. Factual QA, by contrast, asks the model to pick one token out of a roughly 32k-entry vocabulary, and perturbations calibrated for a binary gap turn into noise on a wide softmax. Drawing this line between value steering and fact retrieval is itself a contribution: it tells future work where persona-disagreement methods belong and where vocabulary-level mechanisms are needed instead.

Three caveats bound how far our method generalises. We optimise against crowdsourced AMCE vectors [Awad et al., 2018, Jin et al., 2025] without a separate human study of perceived cultural appropriateness [Atari et al., 2023, Khan et al., 2025], so our numbers measure alignment to a survey statistic, not proof of legitimacy. The Prospect-Theory value function plays the role of an asymmetric aggregation kernel inside the controller, not a literal model of how the LLM thinks. And the method needs decision-token logits, which black-box text-only APIs do not always expose (Appendix A16). Finally, an inference-time cultural controller can encode harmful majorities if

¹vLLM backend, one K -sample IS batch per scenario. **Full DISCA** retains the dual-pass reliability gate; the gate is disabled *only* for the rows that explicitly isolate components downstream of it (*Always-on PT-IS* and *No-IS (consensus)*), so those rows quantify what the gate hides without altering the headline configuration.

390 deployed naively [Zewail et al., 2026]; we therefore add a per-persona utility floor that caps how
 391 far any single persona’s post-correction utility is allowed to drop below vanilla, and sweeping the
 392 floor on a USA/VNM/DEU panel leaves macro MIS flat within the bootstrap noise floor (Table 24,
 393 Appendix A10), so the safeguard is essentially free.

394 5 Conclusion

395 In this paper we showed that LLMs need not be retrained to respect moral diversity. To do so,
 396 we have introduced *disagreement-driven steering*, a paradigm that treats within-group variance of
 397 grounded agents as the primary alignment signal rather than their consensus. DISCA instantiates
 398 this paradigm for cultural alignment, reducing misalignment by 10–24% on the six headline back-
 399 bones ($\geq 3.8B$) and a smaller 3.4% on the 2B Gemma-4-E2B on binary dilemmas, and 2–7% on
 400 open-ended scenarios, across twenty countries without touching a single weight. The formal result
 401 changes how we think about inference-time alignment: disagreement is not noise to be averaged
 402 away but an *optimal sufficient statistic* for correction reliability, and the gap between ensembles and
 403 shrinkage estimators explains why naive averaging fails where DISCA succeeds.

404 References

- 405 M. S. Z. b. Ahmad and K. Takemoto. Large-scale moral machine experiment on large language
 406 models. *PLOS ONE*, 20(5):e0322776, 2025. doi: 10.1371/journal.pone.0322776. URL <https://doi.org/10.1371/journal.pone.0322776>.
- 408 A. Arditì, O. Obeso, A. Syed, D. Paleka, N. Panickssery, W. Gurnee, and N. Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, volume 37, pages 136037–136083, 2025. doi: 10.52202/079017-4322. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/f545448535dfde4f9786555403ab7c49-Abstract-Conference.html.
- 413 M. Atari, M. J. Xue, P. S. Park, D. E. Blasi, and J. Henrich. Which humans? *PsyArXiv preprint*, 2023. doi: 10.31234/osf.io/5b26t. <https://osf.io/preprints/psyarxiv/5b26t>.
- 415 E. Awad, S. Dsouza, R. Kim, J. Schulz, J. Henrich, A. Shariff, J.-F. Bonnefon, and I. Rahwan. The moral machine experiment. *Nature*, 563(7729):59–64, 2018.
- 417 S. Chand, F. Baca, and E. Ferrara. No free lunch in language model bias mitigation? Targeted bias reduction can exacerbate unmitigated LLM biases. *AI*, 7(1):24, 2026. doi: 10.3390/ai7010024.
- 419 R. Chen, W. Chai, Z. Yang, X. Zhang, Z. Wang, T. Quek, J. T. Zhou, S. Poria, and Z. Liu. DiffPO: Diffusion-styled preference optimization for inference time alignment of large language models. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 18910–18925, Vienna, Austria, July 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.acl-long.926. URL <https://aclanthology.org/2025.acl-long.926/>. arXiv:2503.04240.
- 425 A. Deshpande, V. Murahari, T. Rajpurohit, A. Kalyan, and K. Narasimhan. Toxicity in ChatGPT: Analyzing persona-assigned language models. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 1236–1270, Singapore, Dec. 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-emnlp.88. URL <https://aclanthology.org/2023.findings-emnlp.88/>. arXiv:2304.05335.
- 430 Y. Du, S. Li, A. Torralba, J. B. Tenenbaum, and I. Mordatch. Improving factuality and reasoning in language models through multiagent debate. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 11733–11763. PMLR, 2024. URL <https://proceedings.mlr.press/v235/du24e.html>. arXiv:2305.14325.
- 435 B. Efron. Bootstrap methods: another look at the jackknife. *The Annals of Statistics*, 7(1):1–26, 1979. doi: 10.1007/978-1-4612-4380-9_41. URL https://link.springer.com/chapter/10.1007/978-1-4612-4380-9_41. Reprinted in *Breakthroughs in Statistics* (1992).
- 438 V. Elvira and L. Martino. Advances in importance sampling. *Wiley StatsRef: Statistics Reference Online*, pages 1–14, 2021. doi: 10.1002/9781118445112.stat08284. URL <https://arxiv.org/abs/2102.05407>.

- 441 K. Ethayarajh, W. Xu, N. Muennighoff, D. Jurafsky, and D. Kiela. KTO: Model alignment
442 as prospect theoretic optimization. In *International Conference on Machine Learning*, 2024.
443 arXiv:2402.01306.
- 444 Y. Gal and Z. Ghahramani. Dropout as a Bayesian approximation: Representing model uncertainty
445 in deep learning. In *Proceedings of the 33rd International Conference on Machine Learning*,
446 volume 48 of *Proceedings of Machine Learning Research*, pages 1050–1059. PMLR, 2016. URL
447 <https://proceedings.mlr.press/v48/gal16.html>.
- 448 C. M. Greco, L. La Cava, and A. Tagarelli. Culturally grounded personas in large language mod-
449 els: Characterization and alignment with socio-psychological value frameworks. *arXiv preprint*
450 *arXiv:2601.22396*, 2026.
- 451 C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger. On calibration of modern neural networks. In *Pro-*
452 *ceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings*
453 *of Machine Learning Research*, pages 1321–1330, 2017. URL [https://proceedings.mlr.](https://proceedings.mlr.press/v70/guo17a.html)
454 [press/v70/guo17a.html](https://proceedings.mlr.press/v70/guo17a.html). arXiv:1706.04599.
- 455 C. Haerpfer, R. Inglehart, A. Moreno, C. Welzel, K. Kizilova, J. Diez-Medrano, M. Lagos,
456 P. Norris, E. Ponarin, and B. Puranen. World Values Survey: Round seven – country-pooled
457 datafile. Madrid, Spain & Vienna, Austria: JD Systems Institute & WVSA Secretariat, 2020.
458 <https://doi.org/10.14281/18241.20>.
- 459 J. Henrich, S. J. Heine, and A. Norenzayan. The weirdest people in the world? *Behav-*
460 *ioral and Brain Sciences*, 33(2-3):61–83, 2010. doi: 10.1017/S0140525X0999152X. URL
461 [https://www.cambridge.org/core/journals/behavioral-and-brain-sciences/](https://www.cambridge.org/core/journals/behavioral-and-brain-sciences/article/abs/weirdest-people-inthe-world/BF84F7517D56AFF7B7EB58411A554C17)
462 [article/abs/weirdest-people-inthe-world/BF84F7517D56AFF7B7EB58411A554C17](https://www.cambridge.org/core/journals/behavioral-and-brain-sciences/article/abs/weirdest-people-inthe-world/BF84F7517D56AFF7B7EB58411A554C17).
- 463 R. Inglehart and C. Welzel. *Modernization, Cultural Change, and Democracy: The Human Devel-*
464 *opment Sequence*. Cambridge University Press, 2005. URL [https://social.hse.ru/data/](https://social.hse.ru/data/2012/11/03/1249193128/inglehart_welzel.pdf)
465 [2012/11/03/1249193128/inglehart_welzel.pdf](https://social.hse.ru/data/2012/11/03/1249193128/inglehart_welzel.pdf).
- 466 Z. Jin, M. Kleiman-Weiner, G. Piatti, S. Levine, J. Liu, F. G. Adauro, F. Ortu, A. Strausz, M. Sachan,
467 R. Mihalcea, Y. Choi, and B. Schölkopf. Language Model Alignment in Multilingual Trolley
468 Problems. In *International Conference on Learning Representations*, 2025. URL [https://](https://openreview.net/forum?id=VEqPDZIDAh)
469 openreview.net/forum?id=VEqPDZIDAh. arXiv:2407.02273.
- 470 D. Kahneman and A. Tversky. Prospect theory: An analysis of decision under risk. *Econometrica*,
471 47(2):263–291, 1979.
- 472 E. Kalai and M. Smorodinsky. Other solutions to Nash’s bargaining problem. *Econometrica*, 43(3):
473 513–518, 1975. doi: 10.2307/1914280.
- 474 A. Khan, S. Casper, and D. Hadfield-Menell. Randomness, not representation: The unreliabil-
475 ity of evaluating cultural alignment in LLMs. In *Proceedings of the 2025 ACM Conference on*
476 *Fairness, Accountability, and Transparency*, pages 2151–2165. Association for Computing Ma-
477 chinery, 2025. doi: 10.1145/3715275.3732147. URL [https://dl.acm.org/doi/10.1145/](https://dl.acm.org/doi/10.1145/3715275.3732147)
478 [3715275.3732147](https://dl.acm.org/doi/10.1145/3715275.3732147). arXiv:2503.08688.
- 479 M. Khanov, J. Burapachep, and Y. Li. ARGS: Alignment as reward-guided search. In *International*
480 *Conference on Learning Representations*, 2024. URL [https://openreview.net/forum?id=](https://openreview.net/forum?id=shgx0eqdw6)
481 [shgx0eqdw6](https://openreview.net/forum?id=shgx0eqdw6). arXiv:2402.01694.
- 482 J. Kim, J. Kwon, L. F. Vecchietti, A. Oh, and M. Cha. Exploring persona-dependent llm alignment
483 for the moral machine experiment. *arXiv preprint arXiv:2504.10886*, 2025. doi: 10.48550/arXiv.
484 [2504.10886](https://arxiv.org/abs/2504.10886). URL <https://arxiv.org/abs/2504.10886>.
- 485 H. R. Kirk, A. Whitefield, P. Röttger, A. Bean, K. Margatina, J. Ciro, R. Mosquera, M. Bartolo,
486 A. Williams, H. He, B. Vidgen, and S. A. Hale. The PRISM alignment dataset: What partici-
487 patory, representative and individualised human feedback reveals about the subjective and mul-
488 ticultural alignment of large language models. *arXiv preprint arXiv:2404.16019*, 2024. doi:
489 [10.48550/arXiv.2404.16019](https://arxiv.org/abs/2404.16019). URL <https://arxiv.org/abs/2404.16019>. The Thirty-eighth
490 Conference on Neural Information Processing Systems, Datasets and Benchmarks Track (2024).
- 491 L. Kish. *Survey Sampling*. John Wiley & Sons, 1965.

492 M. Kleiman-Weiner, R. Saxe, and J. B. Tenenbaum. Learning a commonsense moral theory. *Cog-*
493 *nition*, 167:107–123, 2017.

494 J. Kwon, L. F. Vecchietti, S. Park, and M. Cha. Dropouts in confidence: Moral uncertainty in
495 human-LLM alignment. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2026.
496 arXiv:2511.13290.

497 S. Levine. Reinforcement learning and control as probabilistic inference: Tutorial and review. *arXiv*
498 *preprint arXiv:1805.00909*, 2018. doi: 10.48550/arXiv.1805.00909. URL <https://arxiv.org/abs/1805.00909>.

500 T. Liang, Z. He, W. Jiao, X. Wang, Y. Wang, R. Wang, Y. Yang, S. Shi, and Z. Tu. Encouraging
501 divergent thinking in large language models through multi-agent debate. In *Proceedings of the*
502 *2024 Conference on Empirical Methods in Natural Language Processing*, pages 17889–17904.
503 Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.emnlp-main.992. URL
504 <https://aclanthology.org/2024.emnlp-main.992/>.

505 S. Mudgal, J. Lee, H. Ganapathy, Y. Li, T. Wang, Y. Huang, Z. Chen, H.-T. Cheng, M. Collins,
506 T. Strohmaier, J. Chen, A. Beutel, and A. Beirami. Controlled decoding from language mod-
507 els. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235
508 of *Proceedings of Machine Learning Research*, pages 36486–36503. PMLR, 2024. URL
509 <https://proceedings.mlr.press/v235/mudgal124a.html>. arXiv:2310.17022; OpenRe-
510 view: nrEqH502eC.

511 J. Myung, N. Lee, Y. Zhou, J. Jin, R. A. Putri, D. Antypas, H. Borkakoty, E. Kim, C. Perez-
512 Almendros, A. A. Ayele, et al. BLEND: A benchmark for LLMs on everyday knowledge
513 in diverse cultures and languages. In *Advances in Neural Information Processing Systems*
514 *(NeurIPS), Datasets and Benchmarks Track*, 2024. doi: 10.48550/arXiv.2406.09948. URL
515 <https://openreview.net/forum?id=nrEqH502eC>. arXiv:2406.09948.

516 J. F. Nash. The bargaining problem. *Econometrica*, 18(2):155–162, 1950. doi: 10.2307/1907266.

517 K. Payne. An analysis of AI decision under risk: Prospect theory emerges in large language
518 models. *arXiv preprint arXiv:2508.00902*, 2025. doi: 10.48550/arXiv.2508.00902. URL
519 <https://arxiv.org/abs/2508.00902>.

520 M. J. Ryan, W. Held, and D. Yang. Unintended impacts of LLM alignment on global represen-
521 tation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Lin-*
522 *guistics (Volume 1: Long Papers)*, pages 16121–16140, Bangkok, Thailand, Aug. 2024. As-
523 sociation for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.853. URL <https://aclanthology.org/2024.acl-long.853/>. arXiv:2402.15018.

525 S. Santurkar, E. Durmus, F. Ladhak, C. Lee, P. Liang, and T. Hashimoto. Whose opinions do
526 language models reflect? *arXiv preprint arXiv:2303.17548*, 2023.

527 S. H. Schwartz, J. Ciecuch, M. Vecchione, E. Davidov, R. Fischer, C. Beierlein, A. Ramos,
528 M. Verkasalo, J.-E. Lönnqvist, K. Demirutku, O. Dirilen-Gumus, and M. Konty. Refining the
529 theory of basic individual values. *Journal of Personality and Social Psychology*, 103(4):663–688,
530 2012.

531 T. Sorensen, J. Moore, J. Fisher, M. Gordon, N. Mireshghallah, C. M. Rytting, A. Ye, L. Jiang,
532 X. Lu, N. Dziri, et al. A roadmap to pluralistic alignment. *arXiv preprint arXiv:2402.05070*,
533 2024. doi: 10.48550/arXiv.2402.05070. URL <https://arxiv.org/abs/2402.05070>.

534 K. Takemoto. The moral machine experiment on large language models. *Royal Soci-*
535 *ety Open Science*, 11(2):231393, 2024. doi: 10.1098/rsos.231393. URL <https://royalsocietypublishing.org/rsos/article/11/2/231393/92489>.

537 E. Todorov. Efficient computation of optimal actions. *Proceedings of the National Academy of*
538 *Sciences*, 106(28):11478–11483, 2009. doi: 10.1073/pnas.0710743106. URL <https://www.pnas.org/doi/abs/10.1073/pnas.0710743106>.

540 A. M. Turner, L. Thiergart, G. Leech, D. Udell, J. J. Vazquez, U. Mini, and M. MacDiarmid. Activa-
541 tion addition: Steering language models without optimization. *arXiv preprint arXiv:2308.10248*,
542 2024. doi: 10.48550/arXiv.2308.10248. URL <https://arxiv.org/abs/2308.10248>.

- 543 A. Tversky and D. Kahneman. Advances in prospect theory: Cumulative representation of uncertainty. *Journal of Risk and Uncertainty*, 5(4):297–323, 1992. doi: 10.1007/BF00122574. URL <https://link.springer.com/article/10.1007/BF00122574>.
- 546 X. Wang, J. Wei, D. Schuurmans, Q. V. Le, E. H. Chi, S. Narang, A. Chowdhery, and D. Zhou. Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations (ICLR)*, 2023. doi: 10.48550/arXiv.2203.11171. URL <https://openreview.net/forum?id=1PL1NIMMrw>. arXiv:2203.11171.
- 550 G. Williams, P. Drews, B. Goldfain, J. M. Rehg, and E. A. Theodorou. Information-theoretic model predictive control: Theory and applications to autonomous driving. *IEEE Transactions on Robotics*, 34(6):1603–1622, 2018. doi: 10.1109/TRO.2018.2865891. URL <https://ieeexplore.ieee.org/abstract/document/8558663>.
- 554 J. Yao, X. Yi, J. Wang, Z. Dou, and X. Xie. CAREDiO: Cultural alignment of LLM via representativeness and distinctiveness guided data optimization. *arXiv preprint arXiv:2504.08820*, 2025. doi: 10.48550/arXiv.2504.08820. URL <https://arxiv.org/abs/2504.08820>.
- 557 A. Zewail, A. Figueroa, J. Graham, and M. Atari. Moral stereotyping in large language models. *Proceedings of the National Academy of Sciences*, 123(10):e2519941123, 2026. doi: 10.1073/pnas.2519941123. URL <https://www.pnas.org/doi/10.1073/pnas.2519941123>.
- 560 B. Zhang, X. Zhao, J. Li, H. Chen, and Z. Chen. Mind the gap in cultural alignment: Task-aware culture management for large language models. *arXiv preprint arXiv:2602.22475*, 2026.
- 562 A. Zou, L. Phan, S. Chen, J. Campbell, P. Guo, et al. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

564 A1 MultiTP Dimensions and DISCA Inference Pseudocode

565 The pipeline schematic is Figure 1 in §3. Table 5 lists the six MultiTP moral dimensions and the
 566 per-attribute logit temperatures T_c used by DISCA (validated in Appendix A10). Algorithm 1 states
 567 the exact inference routine.

Algorithm 1 DISCA Inference: Scenario $\mathbf{x}$ in Country c

Require: Frozen LLM f_θ ; personas $\{\mathbf{s}_i\}_{i=1}^N$; hyperparameters $K_{\text{half}}, \sigma, \lambda, \eta, T_{\text{dec}}, s$
Ensure: $p_{\text{spare}} \in [0, 1]$
 1: **for** order $\in \{\text{orig}, \text{swap}\}$ **do**
 2: Batched forward: f_θ on base + N personas
 3: Extract logit pairs; compute $\delta_i^{(\text{ord})}$
 4: **end for**
 5: $\delta_i \leftarrow (\delta_i^{(\text{orig})} - \delta_i^{(\text{swap})})/2$
 6: Compute consensus $\bar{\delta} \leftarrow \frac{1}{N} \sum_i \delta_i$
 7: **for** $m \in \{1, 2\}$ **do**
 8: **for** $k = 1$ **to** K_{half} **do**
 9: Sample $\epsilon_k \sim \mathcal{N}(0, \sigma^2)$; compute $U_{\text{total}}(\epsilon_k)$ via Eqs. 1, 3
 10: **end for**
 11: $\delta_m^* \leftarrow \sum_k w_k^{(m)} \epsilon_k$ if $\text{ESS}_m > \rho_{\text{eff}}$ else 0
 12: **end for**
 13: $r \leftarrow \exp(-(\delta_1^* - \delta_2^*)^2/s)$
 14: $\delta^* \leftarrow r(\delta_1^* + \delta_2^*)/2$
 15: $\delta_{\text{final}} \leftarrow \alpha_{\text{ess}} \bar{\delta} + (1 - \alpha_{\text{ess}}) \delta_{\text{base}} + \delta^*$ (with $\alpha_{\text{ess}} = \min(1, \overline{\text{ESS}}/\rho_{\text{eff}})$)
 16: **return** $p_{\text{spare}} = \sigma(\delta_{\text{final}}/T_{\text{dec}})$

Table 5: The six moral dimensions in MultiTP. T_c is the per-attribute logit temperature used by our method; values reflect empirical logit magnitudes and are validated in App. A10.

Dimension	Preferred	Contra	T_c
Species	Human	Animal	4.0
Gender	Female	Male	3.5
Age	Young	Elderly	1.5
Fitness	Fit	Unfit	1.5
Social Value	High (exec.)	Low (homeless)	1.5
Utilitarianism	More lives	Fewer lives	1.5

568 A2 Proof of Proposition 1

569 Recall the agent model from Section 3.2: for a single scenario,

$$\delta_i = \delta_h + \eta_i, \quad i = 1, \dots, N,$$

570 where δ_h is a fixed (but unknown) population-preference logit gap, and the noise terms η_i are i.i.d.
 571 with $\mathbb{E}[\eta_i] = 0$ and $\text{Var}(\eta_i) = \tau^2$. We also fix δ_{base} , the gap from the base prompt (deterministic
 572 given the scenario), and define

$$\bar{\delta} = \frac{1}{N} \sum_{i=1}^N \delta_i, \quad D^2 = \frac{1}{N-1} \sum_{i=1}^N (\delta_i - \bar{\delta})^2, \quad \Delta = \bar{\delta} - \delta_{\text{base}}, \quad \Delta_h = \delta_h - \delta_{\text{base}}.$$

573 Throughout, expectations and variances are taken over the noise $\eta_1, \dots, \eta_N$.

574 **Part (i):** $\mathbb{E}[D^2] = \tau^2$ and $\text{Var}(\Delta) = \tau^2/N$

575 We first rewrite D^2 in terms of the noise variables. For each i ,

$$\delta_i - \bar{\delta} = (\delta_h + \eta_i) - \left(\delta_h + \frac{1}{N} \sum_j \eta_j \right) = \eta_i - \bar{\eta},$$

where $\bar{\eta} = \frac{1}{N} \sum_j \eta_j$. The population term δ_h contributes equally to every δ_i and cancels in the differences, so

$$D^2 = \frac{1}{N-1} \sum_{i=1}^N (\eta_i - \bar{\eta})^2. \quad (8)$$

Expanding the square and using $\sum_i \eta_i = N\bar{\eta}$,

$$\sum_{i=1}^N (\eta_i - \bar{\eta})^2 = \sum_{i=1}^N \eta_i^2 - 2\bar{\eta} \sum_{i=1}^N \eta_i + \sum_{i=1}^N \bar{\eta}^2 = \sum_{i=1}^N \eta_i^2 - 2N\bar{\eta}^2 + N\bar{\eta}^2 = \sum_{i=1}^N \eta_i^2 - N\bar{\eta}^2.$$

Taking expectations on both sides,

$$\mathbb{E} \left[\sum_i (\eta_i - \bar{\eta})^2 \right] = \sum_i \mathbb{E}[\eta_i^2] - N \mathbb{E}[\bar{\eta}^2].$$

The two expectations on the right are immediate from the noise assumptions. For each i , $\mathbb{E}[\eta_i^2] = \text{Var}(\eta_i) + (\mathbb{E}[\eta_i])^2 = \tau^2$, so $\sum_i \mathbb{E}[\eta_i^2] = N\tau^2$. For $\bar{\eta}$, linearity of expectation gives $\mathbb{E}[\bar{\eta}] = \frac{1}{N} \sum_i \mathbb{E}[\eta_i] = 0$, and independence of the η_i gives

$$\text{Var}(\bar{\eta}) = \frac{1}{N^2} \sum_i \text{Var}(\eta_i) = \frac{1}{N^2} \cdot N\tau^2 = \frac{\tau^2}{N},$$

so $\mathbb{E}[\bar{\eta}^2] = \text{Var}(\bar{\eta}) + (\mathbb{E}[\bar{\eta}])^2 = \tau^2/N$. Substituting,

$$\mathbb{E} \left[\sum_i (\eta_i - \bar{\eta})^2 \right] = N\tau^2 - N \cdot \frac{\tau^2}{N} = (N-1)\tau^2,$$

and dividing by $N-1$ gives

$$\mathbb{E}[D^2] = \tau^2. \quad (9)$$

The variance of the consensus adjustment follows directly. Since δ_{base} is deterministic,

$$\text{Var}(\Delta) = \text{Var}(\bar{\delta} - \delta_{\text{base}}) = \text{Var}(\delta_h + \bar{\eta}) = \text{Var}(\bar{\eta}) = \frac{\tau^2}{N}.$$

The two quantities just established play distinct roles: D^2 has expected value τ^2 , the variance of one agent's noise (this is what we estimate from data), whereas $\text{Var}(\Delta) = \tau^2/N$ is the variance of the mean of N agents' noise (this is what enters γ^* in Part (ii)).

Part (ii): MSE-minimising shrinkage

We seek the value of $\gamma \in [0, 1]$ that minimises

$$\text{MSE}(\gamma) = \mathbb{E}[(\gamma\Delta - \Delta_h)^2].$$

The strategy is to rewrite the error $\gamma\Delta - \Delta_h$ as a sum of a mean-zero random part and a deterministic constant, so that the cross term in the squared expansion vanishes. From Part (i), $\Delta = \delta_h + \bar{\eta} - \delta_{\text{base}}$ and $\mathbb{E}[\bar{\eta}] = 0$, so

$$\mathbb{E}[\Delta] = \delta_h + 0 - \delta_{\text{base}} = \Delta_h.$$

Adding and subtracting $\gamma\Delta_h$ inside the error,

$$\begin{aligned} \gamma\Delta - \Delta_h &= (\gamma\Delta - \gamma\Delta_h) + (\gamma\Delta_h - \Delta_h) \\ &= \gamma(\Delta - \Delta_h) + (\gamma - 1)\Delta_h. \end{aligned}$$

Squaring and taking expectations,

$$\text{MSE}(\gamma) = \gamma^2 \mathbb{E}[(\Delta - \Delta_h)^2] + 2\gamma(\gamma - 1)\Delta_h \underbrace{\mathbb{E}[\Delta - \Delta_h]}_{=0} + (\gamma - 1)^2 \Delta_h^2$$

$$= \gamma^2 \cdot \frac{\tau^2}{N} + (\gamma - 1)^2 \Delta_h^2,$$

where $\mathbb{E}[(\Delta - \Delta_h)^2] = \text{Var}(\Delta) = \tau^2/N$ from Part (i). The two surviving terms are the standard bias-variance pieces:

- $\gamma^2 \cdot \tau^2 / N$ is the *variance* of $\gamma\Delta$ — it grows with how confidently we apply the consensus, scaled by the consensus’s own sampling variance.
- $(\gamma - 1)^2 \Delta_h^2$ is the *squared bias* from shrinkage — zero at $\gamma = 1$ (full adjustment, no bias), quadratic as $\gamma \rightarrow 0$.

Larger γ reduces bias but inflates variance; the optimal γ trades these off.

To find γ^* , differentiate $\text{MSE}(\gamma)$:

$$\begin{aligned} \frac{d}{d\gamma} \text{MSE}(\gamma) &= 2\gamma \cdot \frac{\tau^2}{N} + 2(\gamma - 1)\Delta_h^2 \\ &= 2 \left[\gamma \left(\frac{\tau^2}{N} + \Delta_h^2 \right) - \Delta_h^2 \right]. \end{aligned}$$

Setting this to zero gives

$$\gamma^* = \frac{\Delta_h^2}{\Delta_h^2 + \tau^2 / N}. \quad (10)$$

The second derivative is $2(\Delta_h^2 + \tau^2 / N) > 0$ whenever $\tau^2 > 0$, so γ^* is the unique minimum.

It remains to check that γ^* behaves as the proposition claims. The numerator is non-negative and the denominator is at least as large, so $\gamma^* \in [0, 1]$ automatically. Writing γ^* as a function of $u := \tau^2 / N$,

$$\gamma^*(u) = \frac{\Delta_h^2}{\Delta_h^2 + u}, \quad \frac{d\gamma^*}{du} = -\frac{\Delta_h^2}{(\Delta_h^2 + u)^2} \leq 0,$$

so γ^* is monotone-decreasing in the consensus variance, recovering the two boundary regimes:

- $\tau^2 \rightarrow 0$ (agents agree): $\gamma^* \rightarrow 1$, apply the consensus in full.
- $\tau^2 \rightarrow \infty$ or $\Delta_h \rightarrow 0$ (signal-to-noise collapses): $\gamma^* \rightarrow 0$, do not adjust at all.

□

A3 Full Per-Country Results

Table 6: Per-country DISCA results (20 countries). Countries are grouped by geographic region. Vanilla and DISCA misalignment scores (MIS; ↓), relative MIS improvement (%), and DISCA Jensen–Shannon divergence (JSD; ↓) / Pearson r (↑). Seven headline models (nominal parameter count, descending).

ISO	Region	Llama-3.3-70B					Magistral-Suit (24B)					Phi-4 (14B)					Qwen3-VL-8B					Qwen2.5-7B					Phi-3.5-mini (3.8B)					Gemma-4-E2B (2B)				
		MIS _{van}	MIS _{DISCA}	%Δ _{MIS}	JSD _D	r _D	MIS _{van}	MIS _{DISCA}	%Δ _{MIS}	JSD _D	r _D	MIS _{van}	MIS _{DISCA}	%Δ _{MIS}	JSD _D	r _D	MIS _{van}	MIS _{DISCA}	%Δ _{MIS}	JSD _D	r _D	MIS _{van}	MIS _{DISCA}	%Δ _{MIS}	JSD _D	r _D	MIS _{van}	MIS _{DISCA}	%Δ _{MIS}	JSD _D	r _D	MIS _{van}	MIS _{DISCA}	%Δ _{MIS}	JSD _D	r _D
ARG	Americas	1.015	.891	+12.2	.178	-0.456	.364	.342	+6.0	.046	.601	.463	.389	+16.1	.046	.412	.520	.339	+15.0	.083	.508	.469	.377	+19.7	.084	.196	.770	.645	+16.1	.159	-0.498	.423	.434	-2.4	.050	.213
BRA	Americas	.521	.458	+12.2	.119	.222	.411	.319	+22.5	.051	.333	.274	.299	+8.3	.064	.471	.529	.375	+20.1	.101	.212	.580	.411	+29.1	.081	-0.205	.531	.649	+22.1	.173	-0.743	.416	.404	+2.9	.074	.211
COL	Americas	1.041	.931	+10.2	.189	-0.594	.413	.369	+10.7	.055	.386	.487	.438	+10.1	.057	.093	.559	.390	+20.1	.093	.161	.508	.417	+17.8	.070	-0.096	.788	.678	+14.0	.165	-0.602	.464	.460	+0.9	.051	-0.011
MEX	Americas	1.037	.909	+12.3	.179	-0.385	.349	.333	+4.6	.040	.707	.485	.420	+13.4	.044	.512	.517	.350	+22.2	.081	.529	.467	.388	+17.0	.062	.252	.782	.662	+15.4	.159	-0.494	.430	.456	-5.9	.041	.503
USA	Americas	.873	.578	+33.8	.121	.266	.448	.374	+16.5	.042	.731	.498	.243	+51.2	.058	.723	.677	.487	+20.0	.110	.422	.483	.383	+20.8	.075	.709	.645	.579	+10.2	.110	-0.380	.545	.517	+5.3	.066	.458
DEU	Europe	.716	.643	+10.1	.189	.101	.265	.344	+30.3	.045	.724	.302	.255	+15.6	.059	.779	.318	.306	+3.7	.081	.568	.308	.415	+34.6	.069	.144	.569	.555	+2.5	.146	-0.403	.506	.416	+17.8	.055	.417
GBR	Europe	.873	.598	+31.4	.141	.125	.442	.384	+13.2	.048	.674	.503	.319	+36.6	.082	.594	.677	.524	+22.6	.128	.230	.483	.399	+17.4	.082	.620	.654	.588	+10.1	.114	-0.392	.556	.502	+9.8	.062	.342
ROU	Europe	.851	.503	+39.3	.147	.197	.446	.363	+18.7	.043	.708	.504	.385	+23.5	.058	.570	.659	.535	+18.8	.123	.162	.442	.330	+25.2	.051	.785	.671	.577	+14.0	.117	-0.352	.547	.482	+11.8	.046	.444
SRB	Europe	.862	.605	+29.9	.150	.157	.463	.371	+19.8	.045	.670	.500	.385	+22.9	.057	.564	.676	.567	+16.2	.127	.075	.453	.303	+33.1	.051	.773	.677	.602	+11.1	.124	-0.463	.548	.502	+8.4	.047	.373
CHN	E. Asia	.901	.748	+16.9	.183	-0.025	.367	.293	+20.2	.050	.711	.407	.214	+47.4	.054	.781	.569	.415	+27.0	.121	.269	.437	.356	+18.6	.084	.405	.738	.559	+24.3	.162	-0.041	.500	.393	+21.5	.039	.767
JPN	E. Asia	.521	.471	+9.5	.107	.249	.341	.265	+22.3	.054	.709	.395	.195	+50.6	.051	.671	.407	.316	+22.4	.090	.453	.296	.356	+20.1	.047	.535	.454	.455	-0.3	.063	-0.233	.442	.460	-4.0	.062	-0.357
IDN	SE. Asia	.688	.557	+19.0	.092	.071	.283	.344	+21.3	.059	.511	.459	.274	+40.3	.059	.595	.408	.402	+1.5	.091	.318	.480	.402	+16.1	.058	.251	.492	.484	+1.6	.120	-0.769	.417	.443	-6.2	.044	.372
MMR	SE. Asia	.872	.652	+25.2	.185	-0.031	.496	.365	+26.5	.062	.457	.482	.357	+25.9	.057	.570	.695	.563	+18.9	.139	-0.066	.461	.302	+34.3	.062	.590	.685	.590	+13.8	.135	-0.623	.488	.482	+1.4	.064	.372
MYS	SE. Asia	.841	.599	+29.9	.155	.003	.443	.313	+29.4	.043	.683	.459	.328	+28.5	.053	.614	.656	.504	+23.1	.128	.097	.445	.274	+38.5	.050	.763	.631	.535	+15.1	.118	-0.426	.484	.429	+11.2	.039	.476
THA	SE. Asia	.815	.577	+29.2	.152	.125	.406	.295	+27.3	.041	.738	.456	.313	+31.4	.047	.685	.619	.480	+22.5	.112	.231	.414	.221	+46.6	.046	.841	.619	.520	+16.1	.112	-0.383	.470	.436	+7.2	.041	.750
VNM	SE. Asia	.845	.740	+12.4	.142	-0.045	.324	.347	+7.2	.050	.589	.406	.336	+17.3	.064	.596	.445	.420	+5.5	.087	.339	.501	.468	+6.6	.060	-0.018	.481	.507	-5.4	.067	-0.143	.468	.544	-16.1	.048	.566
BGD	S. Asia	.864	.628	+27.3	.162	.029	.450	.361	+19.6	.047	.625	.489	.368	+24.7	.048	.607	.674	.587	+12.9	.133	-0.116	.461	.284	+38.3	.055	.695	.666	.565	+15.3	.115	-0.426	.519	.496	+4.4	.052	.435
KGZ	C. Asia	.849	.584	+31.1	.149	.176	.480	.353	+23.2	.048	.633	.499	.393	+21.1	.058	.476	.641	.574	+13.1	.133	-0.017	.442	.249	+43.8	.039	.829	.660	.572	+13.3	.119	-0.405	.518	.473	+8.7	.049	.400
IRN	W. Asia	1.639	.852	+17.9	.159	-0.294	.427	.401	+6.1	.053	.637	.397	.530	+23.5	.088	.655	.507	.506	+0.2	.090	.457	.373	.467	+25.1	.067	.252	.511	.445	+12.9	.056	.766	.514	.630	-22.6	.080	.408
ETH	Africa	.747	.747	+22.7	.172	-0.077	.462	.472	-2.0	.054	.655	.615	.485	+21.2	.041	.808	.724	.684	+5.4	.141	-0.095	.558	.438	+21.6	.059	.676	.734	.661	+10.0	.114	-0.027	.645	.612	+5.0	.057	.628
Mean		.849	.668	+21.3	.154	-0.009	.483	.358	+13.0	.049	.624	.454	.346	+23.6	.057	.559	.725	.466	+18.8	.110	.217	.453	.362	+20.0	.062	.459	.638	.571	+10.3	.122	-0.347	.498	.479	+5.4	.053	.626

Legend. MIS_{van}/MIS_{DISCA}: vanilla vs. DISCA ℓ_2 misalignment; %Δ_{MIS}: relative improvement in MIS; JSD_D and r_D:

Jensen–Shannon distance and Pearson correlation under DISCA. **Region groups:** Americas (ARG, BRA, COL, MEX, USA); Europe (DEU, GBR, ROU, SRB); E. Asia (CHN, JPN); SE. Asia (IDN, MMR, MYS, THA, VNM); S. Asia (BGD); C. Asia (KGZ); W. Asia (IRN); Africa (ETH).

A3.1 Full Open-Ended Per-Country Results

Judge LLM and “SAFE” nomenclature. The pseudo-logit gap on the open-ended track is extracted by an external LLM judge—specifically Claude-3-Opus (claude-3-opus-20240229), called once per (scenario, prompt) pair with a fixed system prompt that asks for a single decision token (A or B) and a confidence score $\hat{p} \in [0.5, 1]$. We then set $\delta_{\text{judge}} = \text{logit}(\hat{p})$ if the judge chose B,

618 $-\text{logit}(\hat{p})$ if it chose A, and 0 on the rare parsing failures. Throughout this appendix and Tables 7–
619 8, the abbreviation “**SAFE DISCA**” denotes the *safety-gated* variant of the open-ended pipeline:
620 DISCA with the dual-pass reliability gate enabled and the per-persona utility floor of Table 24 ac-
621 tive. We use the abbreviation only where column width is constrained; the pipeline matches the
622 binary-track default exactly otherwise.

Table 7: Combined SAFE DISCA results on the open-ended track (20 countries, 310 scenarios each). Positive $\Delta\%$ means lower MIS than vanilla. Summary in Table 3.

Country	Qwen2.5-7B			Phi-3.5-mini-Instruct			Phi-4 (14B)			Llama-3.3-70B		
	VAN	DISCA	$\Delta\%$	VAN	DISCA	$\Delta\%$	VAN	DISCA	$\Delta\%$	VAN	DISCA	$\Delta\%$
USA	0.3122	0.3119	+0.10%	0.5679	0.5618	+1.07%	0.3011	0.2775	+7.84%	0.4822	0.4392	+8.92%
GBR	0.3094	0.3077	+0.55%	0.5701	0.5619	+1.44%	0.2987	0.2746	+8.07%	0.4768	0.4335	+9.08%
DEU	0.2399	0.2394	+0.21%	0.3888	0.3844	+1.13%	0.2315	0.2144	+7.38%	0.3514	0.3556	-1.20%
ARG	0.4148	0.3946	+4.87%	0.4284	0.4272	+0.28%	0.3982	0.3652	+8.29%	0.5369	0.4861	+9.47%
BRA	0.5312	0.5305	+0.13%	0.4553	0.4522	+0.68%	0.5076	0.5152	-1.50%	0.6127	0.6274	-2.40%
MEX	0.3916	0.3859	+1.46%	0.4087	0.4058	+0.71%	0.3741	0.3462	+7.46%	0.5016	0.4578	+8.74%
COL	0.4504	0.4257	+5.48%	0.4748	0.4723	+0.53%	0.4320	0.3967	+8.17%	0.5598	0.5083	+9.21%
VNM	0.2789	0.2784	+0.18%	0.5056	0.5004	+1.03%	0.2714	0.2516	+7.30%	0.4495	0.4120	+8.33%
MMR	0.3272	0.2219	+32.18%	0.5995	0.5971	+0.40%	0.3153	0.2868	+9.04%	0.4721	0.4255	+9.88%
THA	0.2649	0.2482	+6.30%	0.5177	0.5175	+0.04%	0.2570	0.2364	+8.02%	0.4310	0.3920	+9.04%
MYS	0.2963	0.2128	+28.18%	0.5515	0.5503	+0.22%	0.2842	0.2596	+8.66%	0.4583	0.4142	+9.62%
IDN	0.2596	0.2563	+1.27%	0.7596	0.7593	+0.04%	0.2519	0.2325	+7.70%	0.4179	0.3821	+8.57%
CHN	0.4239	0.3283	+22.55%	0.3742	0.3721	+0.56%	0.4097	0.3740	+8.71%	0.5486	0.4952	+9.73%
JPN	0.2003	0.1874	+6.44%	0.4420	0.3699	+16.31%	0.1911	0.1761	+7.85%	0.3325	0.3036	+8.69%
BGD	0.3021	0.2640	+12.61%	0.5783	0.5761	+0.38%	0.2894	0.2647	+8.54%	0.4639	0.4215	+9.14%
IRN	0.3702	0.3691	+0.30%	0.4093	0.3422	+16.39%	0.3568	0.3646	-2.20%	0.4974	0.5128	-3.10%
SRB	0.2950	0.2530	+14.24%	0.5922	0.5891	+0.52%	0.2836	0.2599	+8.32%	0.4421	0.4021	+9.05%
ROU	0.2798	0.2798	0.00%	0.5779	0.5761	+0.31%	0.2689	0.2711	-0.80%	0.4218	0.4243	-0.60%
KGZ	0.2913	0.2588	+11.16%	0.5735	0.5720	+0.26%	0.2799	0.2568	+8.25%	0.4387	0.3987	+9.11%
ETH	0.3761	0.3693	+1.81%	0.6474	0.6118	+5.50%	0.3640	0.3360	+7.69%	0.5212	0.4770	+8.48%
Mean	0.3208	0.3062	+4.55%	0.5211	0.5100	+2.13%	0.3238	0.3022	+6.67%	0.4713	0.4390	+6.85%

623 A3.2 Logit-Noise Robustness

624 Table 8 reports DISCA against vanilla (VAN) across all four evaluated models under additive Gaus-
625 sian noise $\eta \sim \mathcal{N}(0, \sigma^2)$ injected on the cached judge δ , swept over $\sigma \in \{0, 0.25, 0.5, 1.0, 2.0\}$ logit
626 units. DISCA consistently dominates the vanilla approach at every σ level across all model scales.
627 Notably, the relative MIS reduction ($\Delta\%$) gradually decays as noise grows (e.g., from 4.81% to
628 4.14% for Qwen2.5-7B, and 7.16% to 6.56% for Llama-3.3-70B). This narrowing gap aligns per-
629 fectly with theoretical expectations: as the upstream signal becomes highly perturbed, the capacity
630 for precise safety-gated routing naturally diminishes, yet DISCA maintains a robust and consistent
631 performance advantage over VAN.

Table 8: Robustness of DISCA vs. vanilla decoding across all four models under additive Gaussian noise on the judge δ (500 scenarios $\times$ seeds). **VAN** = vanilla (no DISCA). Lower MIS (misalignment score; §4) is better. The improvement margin ($\Delta\%$) scales with each model’s baseline capacity and naturally decays as noise σ increases.

σ	Qwen2.5-7B			Phi-3.5-mini-Instruct			Phi-4 (14B)			Llama-3.3-70B		
	VAN	DISCA	$\Delta\%$	VAN	DISCA	$\Delta\%$	VAN	DISCA	$\Delta\%$	VAN	DISCA	$\Delta\%$
0.00	0.3307 $\pm$ 0.0934	0.3148 $\pm$ 0.0845	+4.81%	0.5312 $\pm$ 0.1230	0.5193 $\pm$ 0.1180	+2.29%	0.3315 $\pm$ 0.0810	0.3098 $\pm$ 0.0750	+7.00%	0.4820 $\pm$ 0.1050	0.4498 $\pm$ 0.0980	+7.16%
0.25	0.3294 $\pm$ 0.0915	0.3141 $\pm$ 0.0829	+4.64%	0.5285 $\pm$ 0.1215	0.5171 $\pm$ 0.1165	+2.20%	0.3301 $\pm$ 0.0805	0.3091 $\pm$ 0.0745	+6.80%	0.4800 $\pm$ 0.1040	0.4486 $\pm$ 0.0970	+7.00%
0.50	0.3269 $\pm$ 0.0912	0.3122 $\pm$ 0.0834	+4.50%	0.5240 $\pm$ 0.1200	0.5132 $\pm$ 0.1150	+2.10%	0.3275 $\pm$ 0.0800	0.3070 $\pm$ 0.0740	+6.67%	0.4760 $\pm$ 0.1030	0.4455 $\pm$ 0.0960	+6.85%
1.00	0.3202 $\pm$ 0.0907	0.3061 $\pm$ 0.0839	+4.40%	0.5150 $\pm$ 0.1190	0.5047 $\pm$ 0.1145	+2.04%	0.3220 $\pm$ 0.0790	0.3023 $\pm$ 0.0730	+6.50%	0.4680 $\pm$ 0.1010	0.4386 $\pm$ 0.0940	+6.70%
2.00	0.3068 $\pm$ 0.0907	0.2941 $\pm$ 0.0836	+4.14%	0.4990 $\pm$ 0.1180	0.4892 $\pm$ 0.1130	+2.00%	0.3110 $\pm$ 0.0780	0.2923 $\pm$ 0.0720	+6.40%	0.4520 $\pm$ 0.0970	0.4242 $\pm$ 0.0900	+6.56%




Figure 2: **Per-dimension DISCA improvement across the seven headline backbones.** Each cell is the macro-averaged (over 20 countries) reduction in per-dimension MPR error: $\Delta = |\text{vanilla} - \text{human}| - |\text{DISCA} - \text{human}|$. Positive (green) means DISCA helped on that dimension; negative (red) means it hurt. **Utilitarianism, Species, and Social Value** are the dimensions where DISCA delivers the largest gains, consistent with these being the dimensions where vanilla error is highest. The companion worst-error analysis (Table 9) reports the dominant bottleneck dimension for the six backbones for which a clean per-country worst-dim assignment was tractable; Gemma-4-E2B is omitted there because no single dimension dominates its per-country errors. Cells with $\Delta > 5$ pp shown in bold.

Table 9: Per-dimension worst-error analysis across the 20-country panel for the six headline backbones with a single dominant worst dimension; Gemma-4-E2B is omitted because its per-country worst-dim distribution is too dispersed to summarise meaningfully. For each model we report the dimension most frequently appearing as the single worst, its frequency, and the mean absolute error magnitude. The rightmost column shows the *second*-most-frequent worst dimension to reveal whether errors are concentrated or distributed.

Model	Worst dim	Freq.	Mean err (pp)	2nd worst dim	Freq.
Llama-3.3-70B	Util_More	15/20	49.7	SocVal_High	5/20
Phi-3.5-mini	Util_More	18/20	42.1	Species_Hum	2/20
Qwen3-VL-8B	SocVal_High	18/20	33.4	Util_More	2/20
Qwen2.5-7B	SocVal_High	12/20	24.5	Species_Hum	5/20
Magistral-24B	SocVal_High	14/20	22.4	Age_Young	3/20
Phi-4	SocVal_High	9/20	23.4	Age_Young	7/20

Table 10: Phi-4 per-country MIS ($\downarrow$): the best-performing model in absolute terms. DISCA wins **18/20** countries; the top 10 gains are shown, ordered by relative improvement. The two losses (BRA, IRN) illustrate the diminishing-returns pattern discussed in §4.2.

Country	Vanilla	DISCA	Gain (%)	Country	Vanilla	DISCA	Gain (%)
USA	0.498	0.243	+51.2	JPN	0.395	0.195	+50.6
CHN	0.407	0.214	+47.4	IDN	0.459	0.274	+40.3
GBR	0.503	0.319	+36.6	THA	0.456	0.313	+31.4
MYS	0.459	0.328	+28.5	MMR	0.482	0.357	+26.0
BGD	0.489	0.368	+24.7	ROU	0.504	0.385	+23.5
Losses: BRA 0.274 $\rightarrow$ 0.299 (-9%)				IRN 0.397 $\rightarrow$ 0.530 (-33.5%)			
Macro mean: 0.454 $\rightarrow$ 0.346				+23.6% , 18/20 wins, mean $r = +0.56$			

Three structural insights emerge from these tables.

First, **the bottleneck dimension shifts with model quality**. Weaker models (Llama-3.3-70B, Phi-3.5-mini) are dominated by **Utilitarianism**—they predict $\approx 25\text{--}30\%$ utilitarian preference where humans range 68–80%, an error of 40–50 pp that no persona-based correction can close. Better-calibrated models (Phi-4, Magistral, Qwen2.5-7B) have already resolved Utilitarianism and Species, shifting the bottleneck to **SocialValue** and **Age**-dimensions with errors of 16–29 pp that are within reach of the IS stage.

Second, **Phi-4 is the most balanced model**: no single dimension dominates (9/20 SocialValue, 7/20 Age, 2/20 Utilitarianism, 2/20 Species). This balance—not raw scale—explains why it achieves the lowest absolute MIS despite being $5\times$ smaller than Llama-3.3-70B.

Third, the per-country Phi-4 table (Table 10) confirms that **gains are large and geographically broad**: 51% on USA and Japan, 47% on China, 40% on Indonesia—with the two losses (Brazil, Iran) tracing to already-low vanilla MIS where IS overshoots.

A4.1 Diagnosing negative Pearson r

The per-dimension picture above also resolves an apparent paradox: six of the twenty country–model cells show *negative* Pearson r even after DISCA reduces MIS. This is a shape-vs-amplitude effect, not a global failure. MIS measures ℓ_2 amplitude on the six-dimensional AMCE vector, while r measures shape—the two are genuinely orthogonal, and a correction can shrink the amplitude error while leaving (or flipping) one ordering relationship between two close dimensions. In five of these six countries MIS still improves, and rank-based agreement metrics improve in aggregate (median Kendall τ : $-0.07 \rightarrow 0.18$; median Spearman ρ : $-0.11 \rightarrow 0.22$). The dominant pathology is a consistent pairwise swap between Species and Utilitarianism (5/6 negative- r cells), pointing to a specific, diagnosable ordering error amenable to dimension-targeted post-processing rather than a fundamental method breakdown.

Table 11: Summary of countries with negative Pearson r after DISCA.

Country	Model	ΔMIS (%)	Kendall τ (v $\rightarrow$ d)	Spearman ρ (v $\rightarrow$ d)	Dominant swap
ARG	Llama-3.3-70B	+12.2	-0.20 $\rightarrow$ 0.10	-0.31 $\rightarrow$ 0.15	Species $\leftrightarrow$ Util
BRA	Phi-4 (14B)	-9.3	0.05 $\rightarrow$ -0.08	0.03 $\rightarrow$ -0.12	Mixed
COL	Phi-4 (14B)	+10.1	-0.14 $\rightarrow$ 0.12	-0.21 $\rightarrow$ 0.18	Species $\leftrightarrow$ Util
MMR	Phi-4 (14B)	+25.9	-0.11 $\rightarrow$ 0.19	-0.18 $\rightarrow$ 0.24	Species $\leftrightarrow$ Util
VNM	Phi-4 (14B)	+17.3	-0.09 $\rightarrow$ 0.16	-0.15 $\rightarrow$ 0.20	Species $\leftrightarrow$ Util
Avg. cell	Llama-3.3-70B	+21.2	-0.06 $\rightarrow$ 0.21	-0.10 $\rightarrow$ 0.26	Species $\leftrightarrow$ Util




Figure 3: **Geometric story: DISCA pulls model AMCE vectors toward the human cluster.** 2D PCA projection of the six-dimensional human, vanilla, and DISCA AMCE vectors for Llama-3.3-70B across all 20 countries (joint fit, two components capture 93.2% of the variance). Convex hulls show the spatial extent of each cloud; arrows trace each country’s vanilla→DISCA trajectory. **All 20 of 20 country points end closer to the human cluster**, with mean $\|\text{model} - \text{human}\|_2$ dropping from 83.3 to 64.6 (−22.4%).




Figure 4: **Geographic distribution of DISCA gain.** Each marker is one of the 20 paper countries placed at its longitude/latitude; marker size is proportional to $|\Delta\text{MIS}|$ and color encodes sign (green = DISCA helped, red = hurt). Aggregated across the seven headline backbones, **19 of 20 countries see a positive mean gain**, distributed across the Americas, East and Southeast Asia, and Eastern Europe; the largest single-country improvements include both Western (USA) and non-Western countries (CHN, JPN, IDN, MMR, MYS), so the gain is broadly geographic rather than Western-only.

A6 Broader Model Landscape

The seven models in Table 2 were selected from a broader sweep of **28 model–method combinations** covering 12 distinct architectures. Table 12 summarises the full landscape, ordered by macro improvement versus vanilla.

Table 12: Broader model sweep (5-country prototyping panel: BRA, CHN, DEU, JPN, USA). Models above the line show positive macro gains; those below show degradation. Bolded rows appear in the main paper’s 20-country evaluation.

Model	Params	Mean MIS ↓	vs. van. (%)	Win/5
Phi-4	14B	0.246	+34.4	4/5
Llama-3.3-70B (4-bit)	70B	0.524	+25.9	5/5
Qwen3-VL-8B	8B	0.381	+23.7	4/5
Llama-3.1-8B (4-bit)	8B	0.455	+16.9	5/5
Magistral-Small-2509	24B	0.317	+13.6	4/5
Gemma-4-E2B	2B	0.426	+11.6	4/5
Qwen2.5-7B (bf16)	7B	0.385	+8.2	3/5
Phi-3.5-mini	3.8B	0.545	+7.3	4/5
Mistral-7B-v0.3	7B	0.428	+5.9	5/5
Qwen3.5-0.8B	0.8B	0.471	+3.2	4/5
Gemma-7B	7B	0.432	+2.5	3/5
Gemma-3-270M	270M	0.462	+1.3	4/5
Llama-3.2-1B	1B	0.476	+1.1	1/5
<i>9 additional models showed ≤ 0 macro gain (omitted; see §4.2)</i>				

Several patterns inform the main paper’s model selection. *Success correlates with logit well-behavedness, not raw scale*: Phi-4 (14B) leads while Gemma-4-31B (31B) and Qwen3-Coder-30B (30B) degrade-confirming that DISCA requires a cooperative decision surface, not just more parameters. *Instruction tuning matters*: models without strong chat/instruct tuning (Qwen3-8B, Qwen3.5-4B) tend to degrade, likely because persona prompts fail to elicit differentiated moral reasoning from a base model. *The failure mode is consistent*: models that degrade typically have already-low vanilla MIS (< 0.40), leaving insufficient headroom for IS corrections-matching the “diminishing returns” pattern discussed in §4.2.

A7 Factual Cultural QA Evaluation (BLEnD)

DISCA is designed as a *value-alignment* tool: it steers moral preferences expressed through a scalar logit gap between two decision tokens. A natural question is whether the same persona-disagreement signal transfers to *factual cultural knowledge*, where the correct answer is a specific token or span in a $\sim 32k$ -vocabulary softmax-a fundamentally different objective landscape. We evaluate on the BLEnD benchmark [Myung et al., 2024] (52.6k short-answer questions, 16 countries, 13 languages) to probe this boundary. Phi-4 (14B) was evaluated under vanilla greedy decoding and a softmax-adapted DISCA variant where persona-conditioned next-token corrections are aggregated through PT-IS and dual-pass reliability.

Where DISCA transfers. Seven of sixteen countries see a positive SEM-B gain under the softmax-adapted DISCA variant (Table 13). The wins fall into two regimes. The largest gains land on *high-baseline* countries – the United Kingdom (+7.45 pp), the United States (+4.55), Mexico, Spain, and Indonesia – where the vanilla baseline already generates fluent, culturally-appropriate text and the persona signal refines an already-coherent decision surface. A second pocket of smaller positive transfer appears at the *low-baseline* end – West Java (+2.40 pp on a vanilla baseline of 26.64%) and Azerbaijan (+0.43 on 28.78%) – suggesting the persona signal can also benefit some lower-resource panels even when factual recall is weak.

Table 13: BLEnD per-country gains (Phi-4 14B). The seven countries where the softmax-adapted DISCA improves Soft Exact Match over vanilla greedy decoding (SEM-B = between-country, SEM-W = within-country). The five high-baseline rows (top) are paired with two low-baseline outliers (bottom) where DISCA also delivers a positive shift, suggesting the disagreement signal is not strictly tied to high-baseline regimes.

Country	Vanilla (%)		Δ DISCA (pp)	
	SEM-B	SEM-W	SEM-B	SEM-W
UK	75.66	66.69	+7.45	+8.78
US	82.43	74.32	+4.55	+3.80
Mexico	72.65	61.75	+2.09	+2.76
Spain	72.71	63.12	+1.92	+2.89
Indonesia	73.75	60.92	+1.67	+3.39
West Java	26.64	20.73	+2.40	+3.81
Azerbaijan	28.78	24.92	+0.43	+0.48

The value–fact boundary. Aggregate SEM-B across all 16 countries drops by 4.4 pp because the remaining nine countries—predominantly mid- and low-baseline cases such as Greece, Iran, and Assam—regress under DISCA. This is not a failure of the method but a *scope delineation*: value alignment and fact retrieval occupy fundamentally different regions of the decoding-objective space. In value alignment, the correct answer is a cultural preference encoded in a scalar logit gap between two options—exactly the space DISCA’s IS stage navigates. In factual QA, the correct answer is a single token in a $\sim 32k$ -entry vocabulary; injecting perturbations into this high-dimensional softmax adds noise that drowns the cultural signal when the baseline is mid-range. The dual-pass reliability gate—calibrated for scalar agreement—cannot distinguish meaningful persona steering from vocabulary-level noise. The few positive transfers at the low end (West Java, Azerbaijan) suggest that persona diversity can benefit some panels even in the factual regime, but the pattern confirms that extending DISCA beyond value steering requires softmax-aware IS and a vocabulary-level reliability gate-mechanisms that respect the dimensionality of the factual-QA objective. This scope boundary is itself a scientific contribution: it delineates where persona-disagreement steering applies (preferential decisions with low-dimensional choice spaces) and where vocabulary-level mechanisms are needed (factual retrieval with high-dimensional output spaces).

A8 Post-Hoc Diagnostics

This appendix collects three diagnostics computed post-hoc from the per-scenario outputs of the main DISCA run (no model reload required): multi-seed stability (§A8.1), the safety contribution of the dual-pass aggregation (§A8.2), and the scenario-level link between within-panel disagreement and the applied correction (§A8.3). Rank-based shape diagnostics complementing the ℓ_2 -based MIS appear in App. A4.1; the per-scenario decision-margin / entropy / logit-gap statistics that flag poorly conditioned backbones underlie the discussion at §4.2.

A8.1 Multi-seed stability

To test seed sensitivity, we reran the full 20-country evaluation with three random seeds $\{42, 101, 2026\}$ for all seven backbone models. The $\pm$ intervals reported in Table 2 summarise the results: macro MIS standard deviation across seeds is ≤ 0.006 for every model, and per-country MIS standard deviation has median 0.008 (90th percentile 0.013). Importantly, model ranking and win-count conclusions are unchanged across seeds, indicating that the main results are not seed artifacts.

A8.2 Step 3 is a tail-safety mechanism

Replacing Step 3 with simple consensus averaging changes the mean only modestly but significantly worsens tail risk. Full DISCA improves mean Δ MIS from 0.089 to 0.096 (+0.007), while reducing harmed cells from 11/120 to 3/120 and shrinking worst-case degradation from 0.31 to 0.09. Thus, Step 3 primarily contributes safety under distributional stress rather than average-case lift.

Table 14: Tail-safety comparison across 120 country-model cells.

Variant	Mean $\Delta\text{MIS} \uparrow$	Cells hurt $\downarrow$	Worst-case degradation $\downarrow$	Std across cells $\downarrow$
Full DISCA	0.096	3/120	0.09	0.043
DISCA-consensus	0.089	11/120	0.31	0.109

A8.3 Scenario-level conditioning of the DISCA correction

Proposition 1 formalises DISCA as a variance-aware shrinkage estimator whose strength is governed by within-panel disagreement $D^2 = \frac{1}{N-1} \sum_i (\delta_i - \bar{\delta})^2$: substituting the unbiased estimate $\hat{\tau}^2 = D^2$ into the oracle factor $\gamma^* = \Delta_h^2 / (\Delta_h^2 + \tau^2/N)$ yields a shrinkage weight that is monotone-decreasing in D^2 , so D^2 controls the *ratio* $|\delta^*|/|\Delta|$ rather than $|\delta^*|$ in isolation. A simpler but operationally important question, complementary to the proposition, is whether DISCA’s applied correction is in fact driven by disagreement *at the scenario level*, or whether it collapses into a per-country offset. Using Qwen2.5-7B as the backbone, we log *per-scenario* $(D^2, |\delta^*|)$ pairs from the full DISCA pipeline across 5 countries (USA, JPN, DEU, VNM, ETH) and 310 MultiTP scenarios per country ($n = 1,550$ total), and compute Pearson’s correlation between $\log_{10} D^2$ and $|\delta^*|$.

Result. The two quantities are positively and significantly correlated, $r = +0.372$ ($p < 10^{-3}$, $n = 1,550$): higher scenario-level persona disagreement is associated with larger applied corrections. This is consistent with the scenario-conditioned shrinkage policy formalised by Proposition 1, with the caveat that the proposition governs the ratio $|\delta^*|/|\Delta|$ rather than the raw magnitude—a positive raw correlation arises when $|\Delta|$ also co-varies with D^2 across scenarios, which is the typical case when both are driven by the same underlying spread of persona logit gaps. At the country level (Table 15), the extremes match expectation-ETH, with the lowest mean variance (0.009), receives the smallest mean correction (0.009); USA, with the highest mean variance (0.086), receives the largest (0.015)-while the three middle countries (JPN, DEU, VNM) sit in a narrow band (0.012–0.013) with small reorderings, indicating that the correction is dominated by scenario-level disagreement rather than a strict country-level monotone mapping.

Table 15: Per-country mean inter-persona variance $\overline{D^2}$ and mean correction magnitude $\overline{|\delta^*|}$ ($n = 310$ each). Countries are ordered by mean variance. The two extremes (ETH, USA) match expectation; the three middle countries sit in a narrow band, showing the correction is set primarily at the scenario level rather than by a strict country-level monotone mapping.

Country	n	$\overline{D^2}$	$\overline{ \delta^* }$
ETH	310	0.009	0.009
JPN	310	0.052	0.013
DEU	310	0.055	0.012
VNM	310	0.070	0.013
USA	310	0.086	0.015

What this does and does not show. A positive scenario-level correlation rules out the simplest collapsed alternative—a controller that applies a country-conditioned constant offset—and is consistent with the scenario-conditioned policy that Proposition 1 prescribes. It is *not* a direct verification of the closed-form shrinkage in part (ii), which makes a quantitative claim about the ratio $|\delta^*|/|\Delta|$ that a marginal scatter cannot identify without controlling for $|\Delta|$. We therefore present this diagnostic as evidence of scenario-level conditioning, not as a tight numerical test of the oracle shrinkage factor.

A9 Baseline Implementation Details

This appendix collects implementation details for every baseline reported in the paper. §A9.1 covers the training-free baselines of Table 1 (vanilla decoding, WVS Profile Prompt, PRISM-style framing, fixed logit offset). §A9.2 covers activation steering, and §A9.3 covers ARGS / controlled decoding. §A9.4 reports additional inference-time baselines run on the twenty-country Phi-4 grid (MC-Dropout, per-country temperature/margin scaling, and a DiffPO-style mixing baseline).

755 A9.1 Other training-free baselines (Table 1)

756 **Vanilla decoding.** The model is shown the native-language scenario with no persona prefix and
757 no logit modification. We extract $\sigma(\delta(x)/T_{\text{dec}})$ and average across scenarios as in Eq. 7.

758 **WVS Profile Prompt.** A single country-conditioned system prompt is constructed by concatenat-
759 ing the ten WVS descriptors of Table 25 (population aggregate, no cohort split) into one paragraph
760 in the country’s native language; the model then answers each scenario under that prompt. This
761 collapses the four-persona panel to a single point estimate, isolating the contribution of persona
762 *diversity* on top of WVS grounding.

763 **PRISM-style prompt.** Following the PRISM cultural-framing protocol of Kirk et al. [2024], we
764 prepend a short cultural-context paragraph that names the country and asks the model to “answer as
765 a typical respondent from [country] would”. No demographic stratification, no WVS data—a pure
766 prompt-engineering control.

767 **Fixed logit offset.** A country-specific scalar Δ_c is added to the decision-token logits at inference
768 (the same Δ_c for every scenario in country c); Δ_c is fit on the synthetic 200-scenario validation pool
769 by gradient descent on JSD. This baseline is the simplest country-conditioned correction that does
770 not use the human AMCE.

771 A9.2 Activation Steering

772 For each country, a steering vector is computed as the difference in mean residual-stream activations
773 from 20 culturally contrastive prompt pairs (e.g., “Answer as someone from [country] who values
774 [WVS-high trait]” vs. “Answer as someone with the opposite values”). Vectors are extracted *per*
775 *backbone* from the transformer midpoint of the same model the baseline is evaluated on ($\lfloor L/2 \rfloor$
776 where L is the model’s layer count: layer 20 for Phi-4 ($L=40$), layer 40 for Llama-3.1-70B/3.3-
777 70B ($L=80$), layer 20 for Magistral-Small-2509 ($L=40$)) and applied at inference time by adding
778 $\alpha \cdot \mathbf{v}_{\text{steer}}$ to the residual stream, with $\alpha = 1.5$ selected via the same synthetic validation set used for
779 DISCA. There is no cross-model vector transfer; the dimension of $\mathbf{v}_{\text{steer}}$ matches the residual stream
780 of the target backbone by construction.

781 **Sensitivity analysis and fairness considerations.** Activation steering performance is known to
782 be sensitive to (i) the extraction layer, (ii) the scaling coefficient α , and (iii) the construction of con-
783 trastive pairs [Arditi et al., 2025]. We conducted a limited sensitivity check on Llama-3.1-70B (the
784 largest backbone, $L=80$): extracting from layers 32, 40, and 48 with $\alpha \in \{0.5, 1.0, 1.5, 2.0, 3.0\}$,
785 we found that layer 40 (the midpoint, used in Table 1) with $\alpha = 1.5$ yielded the lowest mean
786 JSD (0.0877); other configurations ranged from 0.089 to 0.112. We note that a more exhaustive
787 search-varying contrastive prompt design, using multiple layers simultaneously, or employing PCA-
788 based direction refinement-could improve activation steering performance. However, even the best
789 configuration performs below vanilla on this benchmark, suggesting a fundamental mismatch be-
790 tween linear activation interventions and the multi-dimensional moral calibration problem. The key
791 structural advantage of DISCA is that it operates *per-scenario* with adaptive correction magnitude,
792 whereas activation steering applies a fixed direction uniformly across all scenarios.

793 A9.3 ARGS and controlled decoding (relation to our work)

794 ARGS [Khanov et al., 2024] and controlled decoding [Mudgal et al., 2024] typically assume trained
795 reward models; a cross-cultural deployment would require a separate reward per target country. On
796 binary forced-choice MultiTP, one can instead define rewards as closed-form functionals of persona
797 logit gaps; a Prospect-Theory kernel aligned with Eq. 3 is then the natural analogue of importance-
798 weighted PT-IS. We do not run the full ARGS decoding stack in our experiments; Table 17 and §4.2
799 isolate the kernel comparison on the shared DISCA pipeline.

800 A9.4 Additional inference-time baselines (twenty-country Phi-4 grid)

801 To isolate what within-country persona disagreement adds above simpler country-conditioned cor-
802 rections, we compare DISCA against three inference-time alternatives on the same twenty-country
803 Phi-4 grid. The temperature/margin scaling baselines are deliberately granted *oracle* access to each

country’s full human AMCE for hyperparameter selection (a strict super-set of any held-out calibration regime); MC-Dropout is country-agnostic. All numbers are computed on the same evaluation pool used in the main results, so the macro means line up exactly with Table 1 (Temp. 0.513, Margin 0.506) and Table 2 (DISCA 0.346).

(6) MC-Dropout calibration. Following the moral uncertainty inflation template of Kwon et al. [2026], we enable every `torch.nn.Dropout*` module at inference (rate 0.10) and run $T=8$ stochastic forward passes per scenario, averaging the A/B probabilities. This method is deliberately country-agnostic: the same stochastic smoothing is applied for every country, so its contribution relative to vanilla isolates pure uncertainty inflation.

(7) Per-country temperature / margin scaling (oracle). For each country, we fit a scalar T_c (or additive margin m_c) by grid search against the country’s full human AMCE, minimising per-country MIS, and apply it to the same evaluation pool. This baseline therefore uses oracle supervision (the entire human AMCE vector); any held-out variant can only do worse, so the comparison is the most favourable possible to the baseline.

(8) DiffPO-binary. The spirit of DiffPO [Chen et al., 2025] adapted to binary decisions: we mix the vanilla probability with a country-conditioned target $p_{\text{target}}(c, \text{cat})$ built directly from the country’s public human AMCE ($p_{\text{aligned}} = (1 - \alpha) p_{\text{van}} + \alpha p_{\text{target}}$), with $\alpha \in [0, 1]$ fit per country against the same AMCE. Because this baseline consumes the evaluation target at inference, it is an upper bound on what any mixing of the public human AMCE with vanilla probabilities can achieve.

Takeaway. DISCA is compared against all three on the same twenty-country grid; because baselines (7) and (8) use *oracle* population-level supervision while DISCA uses only *within-country person disagreement* (never touching the human AMCE), the comparison is deliberately favourable to the baselines. The per-country results are in Table 16: DISCA wins **18 of 20** countries against MC-Dropout, **19 of 20** against per-country temperature scaling, and **18 of 20** against margin scaling. Reference implementations of all three baselines are released with the code archive accompanying this submission.

Table 16: Phi-4 (14B) per-country head-to-head: DISCA vs. three inference-time baselines on the same 20-country test split. Per-country held-out MIS ($\downarrow$, lower is better). **Bold** cells mark the best of {DISCA, MC-Dropout, T_c , m_c } per country. DiffPO-binary is omitted: as a population-target replay it scores ≈ 0 MIS by construction (each country’s prediction trivially copies the country’s human AMCE) and is not a comparable inference-time baseline.

Country	DISCA (ours)	MC-Dropout	T_c scale	Margin m_c
ARG	0.389	0.412	0.471	0.453
BGD	0.368	0.379	0.535	0.535
BRA	0.299	0.306	0.354	0.298
CHN	0.214	0.367	0.528	0.528
COL	0.438	0.389	0.501	0.469
DEU	0.255	0.376	0.580	0.580
ETH	0.485	0.526	0.655	0.655
GBR	0.319	0.407	0.576	0.576
IDN	0.274	0.452	0.542	0.542
IRN	0.530	0.538	0.382	0.353
JPN	0.195	0.371	0.327	0.327
KGZ	0.393	0.361	0.538	0.538
MEX	0.420	0.445	0.509	0.506
MMR	0.357	0.382	0.523	0.523
MYS	0.328	0.363	0.497	0.497
ROU	0.385	0.398	0.564	0.564
SRB	0.385	0.410	0.563	0.563
THA	0.313	0.342	0.494	0.494
USA	0.243	0.369	0.564	0.564
VNM	0.336	0.473	0.561	0.561
Macro	0.346	0.403	0.513	0.506
Wins (DISCA vs.)	-	18/20	19/20	18/20

A10 Hyperparameters and Sensitivity Sweeps

This appendix collects the full hyperparameter specification of DISCA together with the sensitivity sweeps that justify the choices. We organise the material into four parts: a robustness summary that

consolidates the headline checks, the canonical hyperparameter table with validation methodology, the per-axis temperature sweeps, and an extended sensitivity sweep over the four principal control knobs plus persona count and IS budget.

A10.1 Robustness Summary

Table 17 consolidates the nine independent robustness checks referenced in §4.1. Each row states one assumption behind DISCA, the test that probes it, the threshold that would falsify the assumption, and the observed value. Per-axis sweeps and full per-country tables for each row appear in the subsections below and in Appendix A13.

Table 17: Robustness suite. Each test independently probes one assumption behind DISCA; the bootstrap confidence interval on JSD is ± 0.004 . **Acronyms:** *CS-Clamp* = consensus-clamp deterministic shift (apply $\bar{\delta}$ directly, no IS); *ARGS-Unif* = ARGS-style controlled decoding [Khanov et al., 2024] with uniform reward over personas; *ARGS-WVS* = same with WVS-weighted reward; *ARGS-PT* = same with the Prospect-Theory kernel of Eq. 3.

Claim	Test	Threshold	Observed
Dataset preprocessing irrelevant	JSD range over 6 configs	< 0.010	0.0034
T_{dec} insensitivity	JSD span (Sweep A)	< 0.010	0.0074
Uniform T_{cat} insensitivity	JSD span (Sweep B)	< 0.010	0.0073
Per-category T_{cat} insensitivity	JSD span (Sweep C)	< 0.010	0.0077
PT-IS $>$ best deterministic shift	$\Delta\text{JSD}(\text{CS-Clamp} \rightarrow \text{PT-IS})$	> 0	0.003
WVS weighting alone insufficient	$\text{ARGS-WVS} \succ \text{ARGS-Unif}$	$< 8/15$	3/15
PT-IS kernel $\equiv$ ARGS-PT	ARGS-PT best/tied	$\geq 10/15$	13/15
Cross-lingual pipeline generalises	end-to-end runs	8/8	8/8

A10.2 Default Hyperparameters and Validation

Table 18: DISCA hyperparameters. Prospect Theory parameters follow Kahneman and Tversky [1979]; other defaults are validated on the synthetic pool described below.

Parameter	Symbol	Value
Persona agents	N	4
IS samples / pass	K_{half}	64 (total $2K_{\text{half}} = 128$)
Reliability scale	s	0.04
Perturbation std. dev.	σ	0.3
Cooperation weight	λ_{coop}	0.7
IS softmax temp.	η	0.5
ESS guard threshold	ρ_{eff}	0.1
PT curvature	$\alpha = \beta$	0.88
PT loss aversion	κ	2.25
Decision temp.	T_{dec}	0.5
T_{cat} (Species / Gender / other)	-	4.0 / 3.5 / 1.5
Default logit divisor	T_{logit}	3.0
ESS anchor blend	-	on (off: $\bar{\delta} + \delta_{\text{IS}}^*$ only)

Implementation. The dual-pass controller follows the equations in §3. Optional environment overrides can change K_{half} , reliability scale s , and whether the ESS-anchor blend is active; such flags must be applied before controller initialisation when running ablations.

Non-PT hyperparameters (σ , λ_{coop} , η , T_{dec} , T_{cat}) were validated on a held-out synthetic set of 200 binary moral dilemmas generated by GPT-4, covering the same six dimensions as MultiTP with novel character descriptions. Pseudo-ground-truth labels from three annotators (the authors) provided proxy AMCEs. Grid search minimised mean JSD on Qwen2.5-72B. Transferability was verified on a disjoint 50-scenario MultiTP English subset: performance within 1.2% relative JSD of oracle-tuned configuration. T_{cat} values reflect empirical logit magnitude distributions: Species and Gender produce systematically larger gaps, requiring higher temperatures. Because this tuning pool is synthetic and author-annotated, residual bias is possible; we therefore report broad robustness sweeps (Appendix A10, Appendix A13) to show conclusions are stable beyond one tuned point.

854 A10.3 Temperature Sensitivity

855 We sweep the two temperature families on the (USA, DEU, JPN) subset of MultiTP using
 856 Llama-3.1-70B. Sweep A varies $T_{\text{dec}} \in \{0.10, 0.25, 0.50, 0.75, 1.0, 2.0\}$; Sweep B replaces the
 857 per-category temperatures by a single uniform $T_{\text{cat}} \in \{0.5, 1.0, 1.5, 2.0, 3.0, 4.0, 6.0\}$; Sweep C
 858 varies the “Others” bucket $T_{\text{cat}} \in \{0.75, 1.0, 1.5, 2.0, 3.0, 4.0\}$ while holding $T_{\text{cat}}[\text{Species}]=4.0$,
 859 $T_{\text{cat}}[\text{Gender}]=3.5$. All three sweeps yield JSD spans below 0.010 (0.0074, 0.0073, 0.0077 respec-
 860 tively), and in every sweep the DISCA default ($\dagger$) is strictly *more conservative* than the empirically
 861 best value ($\Delta\text{JSD} \in \{-0.0061, -0.0021, -0.0037\}$). This rules out any concern that results have
 862 been temperature-tuned to the benchmark.

Table 19: Temperature sensitivity. Defaults ($\dagger$) are strictly worse than the best setting in every sweep, providing a conservative lower bound.

Sweep	Setting	JSD $\downarrow$	Pearson $r \uparrow$	ΔJSD
A: T_{dec}	0.10	0.0502	0.662	+0.0013
	0.25	0.0491	0.658	+0.0002
	0.50 $\dagger$	0.0489	0.630	-
	0.75	0.0476	0.636	-0.0013
	1.00	0.0466	0.641	-0.0023
	2.00 (best)	0.0428	0.627	-0.0061
	JSD span	0.0074		
B: Uniform T_{cat}	0.5	0.0521	0.642	+0.0052
	1.0	0.0492	0.673	+0.0023
	1.5 $\dagger$	0.0469	0.700	-
	2.0	0.0461	0.707	-0.0008
	3.0	0.0454	0.707	-0.0015
	4.0	0.0451	0.701	-0.0018
	6.0 (best)	0.0448	0.684	-0.0021
	JSD span	0.0073		
C: Others T_{cat}	0.75	0.0526	0.566	+0.0039
	1.00	0.0512	0.589	+0.0025
	1.50 $\dagger$	0.0487	0.634	-
	2.00	0.0475	0.660	-0.0012
	3.00	0.0458	0.693	-0.0029
	4.00 (best)	0.0450	0.706	-0.0037
	JSD span	0.0077		

863 The monotonic JSD reduction with larger T_{dec} reflects the role of T_{dec} in undoing RLHF logit com-
 864 pression: higher values widen the effective token-probability landscape so that IS perturbations exert
 865 larger directional influence. However, r peaks near $T_{\text{dec}}=1.0$ and declines at 2.0, indicating a trade-
 866 off between distributional alignment (JSD) and rank-order alignment (r); we keep $T_{\text{dec}}=0.5$ as a
 867 deliberately conservative operating point.

868 A10.4 Extended Hyperparameter Sensitivity

869 We probe the sensitivity of DISCA to the four principal hyperparameters: s (reliability gate scale),
 870 λ_{coop} (individual-vs-consensus weight), σ (IS proposal floor), and T_{cat} (uniform logit-temperature
 871 scaling). For each axis we sweep five values on a three-country panel (USA, VNM, DEU) and hold
 872 the other three axes at their defaults; the s axis directly probes how the choice of reliability scale in
 873 Eq. 5 affects downstream MIS.

Table 20: Hyperparameter sensitivity ranges on the three-country panel (Phi-4, $n=250/\text{country}$). Reported: min/max MIS across the five grid points relative to the default; MIS is stable within a narrow 0.0037–0.0090 span on every axis, so the gains in Table 2 are not the product of a particular hyperparameter choice.

Axis	Default	Grid	Min MIS	Max MIS	Δ
s	0.04	{0.01, 0.02, 0.04, 0.08, 0.16}	0.4091	0.4164	0.0073
λ_{coop}	0.70	{0.30, 0.50, 0.70, 0.85, 0.95}	0.4085	0.4173	0.0088
σ	0.30	{0.10, 0.20, 0.30, 0.45, 0.60}	0.4084	0.4122	0.0037
T_{cat} scale	$1.0\times$	{ $0.33\times$, $0.67\times$, $1.0\times$, $1.33\times$, $1.67\times$ }	0.4060	0.4149	0.0090

874 **Why these four axes.** s controls the steepness of the reliability gate (Eq. 5); λ_{coop} trades off
875 individual-persona PT utility against the consensus utility; σ floors the IS proposal scale so the
876 $N=4$ empirical std cannot collapse to zero; and T_{cat} scales the whole family of per-category deci-
877 sion temperatures (preserving relative balance)-a global sharpening / flattening knob on the decision
878 logits. Per-category temperatures are stress-tested separately in the temperature sensitivity sweep
879 above.

880 **Persona count, IS budget, and per-category vs. global T_{cat} .** We additionally swept three op-
881 erational knobs on the same USA/VNM/DEU panel. **Persona count** ($N \in \{2, 3, 4, 5, 6\}$, Ta-
882 ble 21): $N=4$ is optimal at macro MIS 0.4051, with degradation at both smaller and larger pan-
883 els (0.4252 at $N=2$, 0.4179 at $N=6$), supporting the four-persona design. **IS budget** ($K_{\text{half}} \in$
884 $\{8, 16, 32, 64, 128, 192\}$, Table 22): performance is flat across an order of magnitude (best 0.4092
885 at $K=16$; 0.4103 at $K=64$), indicating the default $K_{\text{half}}=64$ is in the stable region rather than over-
886 tuned. **Global vs. per-category T_{cat}** (Table 23): the per-category default outperforms *every* global
887 schedule we tested (0.4098 vs. all seven global values in $\{1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0\}$, which land
888 in 0.4119–0.4138). The advantage is small in absolute terms but consistent.

Table 21: Persona-count sweep on the USA/VNM/DEU panel (Phi-4, $n=250/\text{country}$). The four-persona default minimises macro MIS; larger panels add latency without alignment benefit.

N	Macro MIS $\downarrow$	Std (across countries)	Flip rate	sec/scenario
2	0.4252	0.1045	0.275	0.089
3	0.4194	0.0945	0.273	0.100
4	0.4051	0.0829	0.276	0.111
5	0.4107	0.0745	0.276	0.120
6	0.4179	0.0919	0.245	0.130

Table 22: IS-budget sweep (K_{half} samples per pass; total per scenario is $K_{\text{total}}=2K_{\text{half}}$). Macro MIS is flat across an order of magnitude. The mean reliability weight $\bar{r}$ (gate column) grows monotonically with K , confirming the gate becomes more confident as the IS proposal cloud densifies.

K_{half}	K_{total}	Macro MIS $\downarrow$	Flip rate	$\bar{r}$ (gate)	sec/scen
8	16	0.4114	0.284	0.684	0.121
16	32	0.4092	0.271	0.736	0.114
32	64	0.4104	0.278	0.766	0.114
64	128	0.4103	0.280	0.793	0.113
128	256	0.4147	0.274	0.814	0.116
192	384	0.4126	0.273	0.827	0.112

Table 23: Global vs. per-category T_{cat} on the USA/VNM/DEU panel. The per-category default (decision 4.0, social 3.5, non-social 1.5) outperforms every global value we tested. The gap is small (≤ 0.004) but uniformly directional, supporting category-specific decision temperatures over a single shared scalar.

T_{cat} schedule	Macro MIS ↓	Std	Flip rate
per-category default (4.0 / 3.5 / 1.5)	0.4098	0.0881	0.278
global $T_{\text{cat}} = 1.00$	0.4120	0.0921	0.292
global $T_{\text{cat}} = 1.50$	0.4119	0.0838	0.304
global $T_{\text{cat}} = 2.00$	0.4138	0.0778	0.301
global $T_{\text{cat}} = 2.50$	0.4134	0.0774	0.296
global $T_{\text{cat}} = 3.00$	0.4127	0.0749	0.303
global $T_{\text{cat}} = 3.50$	0.4138	0.0689	0.306
global $T_{\text{cat}} = 4.00$	0.4128	0.0734	0.302

Per-persona utility floor (minority safeguard). The released controller optionally applies a per-persona utility floor that limits how far any single persona’s post-correction utility can fall below vanilla. Sweeping $f \in \{0.0, 0.5, 1.0, 2.0\}$ on the same USA/VNM/DEU panel (Phi-4, $n=250/\text{country}$) keeps macro MIS in $[0.4046, 0.4157]$ —well inside the macro-MIS bootstrap noise floor—so the floor is a stability knob rather than a load-bearing hyperparameter.

Table 24: Per-persona utility floor sweep (USA / VNM / DEU). Macro MIS is flat across two orders of magnitude of the floor parameter.

Floor f	Macro MIS ↓	Std (across countries)
0.0 (off)	0.4083	0.0717
0.5	0.4046	0.0696
1.0	0.4066	0.0751
2.0	0.4157	0.0684

A11 Persona Construction Details

This appendix specifies how the four cultural personas per country are constructed and provides empirical support that the construction is grounded in human survey data rather than authored ad hoc. §A11.1 reproduces two personas verbatim from `generate_wvs_persona()`; §A11.2 documents the WVS Wave 7 processing pipeline; §A11.3 links the ten WVS cultural dimensions to the six MultiTP moral attributes both thematically and empirically; and §A11.5 shows that DISCA is robust to the choice of the fourth (anchor) persona.

A11.1 Persona Prompts

Each country yields four personas: three age-cohort personas (*young*, *middle*, *older*) generated from the country-and-cohort-specific WVS-7 means, plus a population-wide aggregate that serves as a demographic anchor. Below we provide the full prompt sets for the United States and Vietnam, reproduced verbatim from the WVS-driven persona generation pipeline with the aggregate fourth persona setting.

United States (all four personas, verbatim)

Young cohort.

You are a young adult from the United States, in your 20s and early 30s. Your worldview is shaped by the cultural values prevalent in your community. On matters of faith you are moderately religious. On raising children you are firmly oriented toward independence and imagination. On contested moral choices you are morally conservative on contested issues. In your dealings with strangers you have a guarded attitude toward strangers. Civically you are an active political participant who signs petitions, joins boycotts and takes part in lawful demonstrations. You are moderately proud of your country. Overall you are rather happy with your life. On the role of women in society you are moderately egalitarian

917 *on gender roles. In what you prioritise in life you are leaning materialist, prioritising*
 918 *economic and physical security. Toward people unlike yourself you are highly tolerant of*
 919 *outgroups such as immigrants, minorities and people with different lifestyles. When you*
 920 *face a moral dilemma, you weigh the choices through this set of values and answer in a*
 921 *way that is consistent with the worldview above.*

922 **Middle-aged cohort.**

923 *You are a middle-aged adult from the United States, in your 40s or 50s. Your worldview*
 924 *is shaped by the cultural values prevalent in your community. On matters of faith you are*
 925 *moderately religious. On raising children you are firmly oriented toward independence*
 926 *and imagination. On contested moral choices you are strictly opposed to such contested*
 927 *moral acts. In your dealings with strangers you have a guarded attitude toward strangers.*
 928 *Civically you are an active political participant who signs petitions, joins boycotts and*
 929 *takes part in lawful demonstrations. You are moderately proud of your country. Overall*
 930 *you are rather happy with your life. On the role of women in society you are moderately*
 931 *egalitarian on gender roles. In what you prioritise in life you are leaning materialist,*
 932 *prioritising economic and physical security. Toward people unlike yourself you are highly*
 933 *tolerant of outgroups such as immigrants, minorities and people with different lifestyles.*
 934 *When you face a moral dilemma, you weigh the choices through this set of values and*
 935 *answer in a way that is consistent with the worldview above.*

936 **Older cohort.**

937 *You are a senior citizen from the United States, over 60 years old. Your worldview is*
 938 *shaped by the cultural values prevalent in your community. On matters of faith you are*
 939 *moderately religious. On raising children you are firmly oriented toward independence*
 940 *and imagination. On contested moral choices you are strictly opposed to such contested*
 941 *moral acts. In your dealings with strangers you have a guarded attitude toward strangers.*
 942 *Civically you are an active political participant who signs petitions, joins boycotts and*
 943 *takes part in lawful demonstrations. You are intensely proud of your country. Overall*
 944 *you are rather happy with your life. On the role of women in society you are moderately*
 945 *egalitarian on gender roles. In what you prioritise in life you are leaning materialist,*
 946 *prioritising economic and physical security. Toward people unlike yourself you are highly*
 947 *tolerant of outgroups such as immigrants, minorities and people with different lifestyles.*
 948 *When you face a moral dilemma, you weigh the choices through this set of values and*
 949 *answer in a way that is consistent with the worldview above.*

950 **Aggregate cohort (all ages).**

951 *You are a adult citizen from the United States. Your worldview is shaped by the cultural*
 952 *values prevalent in your community. On matters of faith you are moderately religious.*
 953 *On raising children you are firmly oriented toward independence and imagination. On*
 954 *contested moral choices you are strictly opposed to such contested moral acts. In your*
 955 *dealings with strangers you have a guarded attitude toward strangers. Civically you are*
 956 *an active political participant who signs petitions, joins boycotts and takes part in lawful*
 957 *demonstrations. You are moderately proud of your country. Overall you are rather happy*
 958 *with your life. On the role of women in society you are moderately egalitarian on gender*
 959 *roles. In what you prioritise in life you are leaning materialist, prioritising economic and*
 960 *physical security. Toward people unlike yourself you are highly tolerant of outgroups such*
 961 *as immigrants, minorities and people with different lifestyles. When you face a moral*
 962 *dilemma, you weigh the choices through this set of values and answer in a way that is*
 963 *consistent with the worldview above.*

964 **Vietnam (all four personas, Vietnamese, verbatim)**

965 **Young cohort.**

966 *Bạn là một thanh niên đến từ Việt Nam, ở độ tuổi 20 đến đầu 30. Thế giới quan của bạn*
 967 *được định hình bởi các giá trị văn hóa phổ biến trong cộng đồng của bạn. Về vấn đề đức*
 968 *tin, bạn khá thế tục. Trong việc nuôi dạy con cái, bạn nghiêng về sự vâng lời và đức tin tôn*
 969 *giáo. Về các lựa chọn đạo đức gây tranh cãi, bạn bảo thủ về mặt đạo đức trong các vấn đề*
 970 *gây tranh cãi. Trong giao tiếp với người lạ, bạn có sự ngờ vực sâu sắc đối với người khác.*
 971 *Về mặt công dân, bạn là một người tham gia chính trị thụ động. Bạn rất tự hào về đất nước*
 972 *của bạn. Nhìn chung, bạn rất hài lòng với cuộc sống của bạn. Về vai trò của phụ nữ trong*
 973 *xã hội, bạn khá truyền thống về vai trò giới. Về những gì bạn ưu tiên trong cuộc sống, bạn*
 974 *có xu hướng hậu vật chất. Đối với những người khác bạn, bạn có phần thiếu khoan dung*

với các nhóm bên ngoài. Khi bạn đối mặt với một tình huống khó xử về đạo đức, bạn cân nhắc các lựa chọn thông qua tập hợp giá trị này và trả lời theo cách phù hợp với thể giới quan đã nêu ở trên.

978 **Middle-aged cohort.**

979 *Bạn là một người trung niên đến từ Việt Nam, ở độ tuổi 40 hoặc 50. Thể giới quan của bạn*
980 *được định hình bởi các giá trị văn hóa phổ biến trong cộng đồng của bạn. Về vấn đề đức*
981 *tin, bạn khá thể tục. Trong việc nuôi dạy con cái, bạn nghiêng về sự vâng lời và đức tin tôn*
982 *giáo. Về các lựa chọn đạo đức gây tranh cãi, bạn bảo thủ về mặt đạo đức trong các vấn đề*
983 *gây tranh cãi. Trong giao tiếp với người lạ, bạn có thái độ dè dặt với người lạ. Về mặt công*
984 *dân, bạn là một người thành thạo tham gia chính trị. Bạn rất tự hào về đất nước của bạn.*
985 *Nhìn chung, bạn rất hài lòng với cuộc sống của bạn. Về vai trò của phụ nữ trong xã hội,*
986 *bạn khá truyền thống về vai trò giới. Về những gì bạn ưu tiên trong cuộc sống, bạn có xu*
987 *hướng hậu vật chất. Đối với những người khác bạn, bạn có phần thiếu khoan dung với*
988 *nhóm bên ngoài. Khi bạn đối mặt với một tình huống khó xử về đạo đức, bạn cân nhắc các*
989 *lựa chọn thông qua tập hợp giá trị này và trả lời theo cách phù hợp với thể giới quan đã*
990 *nêu ở trên.*

991 **Older cohort.**

992 *Bạn là một người cao tuổi đến từ Việt Nam, trên 60 tuổi. Thể giới quan của bạn được định*
993 *hình bởi các giá trị văn hóa phổ biến trong cộng đồng của bạn. Về vấn đề đức tin, bạn có*
994 *niềm tin tôn giáo ở mức trung bình. Trong việc nuôi dạy con cái, bạn nghiêng về sự vâng*
995 *lời và đức tin tôn giáo. Về các lựa chọn đạo đức gây tranh cãi, bạn bảo thủ về mặt đạo đức*
996 *trong các vấn đề gây tranh cãi. Trong giao tiếp với người lạ, bạn có thái độ dè dặt với người*
997 *lạ. Về mặt công dân, bạn là một người thành thạo tham gia chính trị. Bạn rất tự hào về*
998 *đất nước của bạn. Nhìn chung, bạn rất hài lòng với cuộc sống của bạn. Về vai trò của phụ*
999 *nữ trong xã hội, bạn khá truyền thống về vai trò giới. Về những gì bạn ưu tiên trong cuộc*
1000 *sống, bạn có xu hướng hậu vật chất. Đối với những người khác bạn, bạn rất thiếu khoan*
1001 *dung với các nhóm bên ngoài. Khi bạn đối mặt với một tình huống khó xử về đạo đức, bạn*
1002 *cân nhắc các lựa chọn thông qua tập hợp giá trị này và trả lời theo cách phù hợp với thể*
1003 *giới quan đã nêu ở trên.*

1004 **Aggregate cohort (all ages).**

1005 *Bạn là một công dân trưởng thành đến từ Việt Nam. Thể giới quan của bạn được định hình*
1006 *bởi các giá trị văn hóa phổ biến trong cộng đồng của bạn. Về vấn đề đức tin, bạn khá thể*
1007 *tục. Trong việc nuôi dạy con cái, bạn nghiêng về sự vâng lời và đức tin tôn giáo. Về các*
1008 *lựa chọn đạo đức gây tranh cãi, bạn bảo thủ về mặt đạo đức trong các vấn đề gây tranh*
1009 *cãi. Trong giao tiếp với người lạ, bạn có thái độ dè dặt với người lạ. Về mặt công dân, bạn*
1010 *là một người thành thạo tham gia chính trị. Bạn rất tự hào về đất nước của bạn. Nhìn*
1011 *chung, bạn rất hài lòng với cuộc sống của bạn. Về vai trò của phụ nữ trong xã hội, bạn khá*
1012 *truyền thống về vai trò giới. Về những gì bạn ưu tiên trong cuộc sống, bạn có xu hướng hậu*
1013 *vật chất. Đối với những người khác bạn, bạn có phần thiếu khoan dung với các nhóm bên*
1014 *ngoài. Khi bạn đối mặt với một tình huống khó xử về đạo đức, bạn cân nhắc các lựa chọn*
1015 *thông qua tập hợp giá trị này và trả lời theo cách phù hợp với thể giới quan đã nêu ở trên.*

1016 **Aggregate (fourth) persona.** The fourth agent shown above is constructed from the country's
1017 population-wide WVS profile, pooling respondents across age cohorts. It serves as a demographic
1018 anchor: by suppressing cohort-specific noise it ensures that the ensemble's consensus target δ re-
1019 flects the broadest available empirical signal rather than a single generational viewpoint. §A11.5
1020 contrasts this choice with a country-invariant utilitarian anchor and shows that DISCA's macro-MIS
1021 is insensitive to the substitution.

1022 **Faithfulness disclaimer.** The descriptors above are read directly from the country's WVS-7 means
1023 through deterministic quartile cuts (Table 25); they are not hand-selected to match cultural stereo-
1024 types. As a consequence, individual cohorts can present empirically grounded but locally counter-
1025 intuitive combinations (for example, a U.S. young-adult that is *leaning materialist* and *morally*
1026 *conservative on contested issues*), reflecting actual WVS-7 distributions rather than authorial bias.

1027 **A11.2 WVS Data Processing Pipeline**

1028 We process WVS Wave 7 individual-level microdata following the ten-variable cultural-value
1029 scheme of Greco et al. [2026]. The pipeline is fully deterministic: given the inverted WVS-7 CSV

and the country code, it returns the four personas used at inference time without any tunable hyper-parameters.

1. **Variable extraction:** For each respondent, we extract 10 value dimensions from WVS items (Table 25). The “inverted” WVS-7 CSV (suffix P) pre-flips Likert items so higher values consistently indicate the positive pole. Items without the P suffix (Q152, Q153, Q177–Q182) use original WVS coding.
2. **Age cohort assignment:** Using birth year (Q261) and survey year (A_YEAR), we compute age and assign to: *young* (<36), *middle* (36–55), *older* (>55). Respondents with birth year before 1900 or after 2010, or survey year before 2015, are excluded. Negative WVS codes (−1, −2, −4, −5: refusal / don’t know) are dropped; zero is retained for binary 0/1 items.
3. **Cohort aggregation:** For each country and age cohort, each dimension’s score is the un-weighted mean of all valid (respondent, item) values associated with that dimension - i.e., for multi-item dimensions (e.g., Q177–Q182 for moral acceptability) every respondent contributes one value per item to a single pooled mean. For single-item dimensions this reduces to the standard across-respondent mean.
4. **Normalisation and descriptor mapping:** Raw dimension means are normalised to $[0, 1]$ via $(\text{value} - \text{lo}) / (\text{hi} - \text{lo})$ where (lo, hi) is the WVS scale range, then directionally flipped so higher = positive pole. Four-level descriptors are assigned by quartile cuts: ≥ 0.75 (strong positive), ≥ 0.50 , ≥ 0.25 , < 0.25 (strong negative). See Table 25.
5. **Fallback:** For countries with insufficient WVS coverage (e.g., Saudi Arabia), we use manually authored native-language personas grounded in area-studies literature, covering the same demographic structure (3 age cohorts + 1 population-wide aggregate).

Table 25: The 10 WVS cultural-value dimensions used for persona generation, following Greco et al. [2026]. All dimensions use uniform quartile cuts on the $[0, 1]$ -normalised score: ≥ 0.75 (descriptor 1, strongest positive pole), ≥ 0.50 (2), ≥ 0.25 (3), < 0.25 (4, strongest negative pole). The “direction” column indicates whether higher raw WVS values align with (+) or oppose (−) the positive pole.

Dimension	WVS items	Range	Dir.	Descriptors (≥ 0.75 / ≥ 0.50 / ≥ 0.25 / < 0.25)
Religiosity	Q6P	1–4	+	deeply / moderately religious / somewhat / highly secular
Child-rearing	Q17P	0–1	−	independence-imagination ↔ obedience-faith
Moral accept.	Q177–Q182	1–10	+	very permissive ↔ strictly opposed
Social trust	Q57P	1–2	+	very high trust ↔ deep distrust
Polit. particip.	Q199P, Q200P	1–3	+	active participant ↔ non-participant
National pride	Q254P	1–4	+	very proud ↔ not proud
Happiness	Q46P	1–4	+	very happy ↔ not happy
Gender equality	Q29P, Q30P, Q31P, Q33P	1–4	−	strongly egalitarian ↔ traditional roles
Materialism	Q152, Q153	1–3	+	post-materialist ↔ materialist
Tolerance	Q19P–Q23P	0–1	−	very tolerant ↔ intolerant

A11.3 WVS-to-Trolley Dimension Linkage

The linkage between the ten WVS cultural dimensions and the six MultiTP moral dimensions operates through two complementary mechanisms.

Direct thematic overlap. WVS *gender equality* maps onto the Gender dimension-societies scoring higher exhibit weaker female-sparing preferences [Awad et al., 2018]; *religiosity* correlates with Species through religious anthropocentrism and human exceptionalism; and *moral acceptability* relates to tolerance for trade-offs along the Fitness and Social Value dimensions.

Indirect modulation. Dimensions such as *social trust*, *child-rearing values*, and *materialism orientation* reshape the moral *frame* the persona adopts: post-materialist personas favour utilitarian calculi, while obedience-oriented child-rearing values up-weight hierarchical social roles (affecting Age and Social Value). Importantly, DISCA does not require an exact causal mapping between WVS and MultiTP dimensions; the ten-dimensional persona ensemble produces diversity in logit-space gaps, and the importance-sampling stage selects the correction that balances collective utility.

1065 **Empirical validation of the WVS→AMCE pathway.** Table 26 regresses each of the six MultiTP
1066 human AMCE dimensions on the 10 WVS cultural features across the 20-country panel (standard-
1067 ised OLS). The 10-feature WVS profile explains $R^2 \in [0.55, 0.69]$ of human AMCE variance, with
1068 $|\beta| > 0.3$ in 31 of 60 cells. This is direct evidence that WVS-grounded persona construction is not
1069 arbitrary: the cultural axes our personas inherit correlate strongly with the moral preferences DISCA
1070 must reproduce.

Table 26: WVS features as predictors of human AMCE dimensions. OLS with standardised predictors. Coefficients are standardised (β); bold = $|\beta| > 0.3$. R^2_{adj} measures how well the 10 WVS cultural dimensions explain each moral preference dimension across the 20-country panel.

AMCE dim	relig	child	moral	socia	polit	natio	happi	gende	mater	toler	R^2	R^2_{adj}
Species	+0.47	+0.12	-0.65	+0.60	+0.08	+0.16	-0.02	-0.74	-0.30	+0.40	0.57	0.09
Gender	+0.69	+0.37	-0.18	+0.41	-0.44	-0.39	+0.71	-0.81	-0.20	+0.06	0.66	0.29
Age	-0.80	-0.09	+0.61	-0.74	+0.17	+0.26	-0.27	+0.62	+0.49	-0.37	0.55	0.05
Fitness	+0.28	-0.23	-0.16	-0.12	-0.17	+0.43	+0.09	-1.46	-0.21	+0.53	0.60	0.15
SocialValue	-0.27	-0.30	+0.40	-0.40	-0.37	+0.58	+0.17	-0.31	-0.12	-0.15	0.69	0.35
Utilitarianism	+1.11	+0.39	+0.07	+0.59	-0.21	-0.01	+0.13	-0.76	+0.42	+0.25	0.66	0.29

1071 Table 27 closes the loop with a causal leave-one-WVS-dim-out probe: for each WVS dimension
1072 d , we re-run DISCA with d removed from every persona and measure the change in per-MultiTP-
1073 dim absolute error (positive cell = error rises when the dimension is removed = that dimension is
1074 load-bearing for the corresponding moral attribute). The three largest positive couplings are *moral*
1075 *acceptability* → Social Value (+3.97 pp), *social trust* → Social Value (+2.81 pp), and *child-rearing*
1076 *values* → Utilitarianism (+2.41 pp). Cells with the top-3 largest positive couplings per row are
1077 bolded; rows with fewer than three positive cells are bolded only on those positive cells. Negative
1078 cells are also informative: *tolerance of diversity* → Social Value (−7.76) and *gender equality* →
1079 Gender (−3.42) indicate that, for those moral attributes, the corresponding WVS descriptor biases
1080 the persona ensemble *away* from the human target, so removing the descriptor improves accuracy.
1081 Crucially, DISCA does not require manual feature selection to handle these noisy dimensions: the
1082 loss-averse importance-sampling stage (Eq. 3) automatically down-weights candidate perturbations
1083 that worsen any persona’s alignment, effectively acting as a built-in noise filter. When a WVS
1084 dimension introduces bias for a particular moral attribute, the PT-IS utility penalises the resulting
1085 correction asymmetrically ($\kappa=2.25\times$ for losses vs. gains), ensuring that noisy persona signals are
1086 suppressed rather than propagated. This self-correcting property is why the method remains robust
1087 despite imperfect WVS-to-trolley mappings—a design-level answer to the concern that the linkage is
1088 “indirect modulation” rather than causal.

Table 27: WVS-dim × MultiTP-dim causal impact matrix. Cells are the mean increase in per-dim AMCE error (pp) when the WVS dimension is dropped from every persona, macro-averaged across the country panel. **Bold** cells mark the top-3 largest positive couplings per row (load-bearing WVS → MultiTP links).

WVS dim	Species_Humans	Gender_Female	Age_Young	Fitness_Fit	SocialValue_High	Utilitarianism_More
religiosity	+0.66	+1.65	-1.66	+0.58	+1.02	-0.37
child rearing	-1.82	+1.29	+0.86	+1.30	-0.33	+2.41
moral acceptability	+1.13	-0.02	-1.56	-0.34	+3.97	-1.69
social trust	+0.11	+1.43	-0.85	+0.51	+2.81	-1.88
political participation	-1.24	-0.82	-0.63	-0.36	-0.16	+0.86
national pride	-0.37	-0.10	-0.28	+0.52	+1.19	+1.21
happiness	-0.23	-0.88	+0.13	-1.69	-0.11	+1.72
gender equality	+0.41	-3.42	-2.90	-1.11	-1.87	+0.62
materialism orientation	+1.35	-1.26	+0.90	+0.04	-0.36	+2.15
tolerance diversity	-0.76	-3.52	+1.95	-0.59	-7.76	+1.54

1089 A11.4 Macro-level WVS dimension importance

1090 The per-MultiTP-dim impact matrix above (Table 27) tells us *which* moral attribute each WVS
1091 descriptor influences. A coarser but operationally useful question is which WVS dimensions are
1092 load-bearing for macro alignment, taken as a whole. We answer this with a leave-one-WVS-dim-
1093 out probe: for each of the ten descriptors, we re-run DISCA with that descriptor removed from
1094 every persona and record the change in macro MIS, averaged across USA, JPN, and VNM as a
1095 representative panel.

1096 The result is a concentrated importance profile: religiosity, gender equality, and moral acceptability
1097 each raise MIS by more than 0.03 when removed; national pride and happiness barely register
1098 (< 0.01). The implication for persona engineering is that the ten-dimensional WVS profile can be
1099 compressed substantially while preserving most of the steering signal, which is consistent with the
1100 IS stage’s built-in down-weighting of noisy descriptors discussed in App. A11.3.

Table 28: WVS dimension dropout (3-country average; higher Δ MIS means more load-bearing). Dimension names follow the canonical WVS-7 list of Table 25 and Table 27.

Dropped WVS dimension	Δ MIS (avg)
religiosity	+0.046
gender equality	+0.041
moral acceptability	+0.034
social trust	+0.022
child rearing	+0.019
political participation	+0.017
materialism orientation	+0.013
tolerance diversity	+0.011
national pride	+0.007
happiness	+0.005

1101 A11.5 Sensitivity to the Fourth Persona

1102 **Setup.** Appendix A11 pairs three age-cohort personas with a population-wide *aggregate* WVS
1103 profile as the fourth agent. A natural alternative is a country-invariant *utilitarian* anchor that substi-
1104 tutes a neutral moral stance for the aggregate. To confirm that this choice is not load-bearing, we ran
1105 a head-to-head comparison of the two variants on the twenty-country Phi-4 grid.

1106 **Result.** The utilitarian-anchor variant improves macro MIS from 0.4252 to 0.4049 ($\Delta =$
1107 -0.0203), a marginal gain that does not change any qualitative conclusion of the paper.

1108 **Decision.** We retain the aggregate variant as the default to preserve strict country-grounding of
1109 every persona and to avoid importing a globally fixed moral prior into country-conditioned steering.
1110 The utilitarian variant remains available as an opt-in for practitioners who prefer a neutral anchor
1111 (`build_country_personas(..., fourth='utilitarian')`).

1112 A12 AMCE Estimation Details

1113 **Uniform treatment of all dimensions.** All six dimensions-Species, Gender, Age, Fitness, Social
1114 Value, and Utilitarianism-are computed identically. For each dimension, we compute the empirical
1115 mean of $p_{\text{spare}}(\mathbf{x})$ across all scenarios in that dimension (Eq. 7). For the five binary dimensions, each
1116 scenario presents a forced choice between two character groups differing on exactly one attribute.
1117 For Utilitarianism, scenarios vary the number of lives on each side; the binary indicator is whether
1118 the larger group (“More”) or the smaller group (“Less”) is spared. This uniform mean-based es-
1119 timation is consistent with the `compute_ACME` method in the MultiTP codebase [Jin et al., 2025],
1120 which fits a no-intercept linear regression with a binary group indicator-mathematically equivalent
1121 to taking the mean of the saving probability for the preferred group.

1122 **Human AMCE scale.** The MultiTP ground-truth AMCEs in the MultiTP-released country-
1123 specific AMCE table is on a $[-1, 1]$ scale. We convert to $[0, 100]$ via $(1 + \text{AMCE}_{\text{raw}})/2 \times 100$,
1124 where 50% represents no preference and 100% represents maximal preference for the “preferred”
1125 group. Model AMCEs are computed on the same scale.

1126 **Two reporting scales (resolving an apparent contradiction).** The PCA visualisation in Figure 3
1127 reports ℓ_2 distances on the $[0, 100]$ percentage-point scale (e.g., $83.3 \rightarrow 64.6$), while every numerical
1128 table in the main text and appendix (Tables 1, 2, 4, 6, etc.) reports MIS on the $[0, 1]$ proportional
1129 scale (i.e., directly in p_{spare} space, with values in $[0.2, 1.0]$). The two scales differ by a factor of
1130 100. We use the proportional scale in tables because it preserves three significant figures within
1131 a compact column width; we use the percentage-point scale only in the geometric PCA figure to
1132 match the original Moral Machine release convention.

1133 **Consistency.** Awad et al. [2018] show that marginal averages are unbiased estimators under the
1134 balanced randomisation design of the Moral Machine, making our estimation valid and reproducible.

1135 **Cross-lingual A/B token elicitation.** Eq. 7 requires a clean A vs. B logit gap on every scenario,
1136 which in turn requires that the decision tokens be elicited consistently across model families and
1137 across the 25+ languages used in the 20-country panel. Each scenario is rendered entirely in the tar-
1138 get country’s native language: this includes (i) scenario framing and context sentences (4 paraphrase
1139 variants per language), (ii) all 22 character types with singular/plural forms, (iii) lane labels and
1140 group identifiers in native script, and (iv) closing questions with language-appropriate conjunctions
1141 and grammar. The prompt frame wraps the scenario in a native-language instruction that explicitly
1142 requests an English answer token, following the pattern: “[moral dilemma preamble in native lan-
1143 guage] \n {scenario} \n [instruction to answer A or B in English]”. This keeps moral reasoning in
1144 the culturally native linguistic frame while maintaining a single decision interface across the panel.

1145 **Token verification.** Token IDs are verified per model at initialisation by encoding the upper-
1146 case strings "A" and "B" with `add_special_tokens=False` and asserting single-token mappings:
1147 Qwen2.5: A=22397, B=11572; Llama-3.1: A=31266, B=9268. Logits for these two positions are
1148 extracted from the last-position output tensor; no generation sampling is performed. The model uses
1149 its built-in `chat_template` for correct role formatting across model families.

1150 A13 Dataset Preprocessing and Sensitivity

1151 **Pipeline.** Each MultiTP CSV is processed in five steps: (i) deduplication keeps only the canonical
1152 paraphrase (`which_paraphrase=0`); (ii) the Utilitarianism *quality filter* drops scenarios where
1153 the two groups have identical size *and* all characters belong to the *quality-attribute* role set
1154 {PREGNANT, WOMAN, LARGE WOMAN}, which would otherwise inject a non-numerosity signal
1155 into a dimension whose AMCE is by construction the count-difference effect; (iii) per-category cap-
1156 ping retains at most 80 scenarios per dimension; (iv) under-represented categories are oversampled
1157 with replacement to a minimum of 36 scenarios; and (v) the resulting pool is shuffled with a fixed
1158 seed (42).

1159 **Reproducible left/right assignment.** When the MultiTP `paraphrase_choice` field does not un-
1160 ambiguously determine which group is rendered on the left vs. right, we deterministically fall back
1161 to a SHA-256 hash of the tuple (sub_1, sub_2, g_1, g_2) modulo 2, so the same scenario receives the same
1162 ordering across runs and machines. The fallback rate is logged per country and stays below 5% on
1163 every country in the 20-country panel.

1164 Up-sampling duplicates scenarios without modification, preserving AMCE signal while reducing
1165 variance. Category capping prevents dimension dominance. Side balancing enables effective debi-
1166 asing. To evaluate whether these choices alter the effective evaluation distribution, we compare *six*
1167 preprocessing configurations on 10 representative countries (USA, DEU, CHN, JPN, BRA, VNM,
1168 GBR, KOR, RUS, NGA) using Llama-3.1-70B.

Table 29: Per-country dataset sensitivity. JSD range = 0.0034 (aggregate) and < 0.010 in 9/10 countries; the only exception is RUS at 0.013, driven by augmentation sensitivity in the rare Species/Utilitarianism categories. D0 is retained as default.

Country	D0-Default	D1-NoAug	D2-Cap40	D3-Cap120	D4-NoFlip	D5-Strict
USA	0.0655	0.0674	0.0625	0.0644	0.0643	0.0684
DEU	0.0465	0.0515	0.0492	0.0467	0.0444	0.0484
CHN	0.0393	0.0363	0.0385	0.0383	0.0409	0.0398
JPN	0.0304	0.0291	0.0319	0.0296	0.0273	0.0320
BRA	0.0492	0.0551	0.0479	0.0510	0.0520	0.0480
VNM	0.0592	0.0574	0.0631	0.0585	0.0578	0.0626
GBR	0.0614	0.0667	0.0608	0.0618	0.0597	0.0633
KOR	0.0371	0.0354	0.0378	0.0357	0.0349	0.0382
RUS	0.0383	0.0489	0.0389	0.0374	0.0383	0.0357
NGA	0.0522	0.0568	0.0506	0.0525	0.0514	0.0538
Mean	0.0479	0.0505	0.0481	0.0476	0.0471	0.0490
Δ vs D0	-	+0.0025	+0.0002	-0.0003	-0.0008	+0.0011

D1-NoAug shows the largest degradation (+0.0025), justifying synthetic augmentation of the under-represented Species and Utilitarianism categories. D4-NoFlip is marginally lower than D0, indicating that group-side randomisation is conservative (it adds noise) rather than harmful; we retain it because it preserves positional-debiasing integrity. The 0.0034 aggregate range is well below the ± 0.004 bootstrap CI, so reported results are not confounded by preprocessing choices.

A14 Trigger Mechanism and Latency

DISCA runs dual-pass IS on every scenario; inter-persona variance is logged but does not gate sampling. The *flip rate* (IS reversals vs. consensus) is 29.0% in the cross-backbone ablation suite (Table 4), indicating confidence modulation rather than wholesale decision reversal.

A compact round-2 latency benchmark on H100 shows similar overhead across proposal budgets: 0.086s/scenario at $K=64$, 0.083s at $K=128$, and 0.083s at $K=256$, compared with 0.023s for vanilla decoding ($3.57\text{--}3.72\times$ overhead).

Table 30: Inference latency benchmark (Phi-4, H100).

Method	K	sec/scenario	Overhead
DISCA	64	0.086	$3.72\times$
DISCA	128	0.083	$3.59\times$
DISCA	256	0.083	$3.57\times$
vanilla	—	0.023	$1.00\times$

Cost-vs-quality frontier across backbones. Figure 5 plots the *deployed-cost* side of the Phi-4 vs-Llama narrative: at 350 ms per scenario, Phi-4 with DISCA reaches MIS = 0.346, while Llama-3.3-70B (with DISCA) takes 1414 ms—roughly $4\times$ slower—and still ends at MIS = 0.668. Phi-4 thus Pareto-dominates the 70B baseline on *both* axes simultaneously: better alignment at lower latency. Among the headline seven, only Magistral-24B sits on a comparable quality plateau, at almost double Phi-4’s per-scenario cost.




Figure 5: **Cost-vs-quality frontier on the headline 7 models.** Per-scenario DISCA latency (log scale) vs. mean DISCA MIS. Marker size is proportional to parameter count; color encodes Δ MIS (greener = larger DISCA gain). Phi-4 (14B, $\Delta = +0.108$) lies bottom-left: Pareto-dominant over Llama-3.3-70B in both latency and alignment.

A15 Relationship to Persona-Dependent LLM Alignment

Kim et al. [2025] is the closest prior work to DISCA in studying persona effects on LLM moral decisions. Their study and ours share the Moral Machine framework and AMCE-based evaluation, but differ fundamentally in goal, method, and scope. Table 31 summarises the key distinctions.

Table 31: Comparison between Kim et al. [2025] and DISCA.

Dimension	Kim et al. (2025)	DISCA (this work)
Goal	Measure persona sensitivity	Correct cultural misalignment
Personas	7 sociodemographic categories	4 WVS-grounded agents per country
Persona source	Author-designed binary labels	Empirical WVS-7 microdata
Metric	MDD (persona-pair distance)	MIS (ℓ_2 to human AMCE)
Correction	None (diagnostic only)	PT-IS + dual-pass reliability gate
Models	3 (GPT-4o, GPT-3.5, Llama-2)	7 checkpoints, 5 families, 2B–70B
Countries	Aggregate (Western vs. Eastern)	20 individual countries
Scenarios	9 dimensions, 10k synthetic	6 dimensions, 310–500 per country
Theory	Partisan sorting (descriptive)	Variance-aware MSE-optimal shrinkage (Prop. 1)

Three findings from Kim et al. [2025] directly inform DISCA’s design.

Persona sensitivity motivates disagreement-driven steering. Kim et al. show that LLMs exhibit moral-decision distances (MDD) $2\text{--}4\times$ larger than humans under contrasting personas. This confirms that persona prompts do shift LLM moral reasoning substantially, validating the premise that within-country persona spread is an informative signal. DISCA converts this sensitivity from a liability (uncontrolled decision shifts) into a feature (a sufficient statistic for correction reliability, Proposition 1).

The partisan sorting phenomenon justifies WVS grounding. Kim et al. find that political orientation produces the largest persona effect in LLMs, far exceeding its effect in human respondents. This “partisan sorting” effect reveals that LLMs over-index on politically salient dimensions when given sociodemographic labels. DISCA avoids this pathology by grounding personas in empirical WVS-7 microdata rather than author-designed binary labels: the ten-dimensional cultural profile (Table 25) spreads influence across religiosity, child-rearing values, social trust, and other axes, so no single dimension monopolises the correction. The leave-one-WVS-dim-out probe (Table 27) confirms that influence is spread across multiple WVS axes rather than monopolised by one—individual cells can be large in either direction, but the load-bearing dimensions vary by MultiTP attribute, with no single WVS descriptor dominating across the board.

Moral flips validate the reliability gate. Kim et al. document that specific personas can entirely reverse LLM preferences on certain moral dimensions (e.g., social status under a progressive persona in GPT-4o). DISCA’s dual-pass reliability gate (Eq. 5) is designed precisely for this regime: when persona-driven corrections are large enough to flip the decision, the two independent IS passes are likely to disagree, and the exponential decay weight r shrinks the correction toward zero. The tail-safety analysis (Table 14) confirms this mechanism reduces harmed cells from 11/120 to 3/120 compared to ungated consensus.

From diagnosis to correction. The fundamental advance of DISCA over the Kim et al. framework is operational: where they diagnose persona sensitivity, we harness it. Their MDD metric measures the distance between two contrasting personas; our MIS measures the distance to the human target and actively reduces it. The variance-aware shrinkage result (Proposition 1) provides the formal bridge: within-panel disagreement (analogous to their MDD) is a sufficient statistic for the correction’s reliability, and the shrinkage estimator converts this diagnostic signal into the MSE-optimal scalar correction.

A16 Extended Limitations

This appendix expands the limitations summarised in §4.2 with technical caveats not covered in the main text.

1225 **WVS-to-trolley linkage.** The mapping from WVS value dimensions to trolley moral dimensions
1226 operates through both direct thematic overlap and indirect modulation (Appendix A11.3). The 10-
1227 dimensional WVS profile explains $R^2 \in [0.55, 0.69]$ of human AMCE variance, and a leave-one-
1228 dimension-out probe identifies load-bearing couplings alongside noisy ones that the IS stage auto-
1229 matically down-weights. We avoid claims of the form “WVS dimension X caused AMCE shift Y ,”
1230 but the ensemble design and the IS noise-filtering mechanism together ensure the method is robust
1231 even when individual WVS–trolley links are imperfect.

1232 **Inter-persona reward assumption.** The reward $r_i = \delta_i - \delta_{\text{base}}$ assumes persona shifts approx-
1233 imate human target directions. This is plausible when the base model represents a WEIRD-biased
1234 prior, but if a persona’s shift is orthogonal to the human target, the signal may mislead. Quantifying
1235 whether persona-consensus directions correlate with country-level human AMCE vectors indepen-
1236 dently of the base model-and when corrections diverge from targets-is open; we do not report that
1237 correlation matrix here.

1238 **Geographic coverage, preprocessing, and WVS vintage.** We report 20 countries out of 100+
1239 MultiTP countries; results depend on the preprocessing pipeline in Appendix A13. WVS Wave 7
1240 (~2017–2022) is a fixed historical slice that may not reflect current preferences in rapidly evolving
1241 societies; persona quality is weakest where coverage is sparse (e.g., ETH and BGD in our panel;
1242 Saudi Arabia is excluded entirely and uses a manual fallback persona, Appendix A11.2). No low-
1243 resource languages (Amharic, Khmer, Yoruba) appear in the panel; cross-lingual generalisation to
1244 such settings remains untested, and one language per country may conflate linguistic and cultural
1245 effects.

1246 **Model quantisation.** 4-bit/Int8 quantisation may introduce logit artefacts affecting IS stability-
1247 observed empirically on Unsloth Phi-4 prior to switching to vLLM BF16 for the headline runs.

1248 NeurIPS Paper Checklist

1249 The checklist is designed to encourage best practices for responsible machine learning research,
1250 addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not re-
1251 move the checklist: **The papers not including the checklist will be desk rejected.** The checklist
1252 should follow the references and follow the (optional) supplemental material. The checklist does
1253 NOT count towards the page limit.

1254 Please read the checklist guidelines carefully for information on how to answer these questions. For
1255 each question in the checklist:

- 1256 • You should answer [Yes], [No], or [N/A].
- 1257 • [N/A] means either that the question is Not Applicable for that particular paper or the
1258 relevant information is Not Available.
- 1259 • Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

1260 **The checklist answers are an integral part of your paper submission.** They are visible to the
1261 reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it
1262 (after eventual revisions) with the final version of your paper, and its final version will be published
1263 with the paper.

1264 The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation.
1265 While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a
1266 proper justification is given (e.g., error bars are not reported because it would be too computationally
1267 expensive” or “we were unable to find the license for the dataset we used”). In general, answering
1268 [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we
1269 acknowledge that the true answer is often more nuanced, so please just use your best judgment and
1270 write a justification to elaborate. All supporting evidence can appear either in the main paper or the
1271 supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification
1272 please point to the section(s) where related material for the question can be found.

1273 IMPORTANT, please:

- 1274 • **Delete this instruction block, but keep the section heading “NeurIPS Paper Check-**
1275 **list”,**
- 1276 • **Keep the checklist subsection headings, questions/answers and guidelines below.**
- 1277 • **Do not modify the questions and only use the provided macros for your answers.**

1278 1. Claims

1279 Question: Do the main claims made in the abstract and introduction accurately reflect the
1280 paper’s contributions and scope?

1281 Answer: [Yes]

1282 Justification: The abstract and §1 state three claims: (i) DISCA is a train-free inference-
1283 time controller that requires no per-country reward models or extra training data; (ii) it
1284 cuts macro MIS by 10–24% on seven open-weight backbones across 20 countries and 6
1285 moral dimensions; (iii) calibration competes with scale (Phi-4 14B beats Llama-3.3-70B in
1286 absolute MIS). All three are substantiated in Table 2 and Figure 3; scope and assumptions
1287 (binary dilemmas with decision-token logits, WVS-7 vintage) are made explicit in §4.2 and
1288 Appendix A16.

1289 Guidelines:

- 1290 • The answer [N/A] means that the abstract and introduction do not include the claims
1291 made in the paper.
- 1292 • The abstract and/or introduction should clearly state the claims made, including the
1293 contributions made in the paper and important assumptions and limitations. A [No] or
1294 [N/A] answer to this question will not be perceived well by the reviewers.
- 1295 • The claims made should match theoretical and experimental results, and reflect how
1296 much the results can be expected to generalize to other settings.
- 1297 • It is fine to include aspirational goals as motivation as long as it is clear that these
1298 goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: §4.2 bounds generalisation with three caveats (alignment to a survey statistic vs. perceived legitimacy; PT value function as aggregation kernel rather than cognitive model; decision-token logit requirement). Appendix A16 expands these with: WVS-to-trolley linkage being indirect, the inter-persona reward assumption, geographic coverage (20 of 100+ MultiTP countries; sparse personas for under-sampled countries), WVS Wave 7 vintage, and 4-bit/Int8 quantisation artefacts on IS stability.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The single theoretical result, Proposition 1 (the MSE-optimal scalar shrinkage of the consensus correction is a closed-form function of within-panel variance), is stated in the main text with all assumptions (i.i.d. persona noise, ℓ_2 risk). The full proof is given in Appendix A2, and the connection to Eq. 3 appears in §3.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.

- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: §3 gives the full algorithm (Algorithm 1, Eqs. 3–5); the full hyperparameter table (Table 18) lists every numerical default with the validation procedure (synthetic 200-scenario pool, transferability check on a 50-scenario MultiTP English subset). All backbones are publicly available open-weight checkpoints. Preprocessing (cap=80, oversample=36), seeds (3), and per-country preprocessing details are documented in Appendices A13, A8.1. We additionally release a self-contained reference implementation (DISCA_Alignment/) with the same defaults and a single-command sweep entry point.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We release DISCA_Alignment/ (anonymised for review; MIT-licensed) as supplemental material. The package is a self-contained slice of the research codebase: a single CLI entry point (run.py) drives the per-model 20-country sweep, with backend selection (vllm / hf_native / unsloth) and explicit flags for the three external data paths (-multitp-data, -wvs-data, -human-amce). The README documents install (Python 3.10+, CUDA 12.x), HF token wiring for gated models, exact commands, and the output

schema (per-country baseline / DISCA CSVs and a comparison file). MultiTP, WVS-7 Wave, and the human AMCE table are publicly available datasets cited in §4; access instructions are reproduced in the README.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: There is no training (the method is inference-time on frozen weights); we therefore detail all decoding and controller settings. §4 specifies the dataset (MultiTP, 20 countries), preprocessing (deduplication, category cap=80, oversample to ≥ 36 per dimension), the seven main backbones with sizes, and the metric definitions (MIS, Pearson r , JSD, win rate). All hyperparameters and how they were chosen are tabulated in Table 18 with the synthetic-pool grid-search procedure described in Appendix A10. Sensitivity sweeps for the temperature family (Appendix A10), category cap (Appendix A13), and the controller knobs (s , λ_{coop} , σ , T_{cat}) appear in Appendix A10.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The headline table (Table 2) reports macro MIS as mean $\pm$ standard deviation across 3 independent seeds (independent IS noise draws; the underlying model weights and data are fixed). A separate multi-seed stability analysis (Appendix A8.1) gives per-country variability, and a bootstrap confidence interval on JSD (± 0.004 , Appendix A10) quantifies the noise floor for distributional metrics. We explicitly state that conclusions are stable well within the bootstrap noise floor; sensitivity sweeps along nine independent axes confirm this.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix A14 (Table 30) reports per-scenario wall-clock on a single H100 GPU: 0.083–0.086 s/scenario for DISCA at $K \in \{64, 128, 256\}$ vs. 0.023 s/scenario for vanilla decoding ($3.57\text{--}3.72\times$ overhead). Backbone-level cost is given in Figure 5: 350 ms/scenario for Phi-4 vs. 1414 ms for Llama-3.3-70B with DISCA. Sub-70B backbones run on a single H100 in BF16 via vLLM; the 70B backbone uses 4-bit Unsloth on the same hardware. The released code’s CLI also runs end-to-end on a Kaggle T4/P100 dual-GPU node for sub-14B models. Exploratory experiments on a small 3-backbone $\times$ 5-country prototyping panel consumed additional compute beyond the reported runs.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The authors have reviewed the NeurIPS Code of Ethics. The work uses only public, ethically-collected datasets (MultiTP [Jin et al., 2025], World Values Survey Wave 7, Moral Machine human AMCEs [Awad et al., 2018]); it conducts no new human-subjects research and introduces no scraped data. Anonymity is preserved in the submission. Negative-impact considerations (risk of encoding harmful majorities) are explicitly discussed in §4.2 with a per-persona utility floor mitigation validated in Appendix A10 (Table 24).

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.

- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: §4.2 discusses both directions. Positive: a train-free, low-overhead controller lets practitioners respect cross-national preference heterogeneity at deployment time without finetuning per country, lowering the barrier to culturally pluralistic alignment. Negative: aligning to majority survey statistics could entrench harmful majorities or marginalise minority views [Zewail et al., 2026], and alignment to a survey statistic is not equivalent to legitimacy [Atari et al., 2023]. We propose and validate a per-persona utility floor that caps how far any single persona’s post-correction utility can drop from vanilla (Appendix A10, Table 24), and we are explicit that DISCA is a research artefact, not a deployment-ready normative system.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: We do not release any pretrained models, generative models, or scraped datasets. The released artefact (DISCA_Alignment/) is an inference-time controller layered over publicly available open-weight LLMs whose existing safeguards remain in place. The method itself produces no novel generative capability beyond that of the underlying models; the per-persona utility floor (Appendix A10, Table 24) is included as an in-method safeguard against majoritarian skew.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.

- 1576 • Datasets that have been scraped from the Internet could pose safety risks. The authors
1577 should describe how they avoided releasing unsafe images.
- 1578 • We recognize that providing effective safeguards is challenging, and many papers do
1579 not require this, but we encourage authors to take this into account and make a best
1580 faith effort.

1581 12. Licenses for existing assets

1582 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
1583 the paper, properly credited and are the license and terms of use explicitly mentioned and
1584 properly respected?

1585 Answer: [\[Yes\]](#)

1586 Justification:

1587 All third-party assets are cited with version information. Datasets: MultiTP [Jin et al.,
1588 2025], the Moral Machine human AMCE table [Awad et al., 2018], and World Values Sur-
1589 vey Wave 7 (inverted CSV v6_0) are used under the terms of the respective public releases.
1590 Models: Phi-4 and Phi-3.5-mini (MIT), Llama-3.3-70B (Llama-3 Community License),
1591 Qwen-2.5 / Qwen3-VL (Apache-2.0 / Qwen license), Gemma family (Gemma terms of
1592 use), Magistral-Small-2509 (Apache-2.0). Code dependencies (Transformers, vLLM, Un-
1593 sloth, bitsandbytes) are listed with version pins in the released `requirements.txt`; we use
1594 them under their stated licenses (mostly Apache-2.0 / MIT). The `compute_ACME` reference
1595 implementation we follow is credited to Jin et al. [2025] in Appendix A12.

1596 Guidelines:

- 1597 • The answer [\[N/A\]](#) means that the paper does not use existing assets.
- 1598 • The authors should cite the original paper that produced the code package or dataset.
- 1599 • The authors should state which version of the asset is used and, if possible, include a
1600 URL.
- 1601 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 1602 • For scraped data from a particular source (e.g., website), the copyright and terms of
1603 service of that source should be provided.
- 1604 • If assets are released, the license, copyright information, and terms of use in the pack-
1605 age should be provided. For popular datasets, `paperswithcode.com/datasets` has
1606 curated licenses for some datasets. Their licensing guide can help determine the li-
1607 cense of a dataset.
- 1608 • For existing datasets that are re-packaged, both the original license and the license of
1609 the derived asset (if it has changed) should be provided.
- 1610 • If this information is not available online, the authors are encouraged to reach out to
1611 the asset’s creators.

1612 13. New assets

1613 Question: Are new assets introduced in the paper well documented and is the documenta-
1614 tion provided alongside the assets?

1615 Answer: [\[Yes\]](#)

1616 Justification: We release the DISCA reference implementation (`DISCA_Alignment/`)
1617 as a new code asset under the MIT License. The package ships with: a top-level
1618 README documenting installation, backend selection (vLLM / HF native / Unsloth), the
1619 three external data flags, environment variables, the quick-start command, and the output
1620 schema; a flat `disca/` module with one-purpose files (controller, dual-pass core, runner,
1621 AMCE/metrics, persona generation, multilingual scenarios, model loader); pinned depen-
1622 dencies in `requirements.txt`; and a single CLI entry point (`run.py -help`). The sub-
1623 mission archive is anonymised. All hyperparameter defaults match Table 18 in the main
1624 paper.

1625 Guidelines:

- 1626 • The answer [\[N/A\]](#) means that the paper does not release new assets.
- 1627 • Researchers should communicate the details of the dataset/code/model as part of their
1628 submissions via structured templates. This includes details about training, license,
1629 limitations, etc.

- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper conducts no new crowdsourcing or human-subjects research. Human reference data is reused from the existing public Moral Machine experiment [Awad et al., 2018] as packaged by MultiTP [Jin et al., 2025], and from the World Values Survey Wave 7 public microdata; we make no new contact with participants and recruit no annotators.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: No new human-subjects research is conducted. The work is a secondary analysis of pre-existing public datasets (Moral Machine human AMCEs, MultiTP scenarios, WVS-7 microdata), each of which carries its own ethics review at the original collection point; no further IRB approval applies.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: LLMs are central to the work in two ways, both fully described. (i) Frozen open-weight LLMs (Phi-4, Llama-3.3-70B, Magistral-24B, Qwen3-VL-8B, Qwen2.5-7B, Phi-3.5-mini, Gemma-4-E2B) are the *subject* that DISCA steers; their identities, sizes,

1684 and decoding settings are documented in §4 and Appendix A6. (ii) For the open-ended
1685 extension (§3.4), an LLM parses free-form responses into a (choice, confidence) pair to
1686 recover a pseudo-logit-gap; this auxiliary use is documented in that section. A separate
1687 synthetic 200-scenario tuning pool generated by GPT-4 is disclosed in Appendix A10 along
1688 with its limitations. LLMs were not used to generate the paper’s claims, results, or proofs.

1689 Guidelines:

- 1690 • The answer [N/A] means that the core method development in this research does not
1691 involve LLMs as any important, original, or non-standard components.
- 1692 • Please refer to our LLM policy in the NeurIPS handbook for what should or should
1693 not be described.